

Tree Structures (v5)

Juan Manuel Gimeno Illa

2025/2026

Bibliography

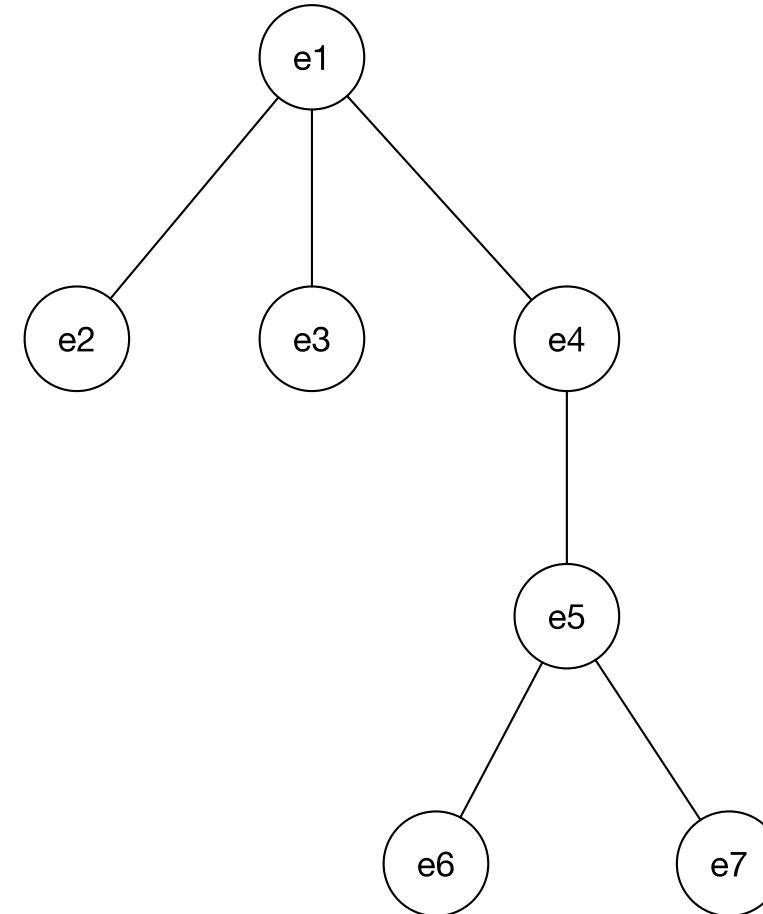
- Robert Sedgewick & Kevon Wayne, Algorithms - Fourth Edition, Pearson Education Inc. (2011)
 - 3.3 (BSTs, 23-Trees and RB-Trees), 6 (B-Trees)
- Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest & Clifford Stein, Introduction to Algorithms – Fourth Edition, Massachusetts Institute of Technology (2022)
 - 6.1, 10.3, 18 (a variant of B-trees that splits descending), B.5
- Josep Maria Ribó, Apropament a les estructures de dades des del programari lliure. Edicions de la Universitat de Lleida. 2018.
 - Chapter 3, 6 (specially B-trees)

Importance of tree structures

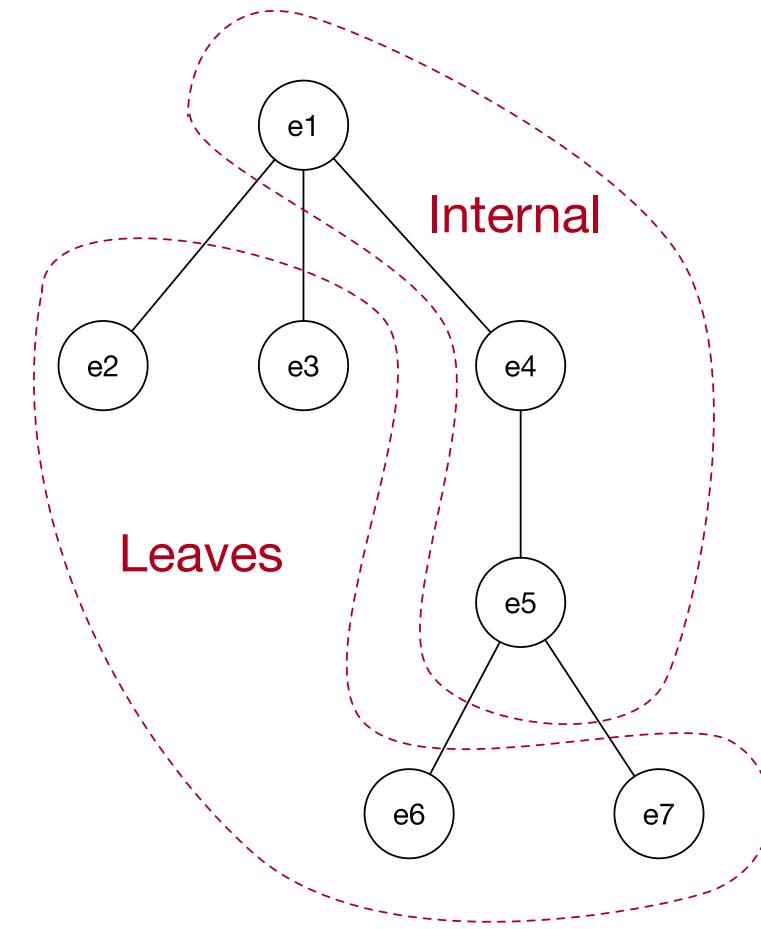
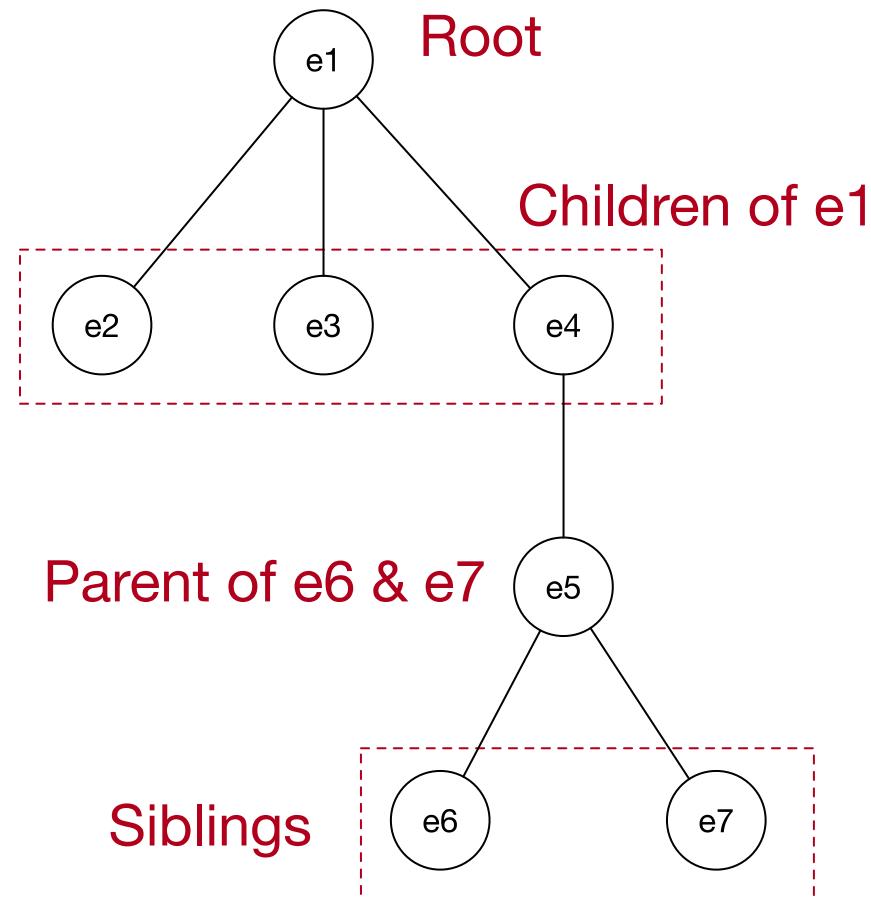
- In computer science, and in data structures in particular, trees are one of the most important structures
- Even that, in most data libraries, we don't find an interface, or a direct implementation, of a data structure named tree
- Why?
 - Because the tree is used as the **representation** (implementation device) used by other data structures to be **efficient**.
 - For example, PriorityQueue uses a heap, which is a type of tree implemented with an array

A definition of Trees

- A **tree**, as a data structure, is a **container of elements** of the same type
 - It can be **empty** (with no elements)
 - Or be a composition of
 - An element (the **root** of the tree)
 - Zero or more **disjoint** (sub)trees named its **children**
- **NOTE:** Usually, empty subtrees are not shown.
- **NOTE:** Sometimes we conflate the nodes of the tree with its elements

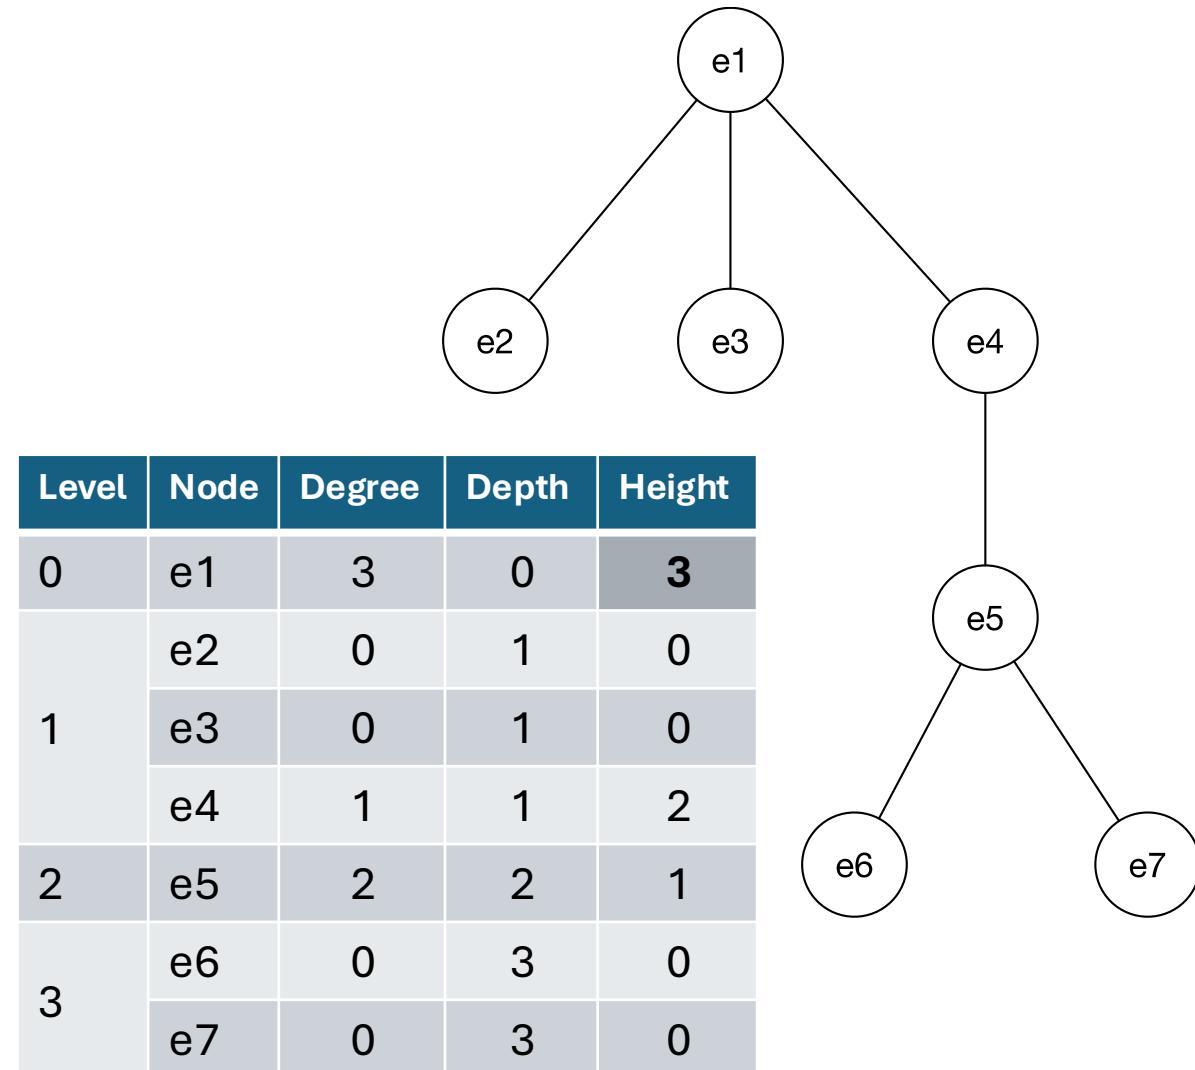


Names, names and more names



Names, names and more names

- The **degree of a node** is the number of non-empty children
- The **depth of a node** is the number of edges from the root to the node.
 - **Level k** is the set of nodes at **depth k**
- The **height of a node** is the number of edges from the node to the deepest leaf.
- The **height of a tree** is the **height of its root**.
 - It's equal to the **max depth** of any node
 - An **empty tree** has height **-1**
- The **size of a tree** is the number of nodes



Trees and recursion

- As the very definition of tree is **recursive**, most of the time, the natural and simplest way to program methods over trees is **recursion**
 1. We must define a **base** case, i.e. a case that can be solved **without recursive calls**.
Usually, the base case is the empty tree.
 2. In the **recursive** case we call the same functions over “**smaller**” trees.
Usually, the recursive calls are made on the children.
 3. To be **efficient**, and this can be difficult sometimes, we should **avoid duplicities** (i.e. calling the same function more than once on the same tree).

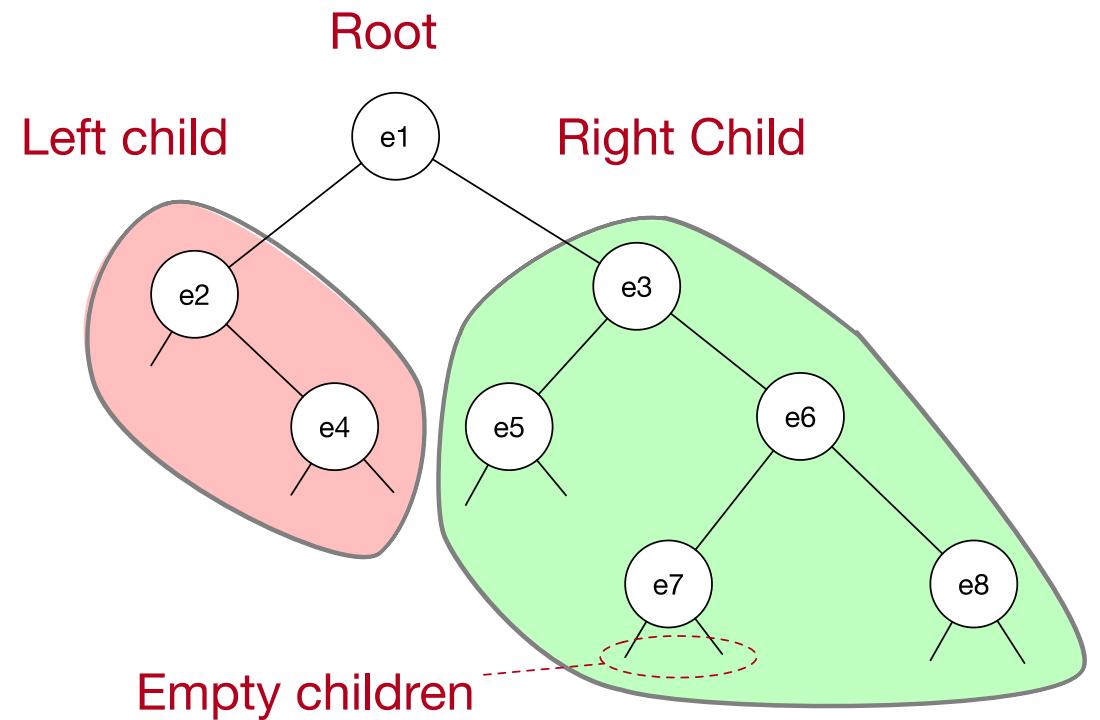
Types of trees

- **Rooted** trees
 - There is node which is the root of the tree
 - Each node has associated a list of its children (can be empty for leaves)
 - The “position” of each (non-empty) children in the list is not relevant
- **Ordered** trees are rooted trees in which
 - Children have an order: first child, second child, ...
 - But if the second disappears, the third occupies its place and so on
- **Positional** trees are ordered trees in which
 - Each child has a definite position
 - So, if the second disappears, its occupied by an empty child
- **N-ary** trees are positional trees in which
 - Any node has at-most N children (some of them can be empty)
- The most important trees are **binary trees** which are **2-ary trees**

Binary Trees

Binary tree

- A **binary tree** is a container of elements of **type E** that can be:
 - An **empty** tree, that is, without elements
 - A **non-empty** tree composed by
 - An element of type E named its **root**
 - Two disjoint binary subtrees (which can be empty) named **left child** and **right child**



NOTE: Usually, empty children are not shown

Properties of binary trees

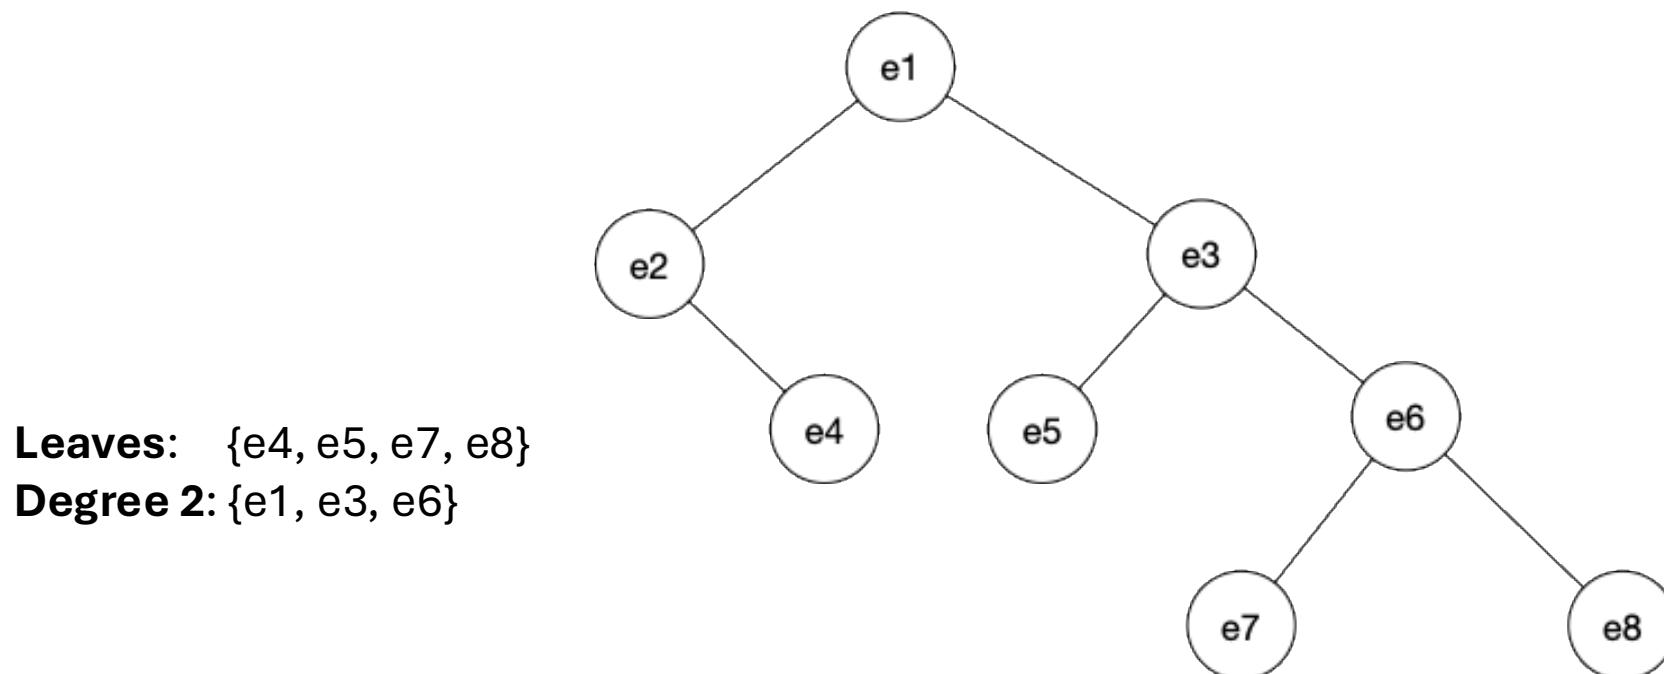
- Binary trees have some **properties** that can be easily **proved by induction**.
- *The maximum number of elements at level k is 2^k*
 - **Base case:**
 - the root, which is at level 0 has a maximum number of elements of $2^0 = 1$, which is true
 - **Inductive case:**
 - the maximum at level $k + 1$ is achieved by adding two children to each of the nodes at level k which, by induction hypotheses, has 2^k maximum nodes
 - so, the maximum number at level $k + 1$ is $2 * 2^k = 2^{k+1}$
 - QED

Properties of binary trees

- *The maximum number of elements in a tree of height h is $2^{h+1} - 1$*
 - **Base case:**
 - An empty tree has height -1 and has $2^{-1+1} - 1 = 2^0 - 1 = 1 - 1 = 0$ elements
 - **Inductive case:**
 - The maximum number of elements in a tree of height $h + 1$ is accomplished by combining a root and two trees of maximum elements of height h
 - This tree has $1 + 2 * (2^{h+1} - 1) = 1 + 2^{h+2} - 2 = 2^{h+2} - 1$ elements
 - QED

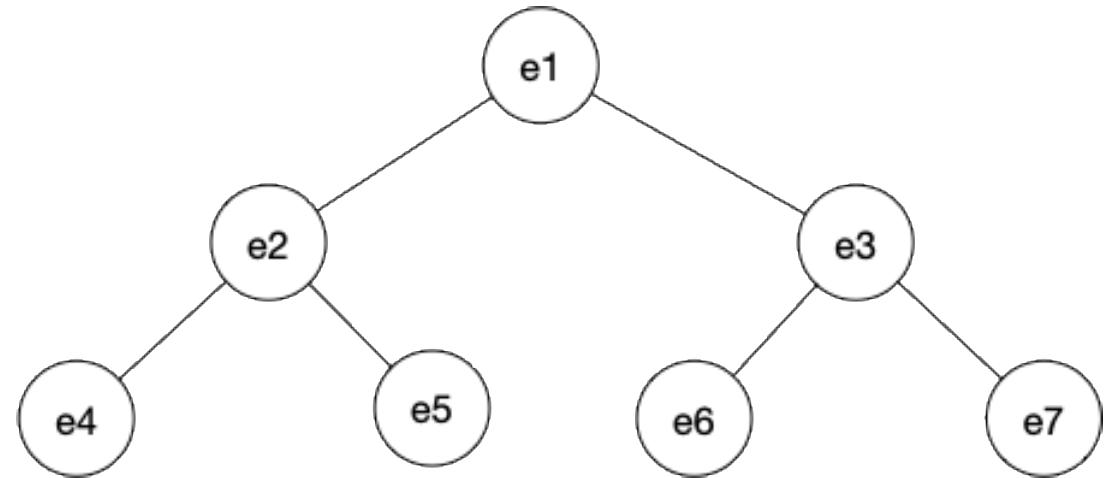
Properties of binary trees

- The number of leaves (degree 0) in a non-empty binary tree is one more than the nodes of degree 2 (i.e. with two children)*



Binary trees

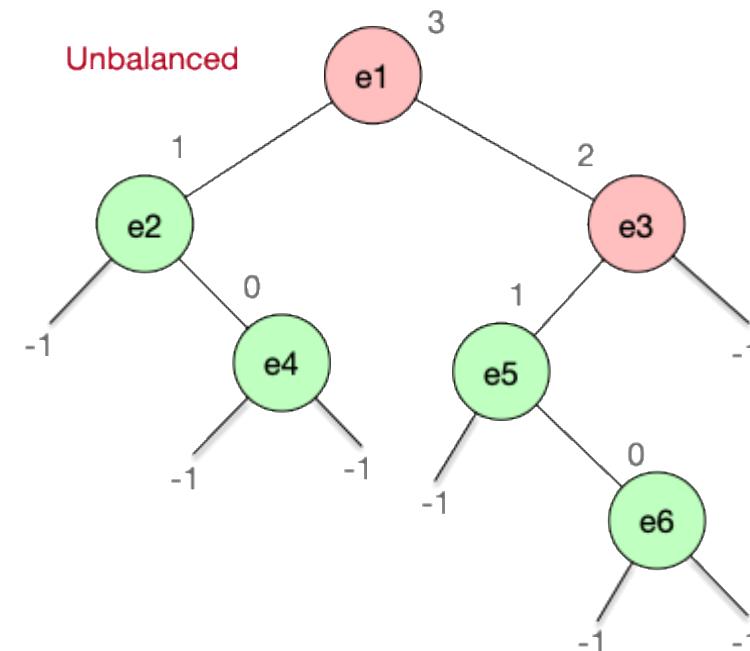
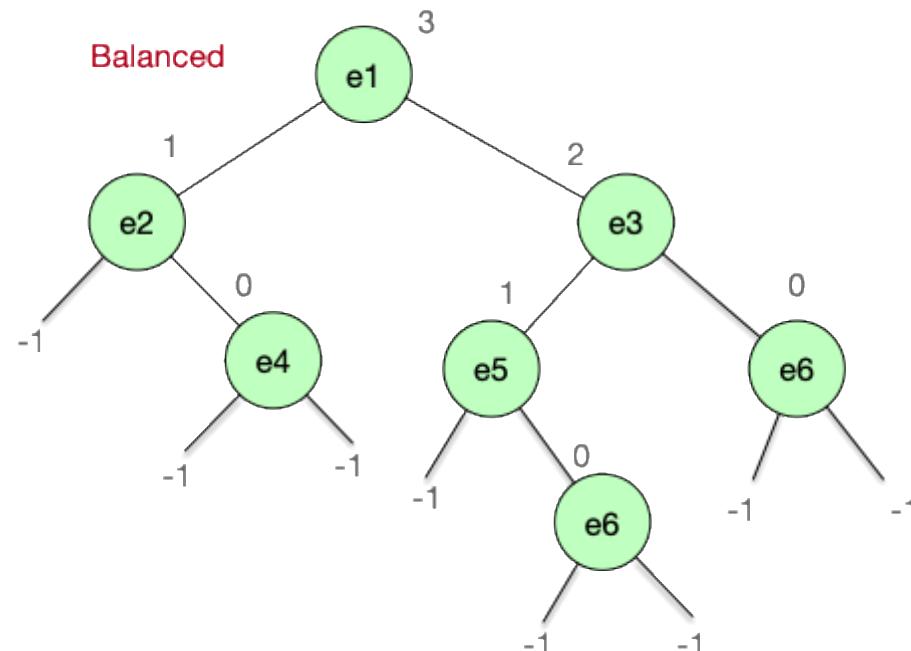
- A **binary tree** is said to be **complete** if
 - It has the maximum number of nodes for its height
- In a **complete binary tree**
 - There are 2^h leaves
 - **The height of a complete tree of size n is $\Theta(\log_2 n)$**
 - There are $2^h - 1$ internal nodes
- **NOTES:**
 - This can be generalized easily to k-ary trees
 - Some authors consider that the last level may not be complete, but with all nodes to the left



Height = 2
Leaves = $2^2 = 4$
Internal = $2^2 - 1 = 4 - 1 = 3$

Binary trees

- A **binary tree** is said to be **balanced** when, for each node the heights of its left and right child differ in at most 1 and both children are balanced
- The **height of a balanced binary tree of size n is $\Theta(\log_2 n)$**
- **NOTE:** This property also generalizes to k-ary trees.

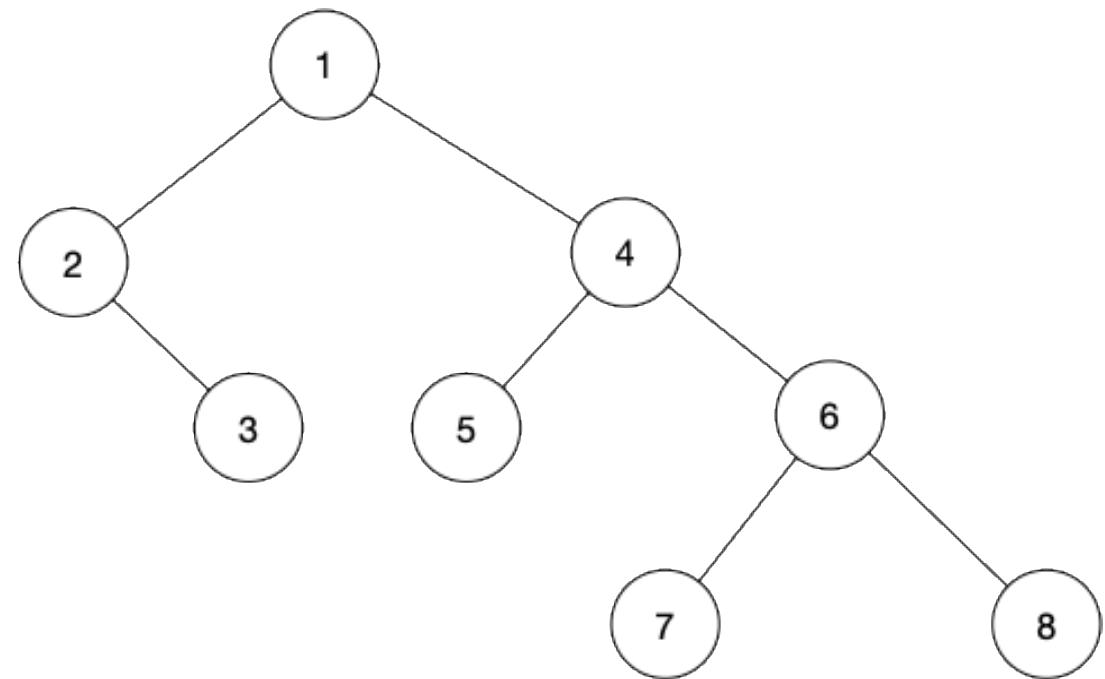


Binary tree traversals

- One of the things we can do with a tree is **traverse** it
 - That is, obtaining a sequence (usually represented as a List) of all the elements in the tree
- But, in which order?
 - Unlike with sequential data structures, here each element has no definite position (i.e. the first element, the second element, ...)
- But there are four “natural” orders
 - Three *recursive* (or in-depth) traversals named: **pre**-order, **in**-order and **post**-order, which are defined recursively and differ in when the **root** node is visited
 - A fourth *non-recursive* traversal named **level**-order

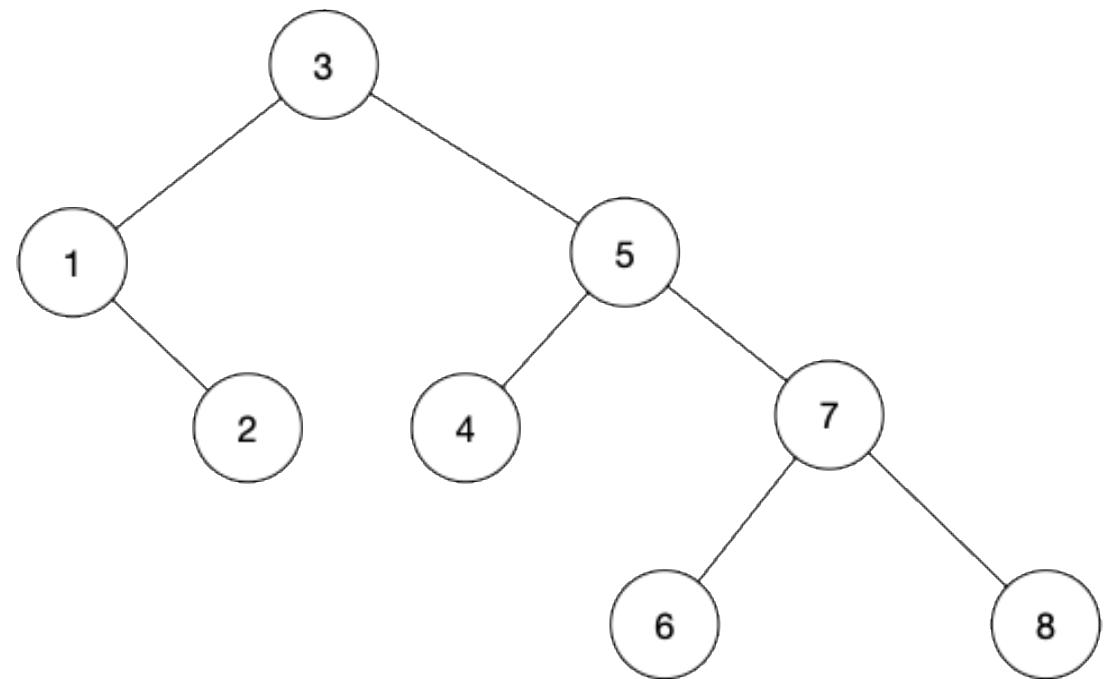
Binary tree traversals

- **Pre-order:**
 1. the **root** of the tree
 2. the left child in pre-order
 3. the right child in pre-order



Binary tree traversals

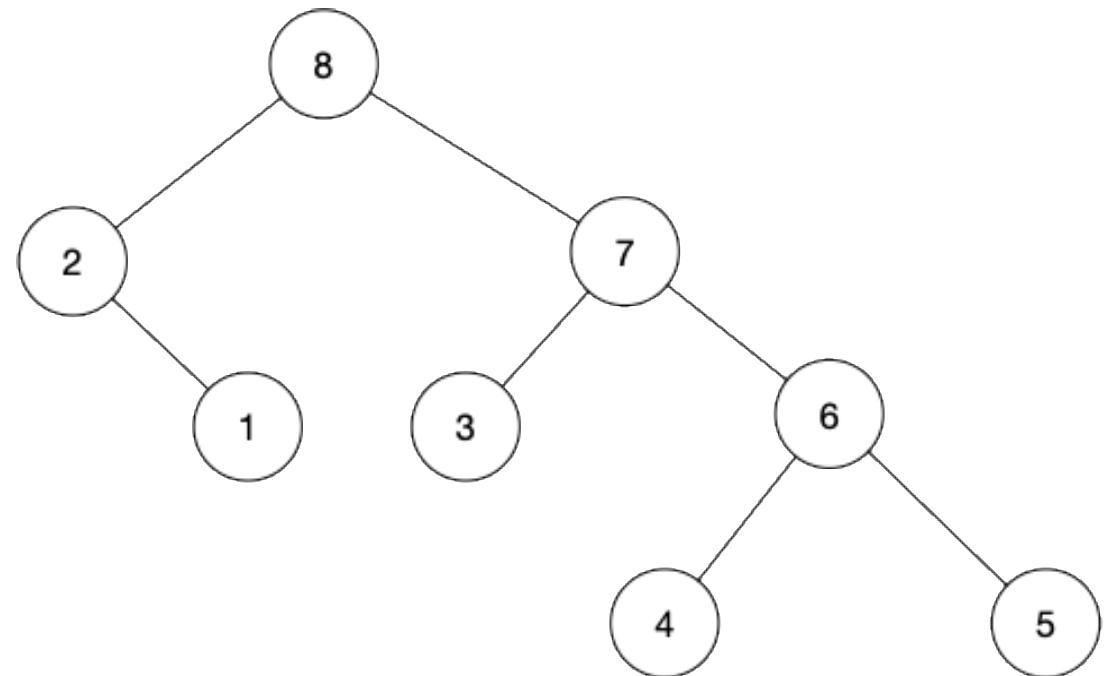
- **In-order:**
 1. the left child in in-order
 2. the **root** of the tree
 3. the right child in in-order



Binary tree traversals

- **Post-order:**

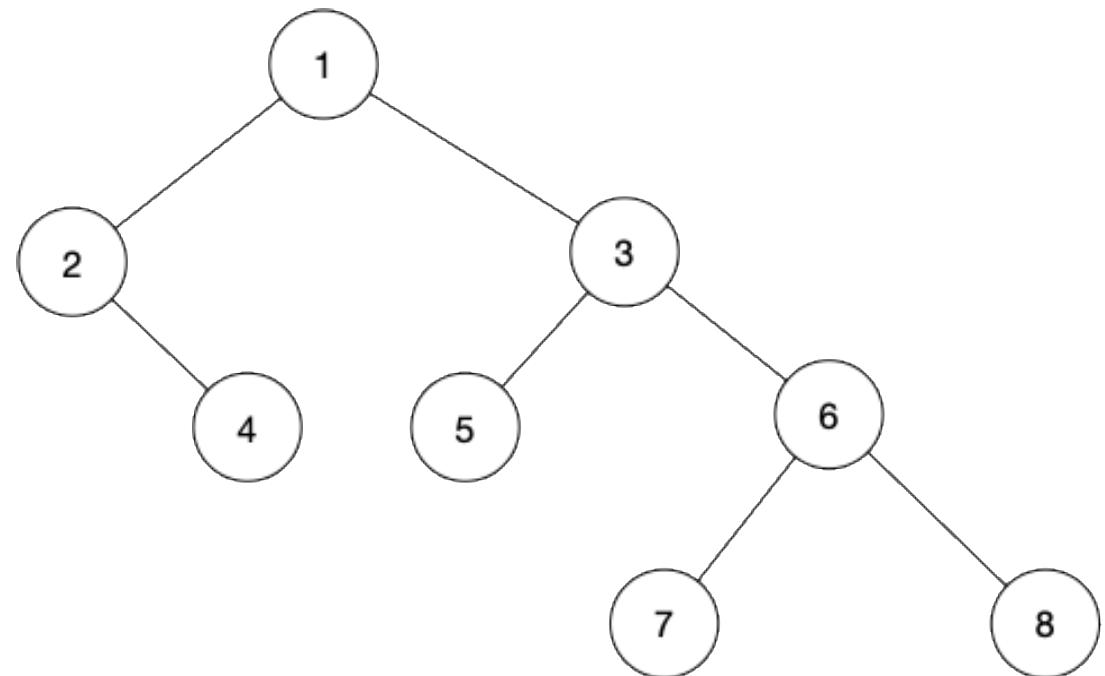
- The left child in post-order
- The right child in post-order
- The **root of the tree**



Binary tree traversals

- **Level-order:**

- First level 0
- Then level 1
- Then level 2
- ...
- (each level from left to right)

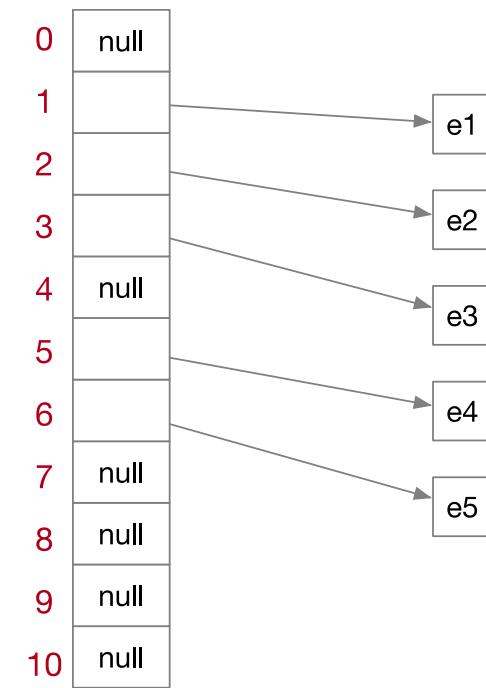
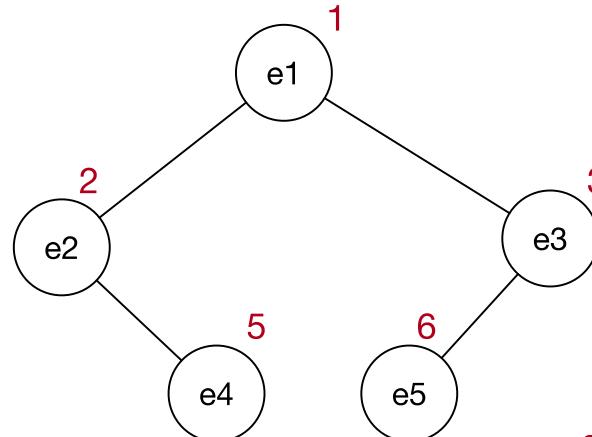


Binary trees

- The first representation we will consider, simply using an **array**
 - But, as we will see, it is **only reasonable in a limited number of cases**
 - Mainly to be used in the implementation of **heaps** and in the **heap-sort** algorithm and in the implementation of **priority queue**
- **In some cases**, as in the examples we will show, the **0-index** position in the array is **not used** (to be able to implement faster arithmetic)
 - For example, in priority queues some implementations ignore the 0-index position
 - But, once you've understood the principles behind the design, moving from one implementation to the other is easy

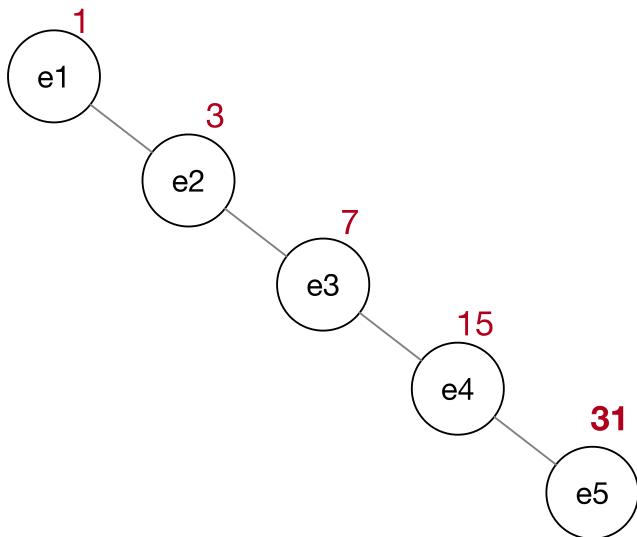
Binary Trees

- Root is at index **1**
- For a **node** at index ***i***
 - Left child: at $2 * i$
 - Right child: at $2 * i + 1$
- How do we know if the **child exists?**
 - Because the **index is out of the array**; or the **value** at its index is **null**
 - So, in this implementation, **null values are not allowed** for the **elements**



Binary Trees

- Why is this implementation not valid in the general case?
 - Because for tree with n nodes, it may need an array of size 2^n

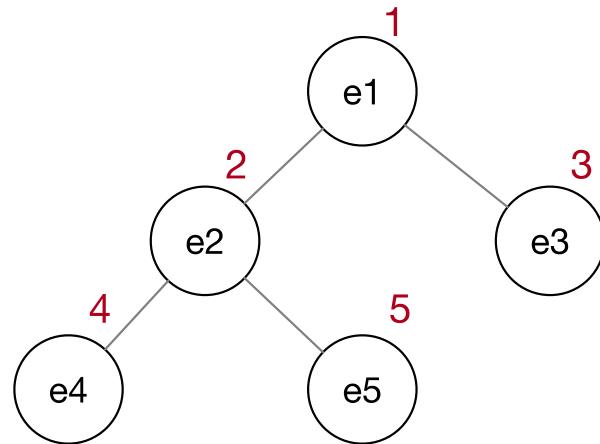


You need an array of $2^5 = 32$ positions !!

- and only 5 of them won't be null

Binary Trees

- But, for trees that are complete (even allowing for an incomplete last level of nodes aligned to the left), it is a very compact and efficient implementation
 - Amortizing resizing operations on the array

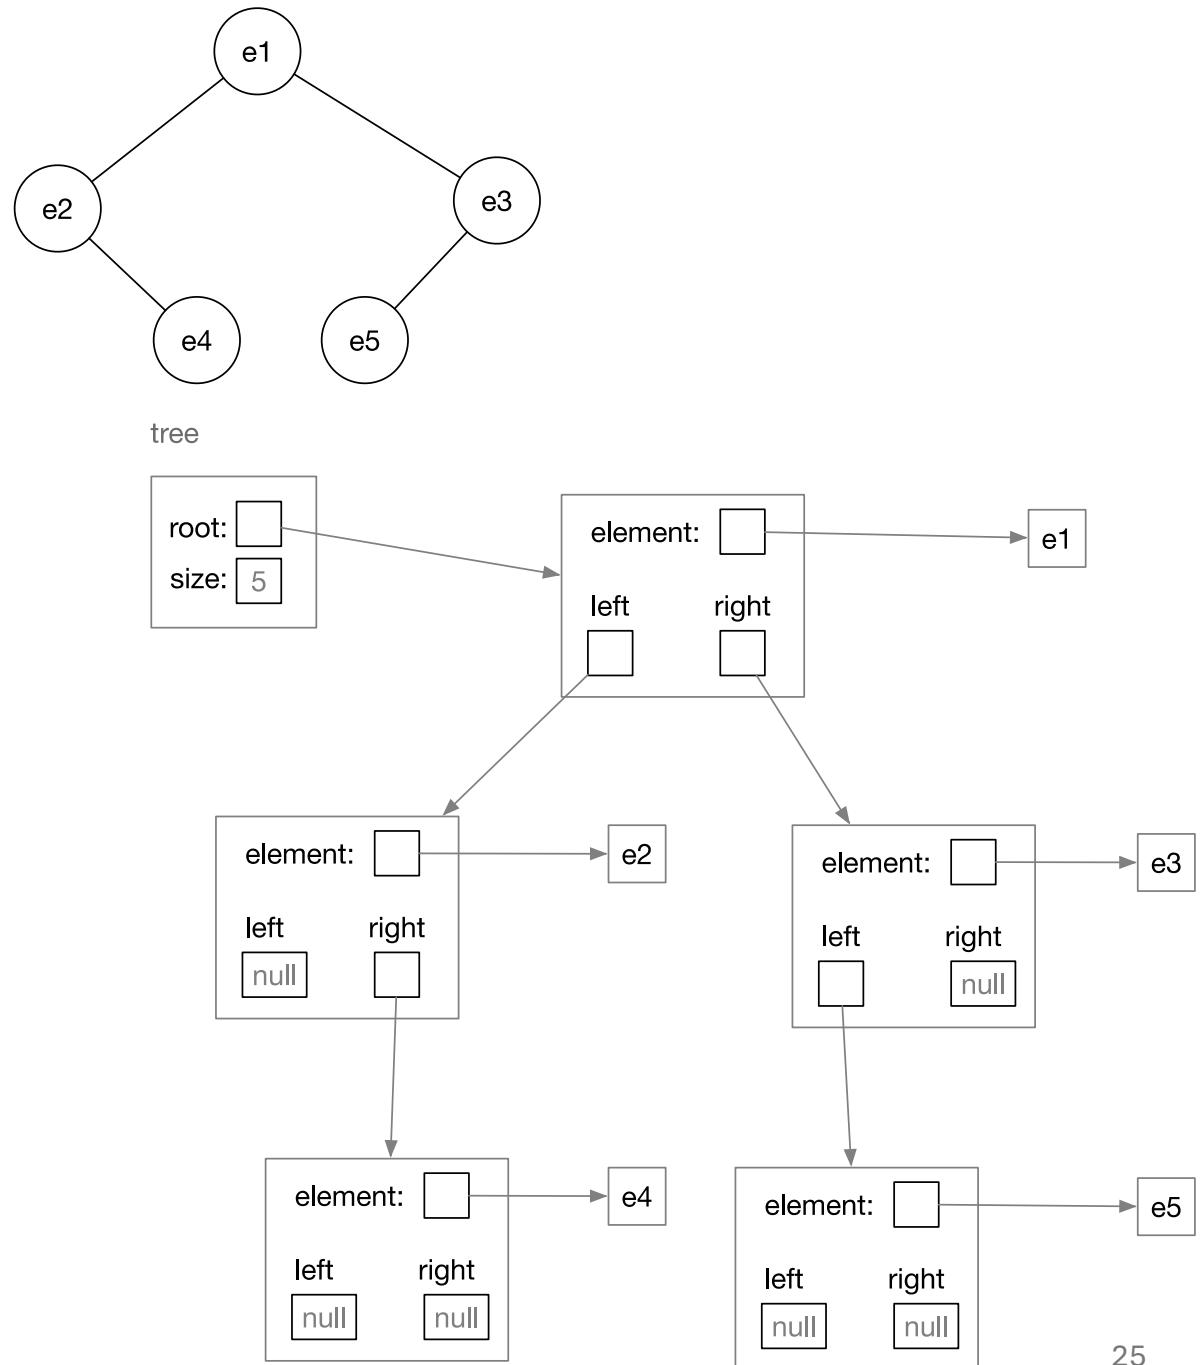


You need an array of only $5 + 1 = 6$ positions !!

- and only 1 of them is null

Binary Trees

- The most used representation of trees uses a **linked structure of nodes**
- A tree has a reference to the node that represents its root
 - **null** if the tree is **empty**
- Each node, apart from the **element** it stores, has two references:
 - one to the **left child**
 - one to the **right child**
- Extra information, such as tree **size**, can be **stored in the tree**



Binary Trees

- As we have said before, trees are used mostly to get efficient implementation of other data structures
 - So, for instance, the JCF does not contain a data structure that presents an interface about trees
- We will add such an interface, and outline a possible implementation that uses this linked representation
- And, when designing this implementation, we will consider important topic such as
 - using more space to not waste time recomputing things
 - mutability and the necessity of copies
 - the java Clonable interface to make copies

Binary Trees

- It's difficult to decide the methods in a `BinaryTree` interface
 - As we have said, most trees are used as implementation devices for other types
 - We have selected some methods that will allow us to comment some trade-offs in the implementation

```
public interface BinaryTree<E> {
```

```
    E root();  
    BinaryTree<E> left();  
    BinaryTree<E> right();
```

```
    default boolean isEmpty() {  
        return size() == 0;  
    }
```

```
    int size();  
    int height();
```

```
    E replaceRoot(E e);  
    void removeLeft();  
    void removeRight();
```

```
    List<E> preOrder();  
    List<E> inOrder();  
    List<E> postOrder();  
    List<E> levelOrder();
```

```
}
```

Binary Trees

- The first important decision we have made is to define the trees as **modifiable**
 - Once a tree is created, we can in-place replace its root or delete (made empty) any of its children
- In the implementation of this type
 - We want the implementations of the **constructors** and **accessors** to be **efficient**, that is, $O(1)$
 - And for the two properties
 - **Size**: needs to be **efficient**
 - **Height**: **no** need of efficiency

```
public interface BinaryTree<E> {
```

```
E root();  
BinaryTree<E> left();  
BinaryTree<E> right();
```

```
default boolean isEmpty() {  
    return size() == 0;  
}
```

```
int size();  
int height();
```

```
E replaceRoot(E e);  
void removeLeft();  
void removeRight();
```

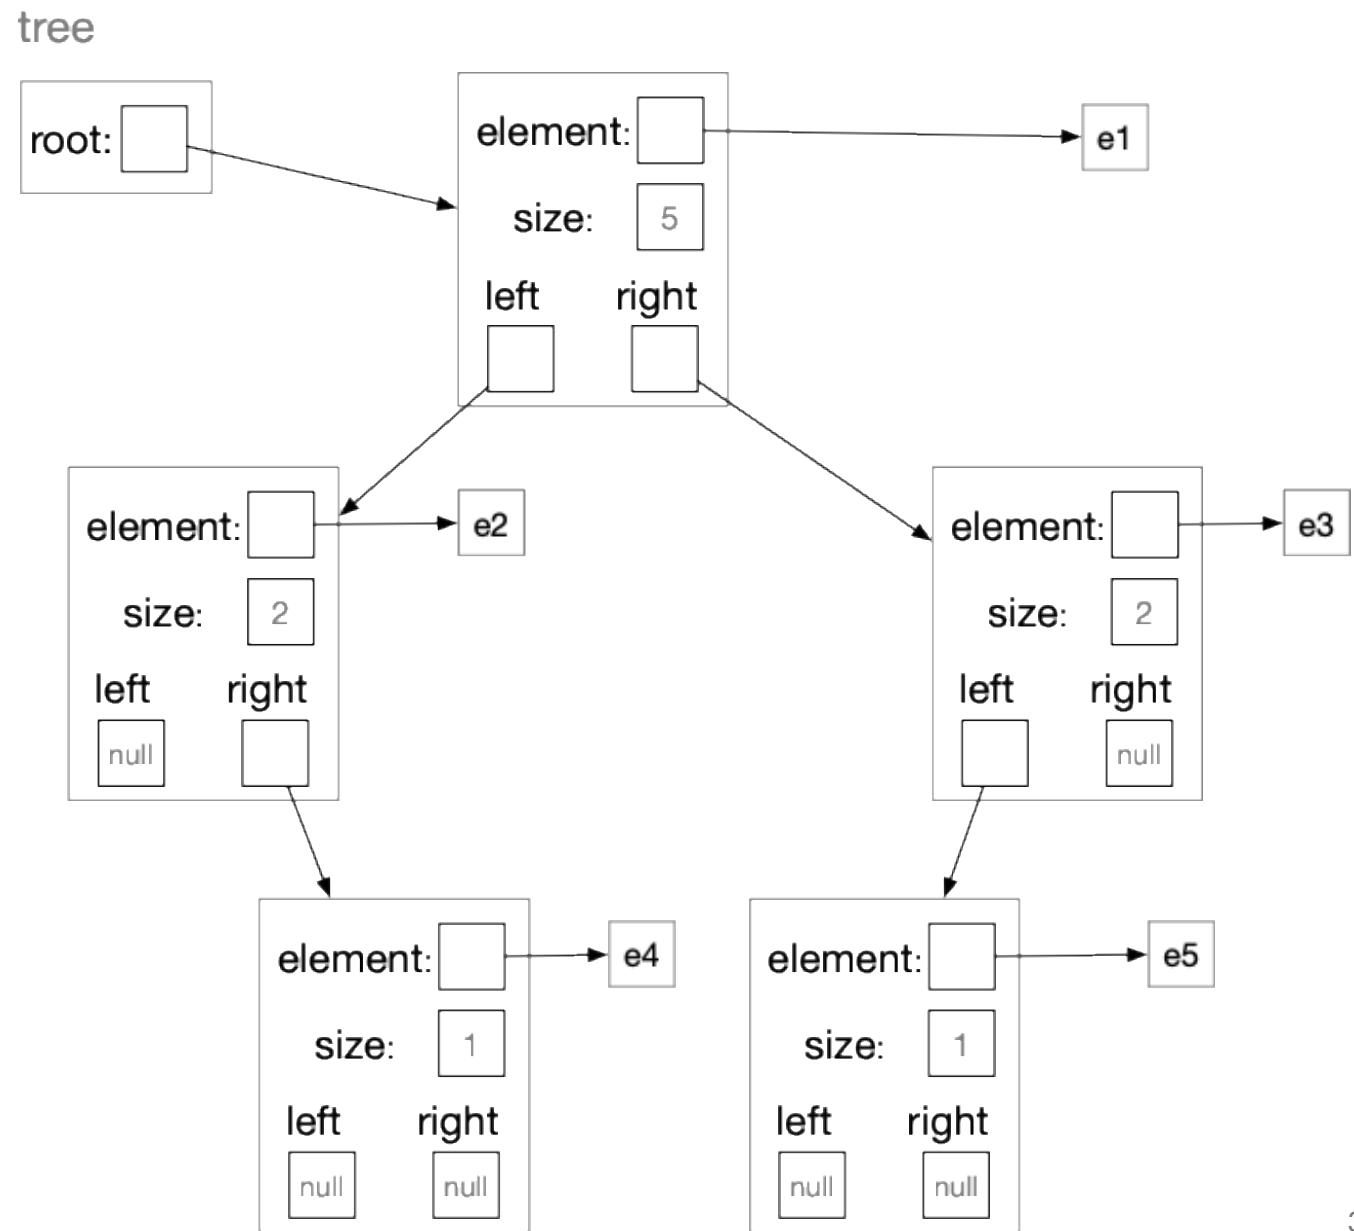
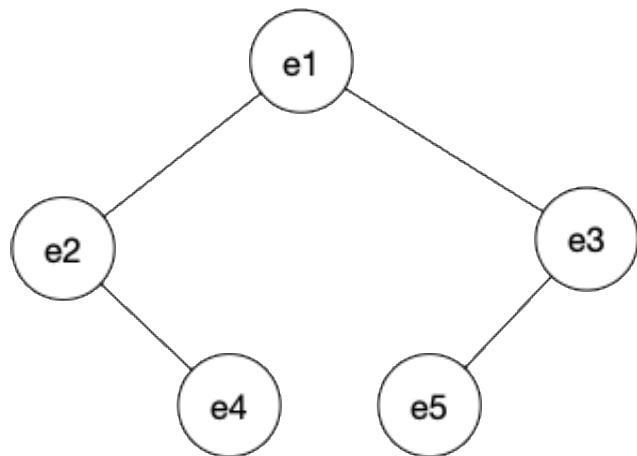
```
List<E> preOrder();  
List<E> inOrder();  
List<E> postOrder();  
List<E> levelOrder();
```

```
}
```

Binary Trees

- The **representation** shown before can be **analysed** given the **efficiency restrictions** we have outlined
- **PROS:**
 - Having the size precomputed in the tree, gives us $O(1)$ for the size property
- **CONS:**
 - But, e.g. returning the left child is $O(n)$, where n is the size of the left child, because the size of the subtree must be computed by traversing it
 - So, the first implementation decision is to make the **size** to be **cached** at the **nodes**
 - NOTE: It's very important to be able to **reason** about this kind of things **before even writing a single line of code**

Binary Trees



```

public class LinkedBinaryTree<E> implements BinaryTree<E> {

    private Node<E> root;                                // Constructors

    private static class Node<E> {
        Node<E> left;
        E element;
        Node<E> right;
        int size;

        Node(Node<E> left, E element, Node<E> right) {
            this.left = left;
            this.element = element;
            this.right = right;
            this.size = 1 + Node.size(left) + Node.size(right);
        }

        static int size(Node<?> node) {
            return node == null ? 0 : node.size;
        }
    }

    public LinkedBinaryTree() { root = null; }

    public LinkedBinaryTree(
        LinkedBinaryTree<E> left,
        E elem,
        LinkedBinaryTree<E> right) {
        Node<E> leftChild = left == null ? null : left.root;
        Node<E> rightChild = right == null ? null : right.root;
        root = new Node<>(leftChild, elem, rightChild);
    }

    private LinkedBinaryTree(Node<E> root) {
        this.root = root;
    }

    // ...
}

```

```
// Accessors
```

```
@Override  
public E root() {  
    if (root == null)  
        throw new NoSuchElementException("root of empty tree");  
    return root.element;  
}
```

```
@Override  
public LinkedBinaryTree<E> left() {  
    if (root == null)  
        throw new NoSuchElementException("left child of empty tree");  
    return new LinkedBinaryTree<>(root.left);  
}
```

```
@Override  
public LinkedBinaryTree<E> right() { ... }
```

```
// Properties
```

```
@Override  
public boolean isEmpty() {  
    return root == null;  
}  
  
@Override  
public int size() { return Node.size(root); }  
  
@Override  
public int height() { return Node.height(root); }
```

```
// Node.height implementation  
static int height(Node<?> node) {  
    if (node == null)  
        return -1;  
    else  
        return 1 + Math.max(height(node.left), height(node.right));  
}
```

```
// Modifiers
```

```
@Override
```

```
public E replaceRoot(E newElement) {  
    if (root == null)  
        throw new NoSuchElementException("the empty tree has no root to replace");  
    E oldElement = root.element;  
    root.element = newElement;  
    return oldElement;  
}
```

```
@Override
```

```
public void removeLeft() {  
    if (root == null)  
        throw new NoSuchElementException("Empty tree");  
    root.size -= Node.size(root.left);  
    root.left = null;  
}
```

```
@Override
```

```
public void removeRight() { ... }
```

```
// Traversals
```

```
@Override  
public List<E> preOrder() {  
    List<E> result = new ArrayList<>(size());  
    if (root != null)  
        root.preOrder(result);  
    return result;  
}
```

```
// Node.preOrder
```

```
void preOrder(List<E> result) {  
    result.add(element);  
    if (left != null)  
        left.preOrder(result);  
    if (right != null)  
        right.preOrder(result);  
}
```

```
// Methods overridden from Object
```

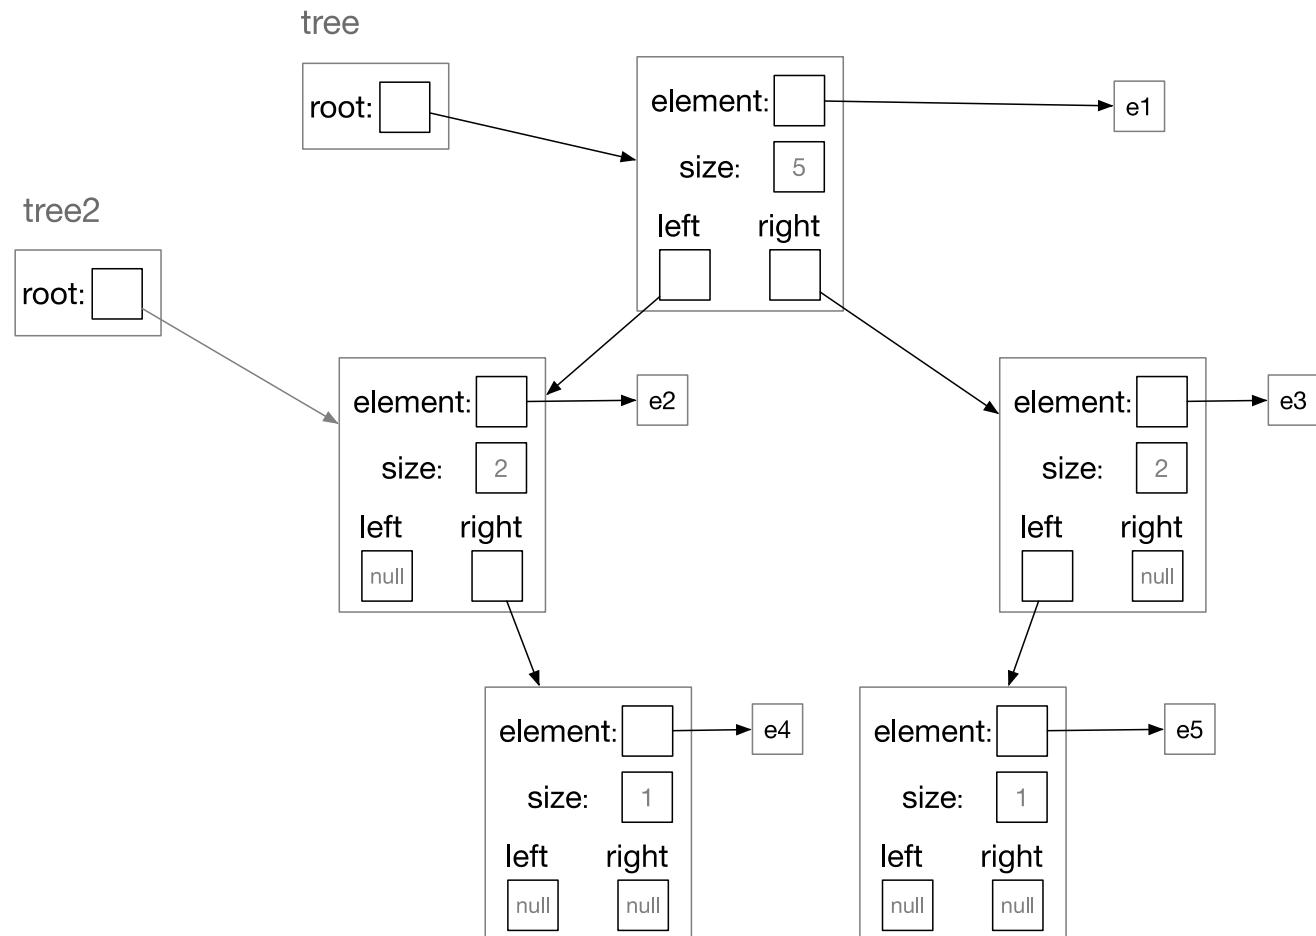
```
@Override  
public boolean equals(Object o) {  
    if (this == o)  
        return true;  
    if (!(o instanceof LinkedBinaryTree<?> bt))  
        return false;  
  
    return Node.equals(root, bt.root);  
}  
  
// Node.equals  
static boolean equals(Node<?> node1, Node<?> node2) {  
    if (node1 == null || node2 == null)  
        return node1 == node2;  
    else  
        return node1.size == node2.size  
            && Objects.equals(node1.element, node2.element)  
            && equals(node1.left, node2.left)  
            && equals(node1.right, node2.right);  
}
```

Binary Trees

- Some comments on the implementation:
 - As we have said we consider that height is not called as often as size, so we do not need to cache it in each node
 - The E in LinkedBinaryTree<E> is not the same E in Node<E> because the later is static, but it's customary to use the same letter
 - The auxiliary methods in class Node are mostly static cause they deal with null references
 - If they were non-static the client code must deal with the nulls
 - The traversal method creates the list for the result, and this list is shared by all the recursive calls
 - If each recursive call returned a list, most of the execution time would be dedicated to copying lists

Binary Trees

- To make the left() operation $O(1)$, the resulting tree and the original one share nodes
- This wouldn't be a problem if the trees and the elements were **unmodifiable**
- This is not the case here:
 - **due to operations on the tree**
 - or to operations on the elements (we cannot control them, because we do not know what E is)



Binary Trees

- If we want to protect the users from the problems derived from sharing, we could
 - make a copy of the whole left tree and return a new tree pointing to it
 - but this would be $O(n)$, where n is the size of the left tree
 - and this could be not needed in all cases !!
- What can we do?
 - add a copy mechanism that will allow a client of the class to get a copy of a tree that doesn't share nodes with the original one
- NOTE:
 - this won't solve the problems with modifiable elements
 - avoiding this kind of problems, by forbidding modifications, is one of the big advantages of functional programming

Binary Trees

- One possibility is to add a method named copy, with signature:

```
public LinkedBinaryTree<E> copy() { ... }
```

- But in Java there exists a construction specifically designed to do that: the marker interface **Cloneable**
 - it's a marker interface, i.e., it defines no methods
 - we have seen it before on the implementations of lists
 - uses extra-linguistic features, i.e. not implemented with the language, but by special machinery in the virtual machine
 - it makes a field-for-field copy of an object (shallow copy)
 - NOTE: Its use is controversial, and most authors recommend to use a method such as copy, but as Java programmers we must know how it works.

Binary Trees

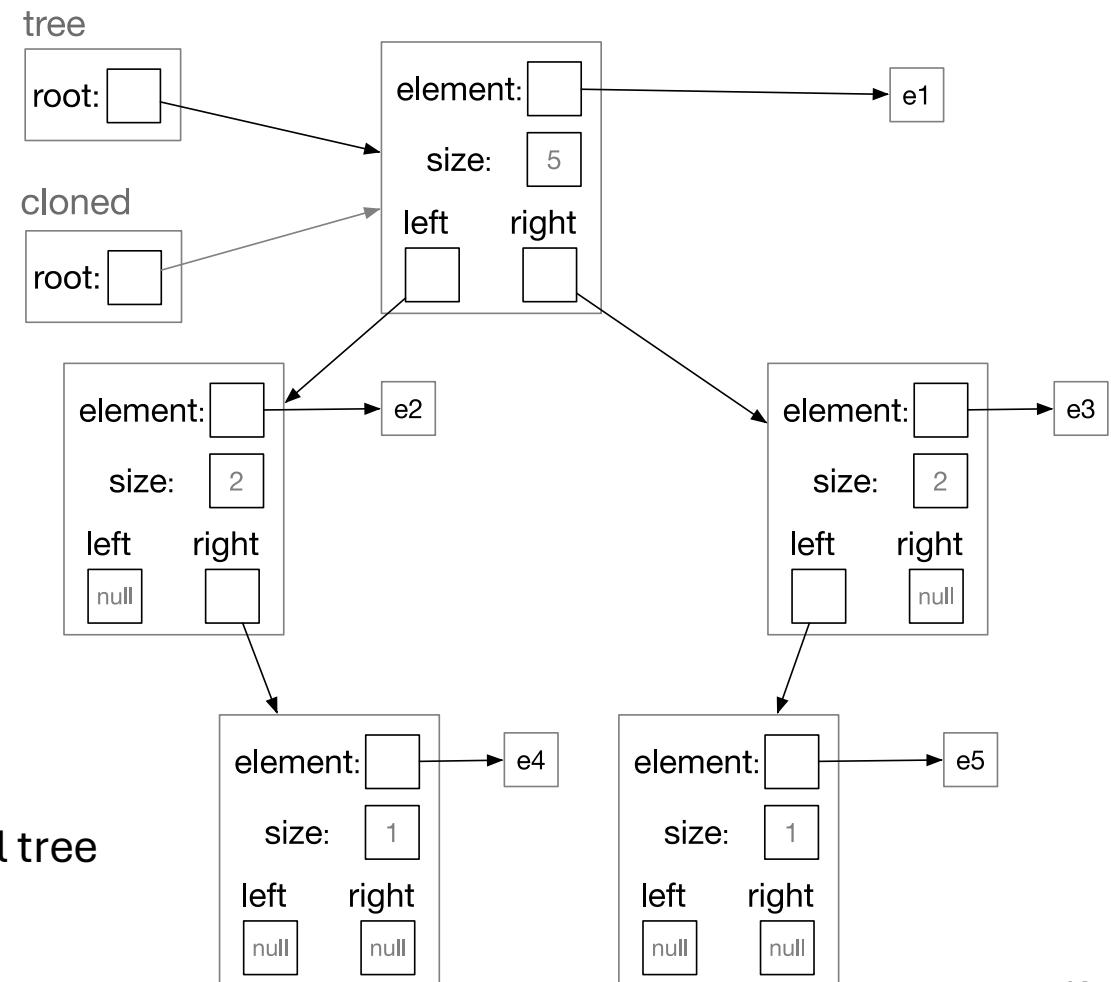
- What do we mean by field-for-field copy?
 - primitive types get its value copied
 - reference types get its reference aliased (that is why is called a shallow copy)
- When you implement the interface cloneable
 - you implement a method `clone()` that will return the copy
 - your implementation will call the one you inherit from your superclass (in our case, `Object`)
 - and then, we must decide whether we must do a deeper copy
- Let's see a minimal implementation and its problems

Binary Trees

```
public class LinkedBinaryTree<E>
    implements BinaryTree<E>, Cloneable {

    @Override
    @SuppressWarnings("unchecked")
    public LinkedBinaryTree<E> clone() {
        try {
            return (LinkedBinaryTree<E>) super.clone();
        } catch (CloneNotSupportedException e) {
            // this shouldn't happen, since we are Cloneable
            throw new InternalError(e);
        }
    }
}
```

- a new instance of LBT is created
- its root points to the root of the original tree
- what if we also clone the root?



```
public class LinkedBinaryTree<E> implements BinaryTree<E>, Cloneable {
```

```
private static class Node<E> implements Cloneable {
```

```
    @Override
```

```
    @SuppressWarnings("unchecked")
```

```
    public Node<E> clone() {
```

```
        try {
```

```
            return (Node<E>) super.clone();
```

```
        } catch (CloneNotSupportedException e) {
```

```
            // this shouldn't happen, since we are Cloneable
```

```
            throw new InternalError(e);
```

```
        }
```

```
}
```

```
    @Override
```

```
    @SuppressWarnings("unchecked")
```

```
    public LinkedBinaryTree<E> clone() {
```

```
        try {
```

```
            LinkedBinaryTree<E> clone = (LinkedBinaryTree<E>) super.clone();
```

```
            if (root != null) clone.root = root.clone();
```

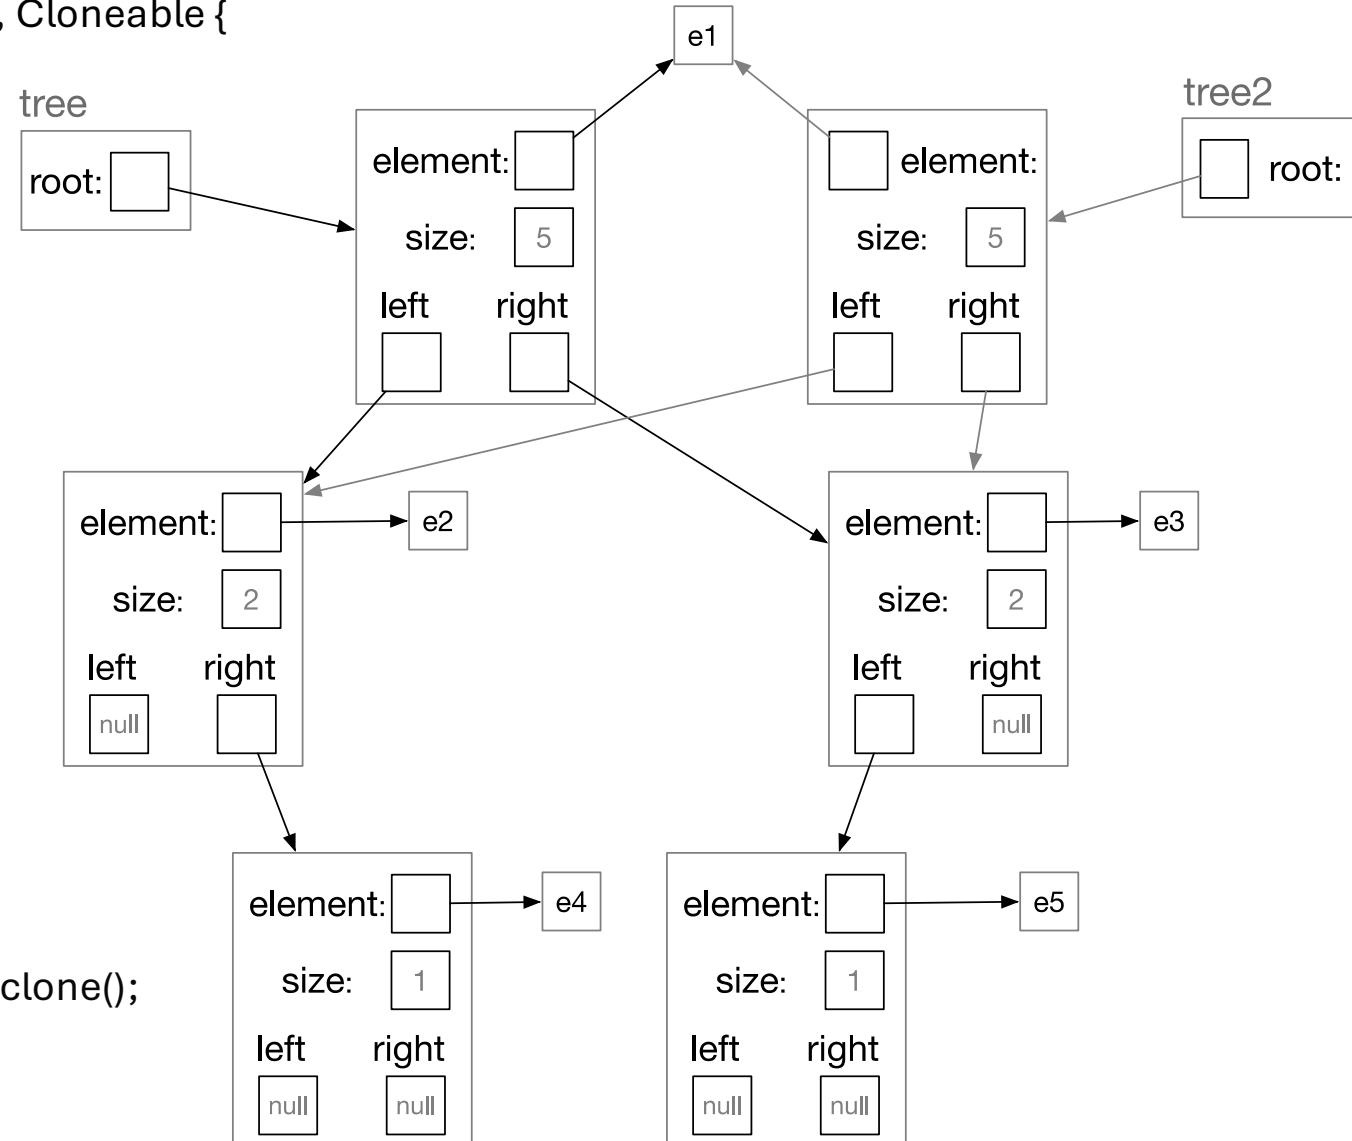
```
            return clone;
```

```
        } catch (CloneNotSupportedException e) {
```

```
            // this shouldn't happen, since we are Cloneable
```

```
            throw new InternalError(e);
```

```
}
```



- You continue sharing most of the tree
- You must clone down the tree !!

Binary Trees

- To do a **deep-copy** of the structure of the tree, the clone method of the class Node must do a **recursive cloning**
- But this does not solve the problem with sharing the instances of elements !!!
 - We can demand the E to be Cloneable as well
 - But then our implementation won't be usable for some types

```
private static class Node<E> implements Cloneable {  
    @Override  
    @SuppressWarnings("unchecked")  
    public Node<E> clone() {  
        try {  
            Node<E> clone = (Node<E>) super.clone();  
            if (left != null) clone.left = left.clone();  
            if (right != null) clone.right = right.clone();  
            return clone;  
        } catch (CloneNotSupportedException e) {  
            // this shouldn't happen, since we are Cloneable  
            throw new InternalError(e);  
        }  
    }  
}
```

Binary Trees

- You can access this implementation at this [repository](#)
- The implementation goes a little bit farther than the one described here because, as the empty tree is unmodifiable, it tries to create only one instance for it
- So, all references to the empty tree reference this instance
- As said before, the empty tree is unmodifiable (and immutable), so this is safe to do
 - And it saves some space !!!

Binary Trees

- But this implementation is not perfect because it allows the creation of **non-tree** things
 - But, at least, no cycles are allowed
- Can be avoid this?
 - Yes, but then the constructor won't be $\mathcal{O}(1)$ but $\mathcal{O}(n)$ because we'll need to clone the tree we receive as parameter
 - As always, if we can live with that, a $\mathcal{O}(1)$ cost is desirable, if not we must pay the price of cloning
- For example, draw the representation of the trees after:

```
var tree1 = new LinkedBinaryTree<Integer>(null, 1, null);
var tree2 = new LinkedBinaryTree<Integer>(tree1, 2, tree1);
```

Binary Search Trees (BSTs)

Binary Search Trees

- At the beginning, we've said that trees are usually used as an **implementation device** to implement other data structures
- This will be the case in hand, in which we'll use trees to implement an **associative data structure**
 - In this structure, we'll **associate keys to values**
 - **Keys are comparable**
- Possible implementations
 - List of pairs key-value
 - all operations linear
 - Sorted list (by key) of pairs key-value
 - search logarithmic
 - insertion/deletion linear

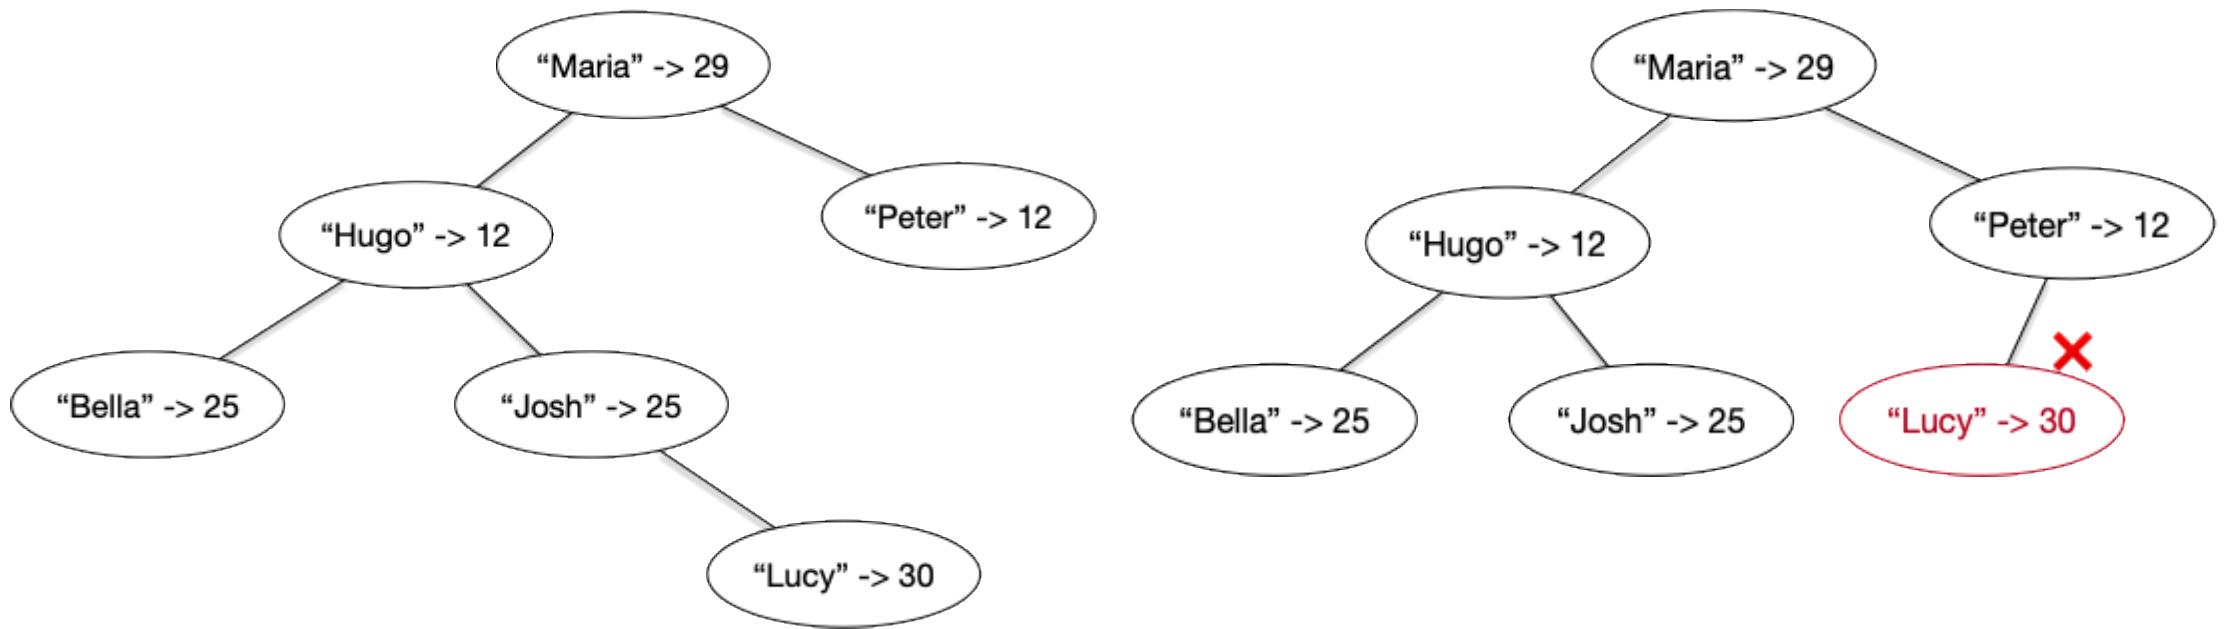
Binary Search Trees

- Binary Search Trees are binary trees which allow to implement these operations in **logarithmic time**
 - Well, **only if the tree is balanced**
- The **structure** of the tree is only **determined by its keys** (the values associated to them are only a payload)
 - So, we'll only show the keys in the diagrams
 - But each node carries a pair key-value
 - Well, when you use a BST to implement a Set, you do not have a value

Binary Search Trees

- A **binary search tree** is either
 - An **empty** binary tree
 - A non-empty binary tree in which
 - The **key** in the **root** is **greater** than all the **keys** in the **left** subtree
 - The **key** in the **root** is **lower** than all the **keys** in the **right** subtree
 - Both **left** and **right** subtrees are **binary search trees**
- So,
 - There are **no duplicated** keys
 - **Keys must be comparable**
 - If the BST implements an associative data structure, the nodes contains both a key and a value
 - Values can be duplicated in the tree and do not need to be comparable

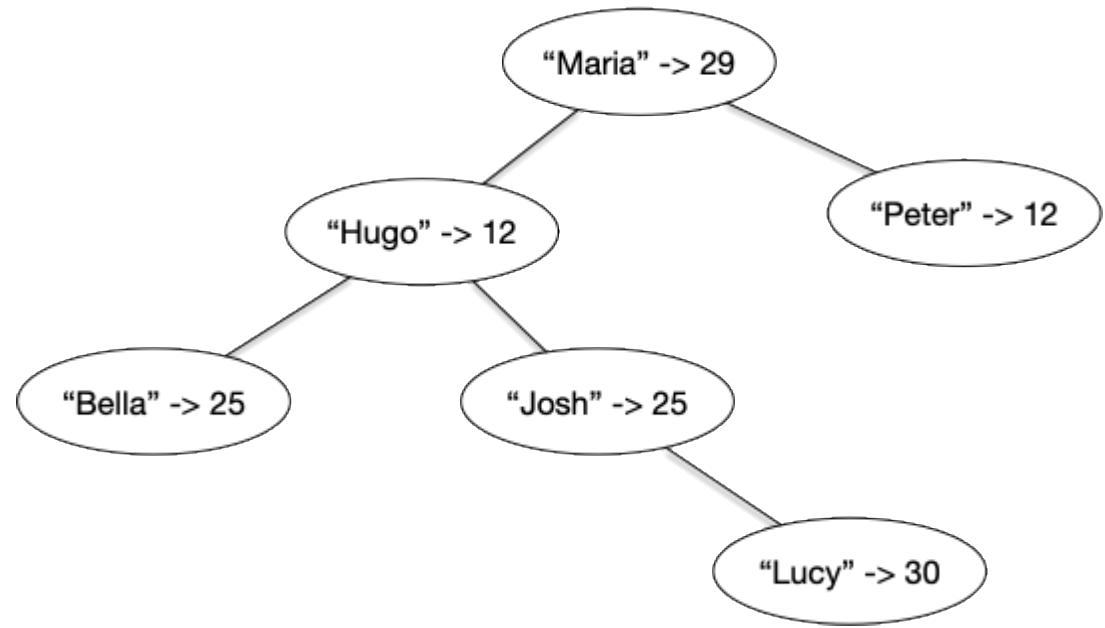
Binary Search Trees



Binary Search Trees

Search for the value associated to a key k

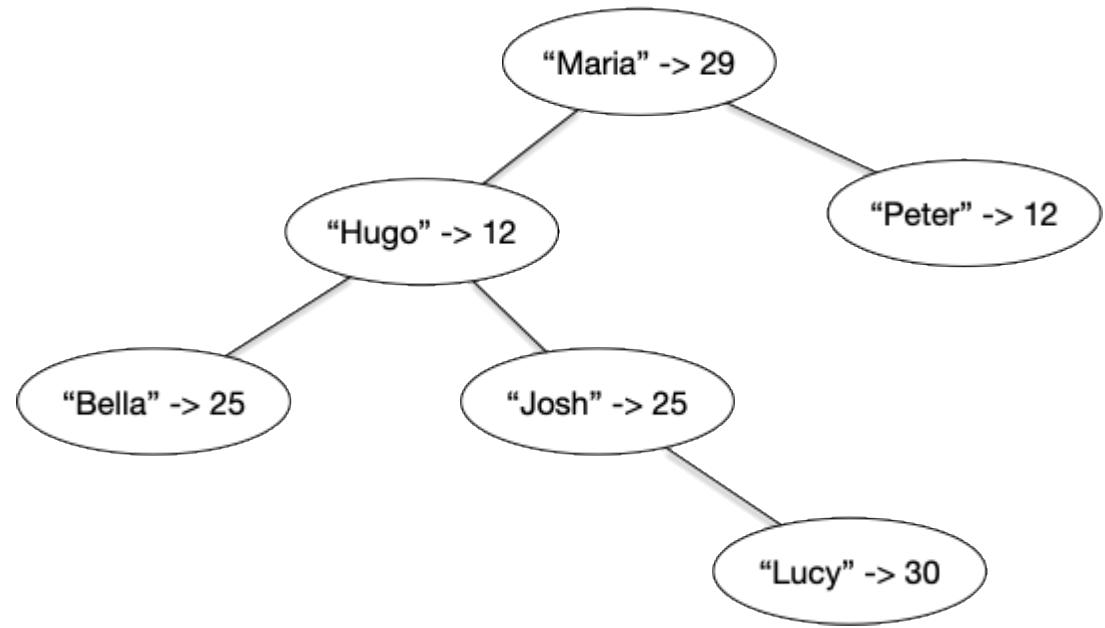
- If at, any point, the tree is empty, we know that the key k does not belong to the tree and we're finished.
- Beginning at the root node of the tree
 - If the key in the node is equal to k , we've found it, and we're finished
 - If the key in the node is greater than k , continue the search on the left child
 - If the key in the node is lower than k , continue the search on the right child.



Binary Search Trees

Insert (associate) the key k with the value v

- If the tree is empty, create a node to be the new root with the new pair
- If not, search for the key k and
 - if is found, change its associated value to v
 - If not, the node where the search failed was a leaf or had only one child
 - Add a new node with the new pair to it

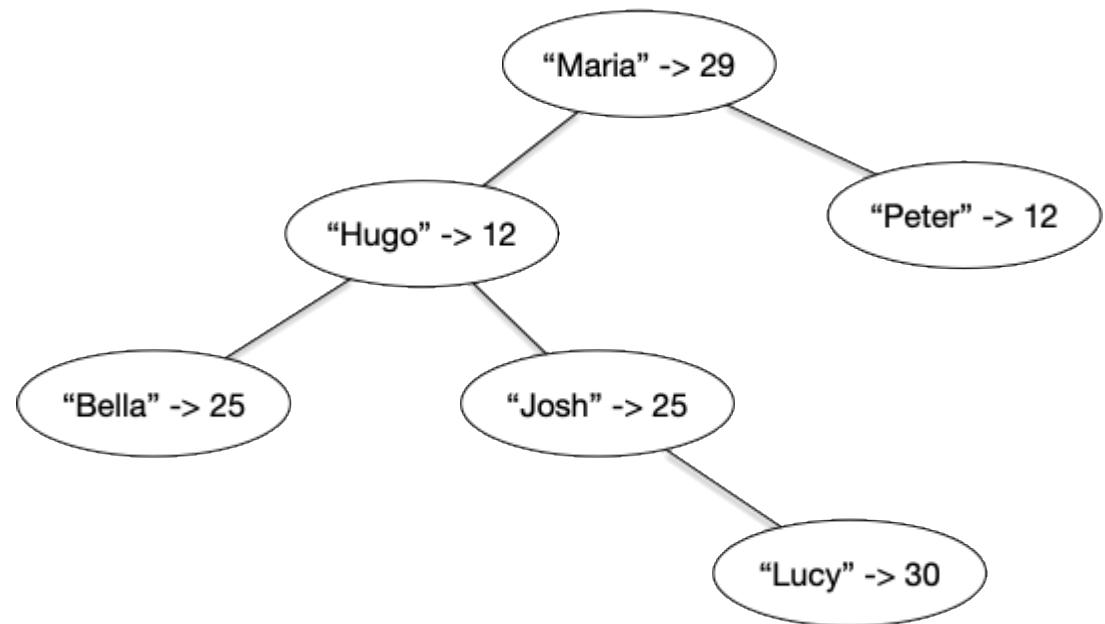


Binary Search Trees

Delete key k

There are four cases:

- A. k does not exist in the tree
- B. It's in a leaf
- C. It's in a node of degree 1 (with a single child)
- D. It's in a node of degree 2 (with two children)

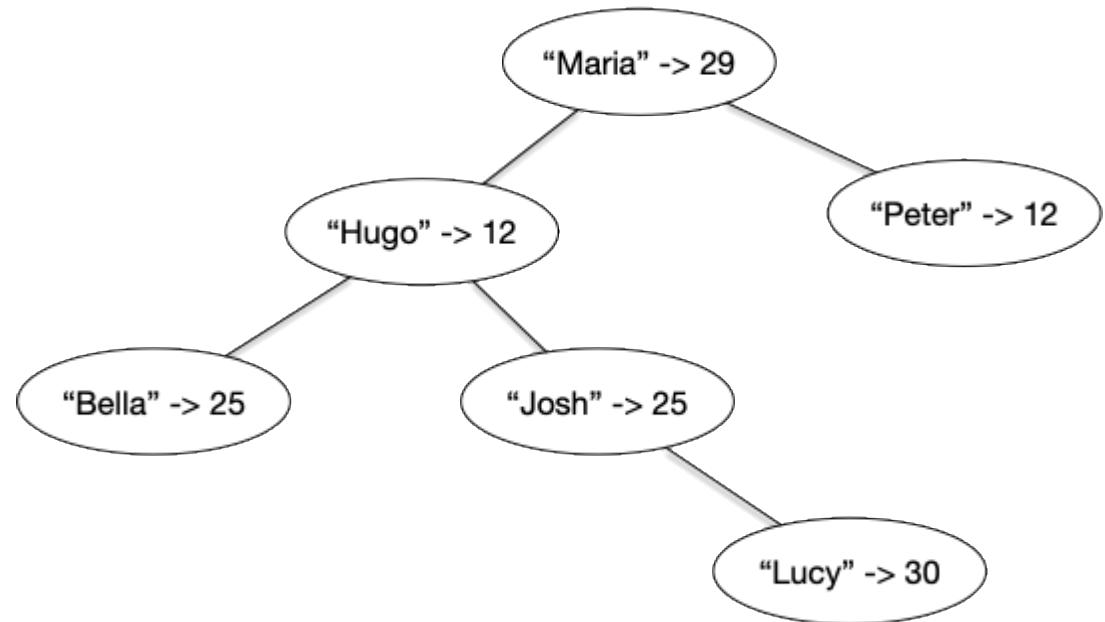


Binary Search Trees

Delete key k

There are four cases:

- A. k does not exist in the tree
 - No further action is needed

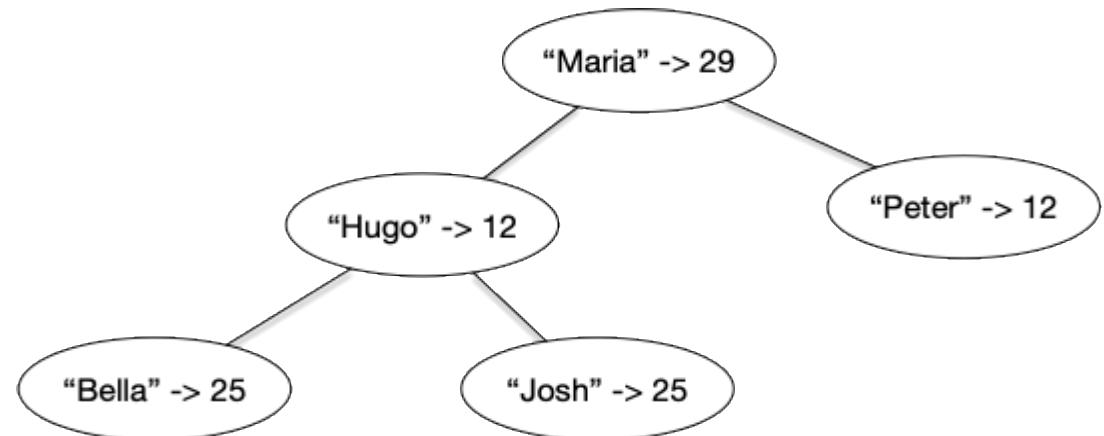
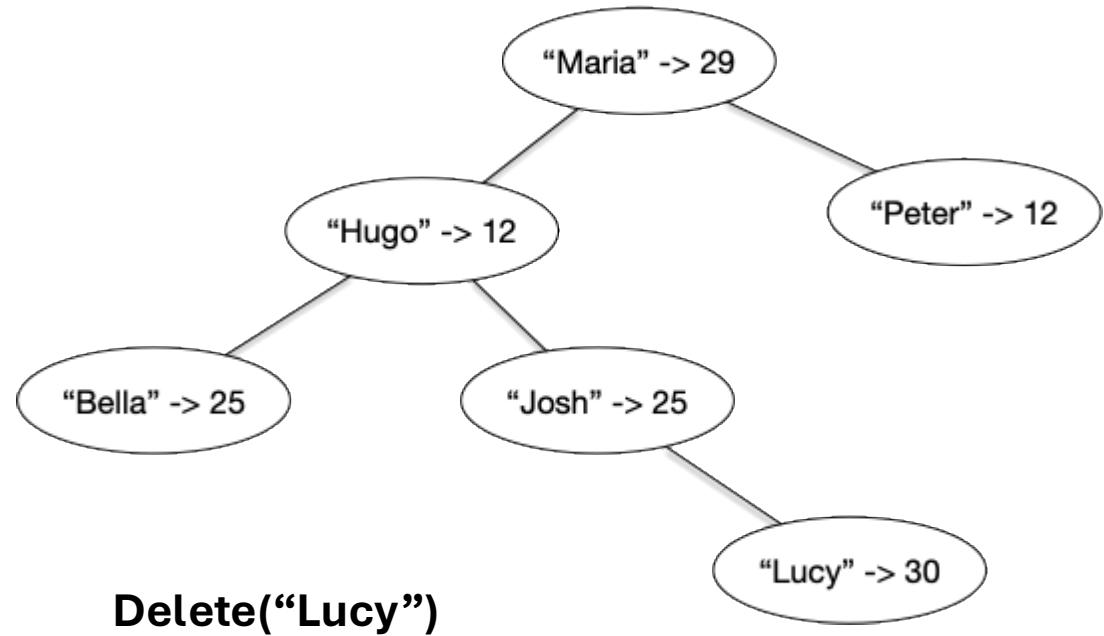


Binary Search Trees

Delete key k

There are four cases:

- B. It's in a leaf
 - We delete the leaf

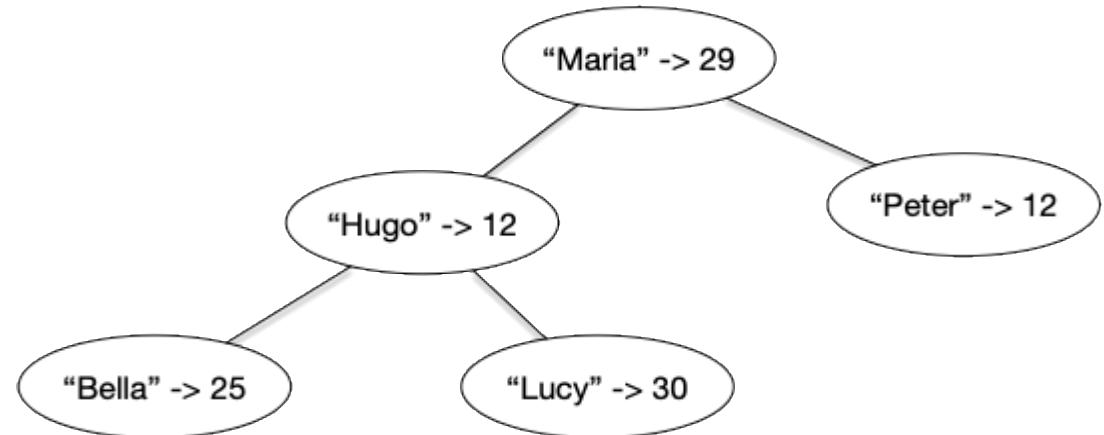
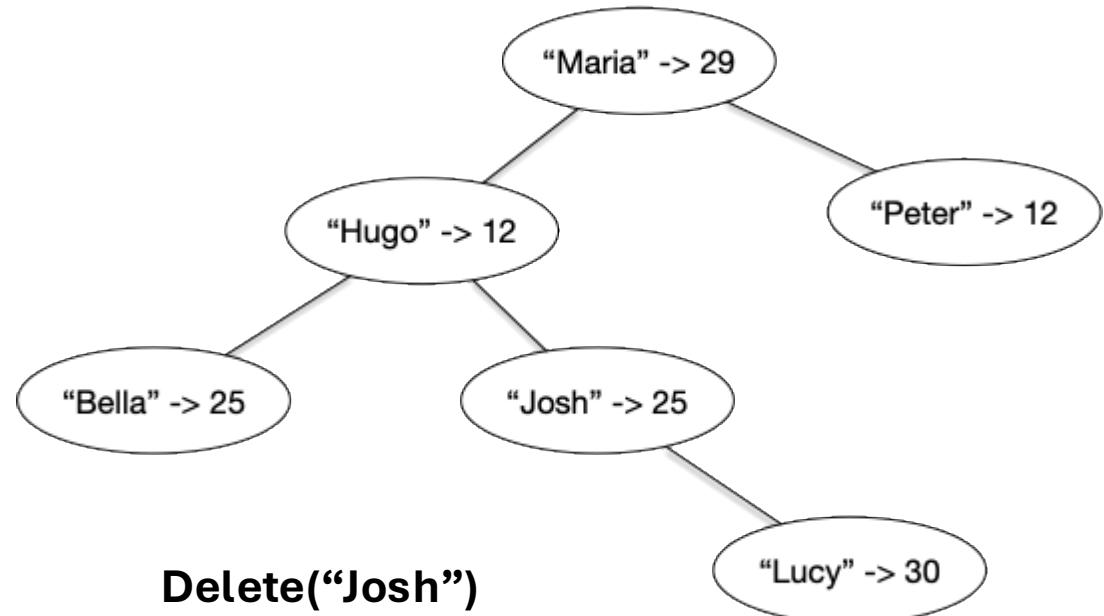


Binary Search Trees

Delete key k

There are four cases:

- C. It's in a node of degree 1 (with a single child)
 - Its single child substitutes it

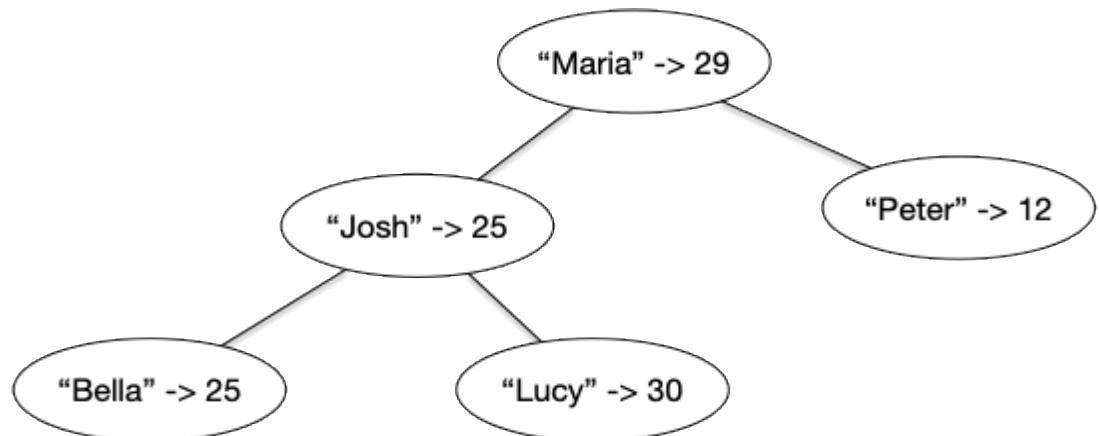
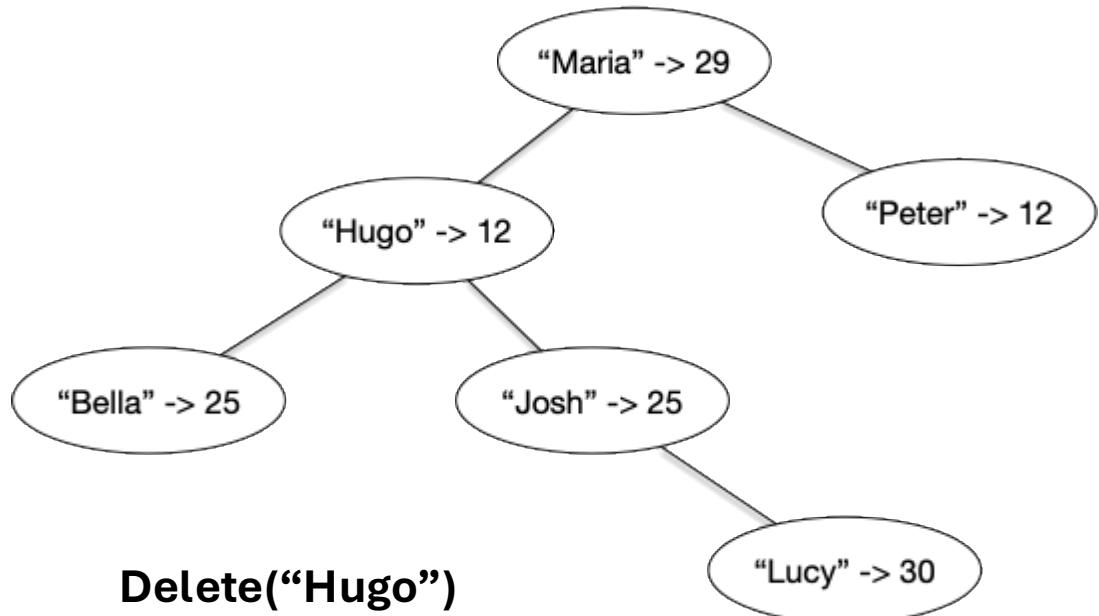


Binary Search Trees

Delete key k

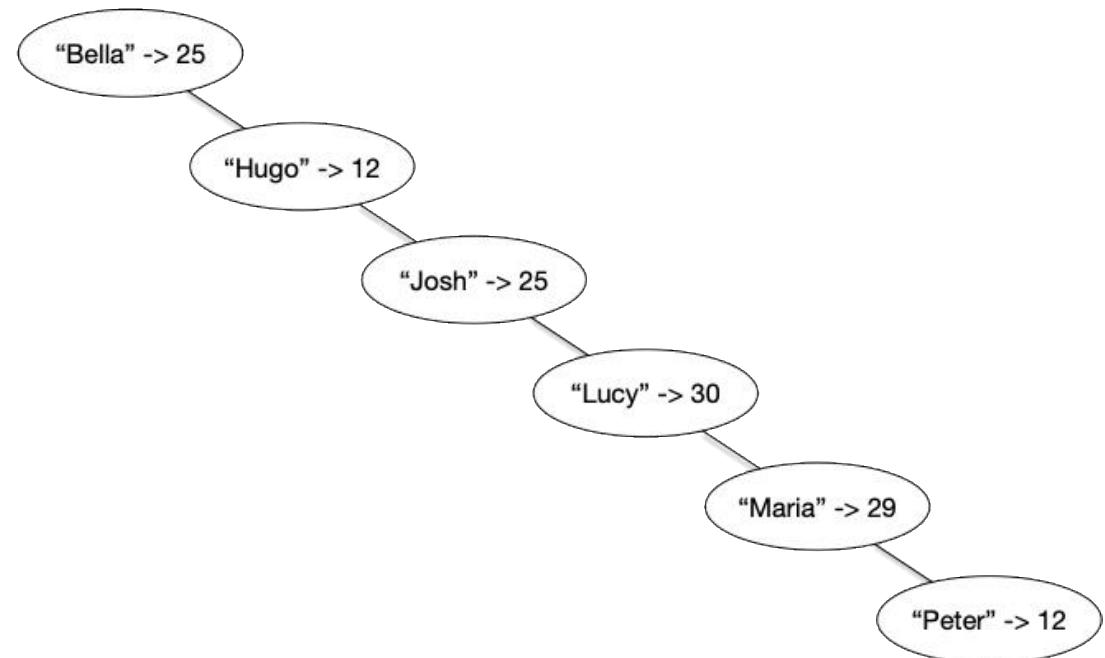
There are four cases:

- D. It's in a node of degree 2 (with two children)
 - We substitute the key pair with that of the lowest key in the right subtree
 - And we remove it from the right subtree
- NOTES:
 - The key is the next to k in sorted order
 - We could have used the biggest key in the left subtree (the previous in sorted order)



Binary Search Trees

- All the operations described above have cost that depends on the **height of the tree**
 - But this make the **linear on the size of the tree**
- Why?
 - Because nothing makes the tree to be balanced !!



Binary Search Trees

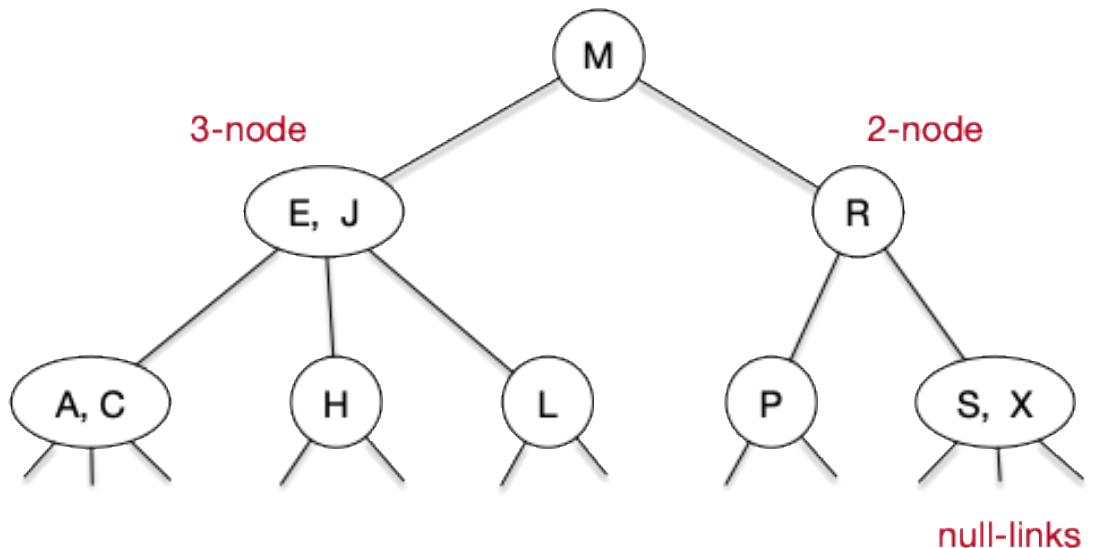
- So, to make the operations in a tree **logarithmic** on its size, one must guarantee that the tree is **balanced** !!
- AVL Trees
 - Invented by Georgii Adelson-Velsky and Yevgeniy Landis in 1962
 - The tree is perfectly balanced all time
 - But operations are complex to implement
- Red-Black Trees
 - Invented by Leonidas Guibas and Robert Sedgewick in 1978, based on previous work by Rudolf Bayer in 1972
 - The tree is not perfectly balanced but guaranteed logarithmic time
 - Operations are simpler to implement

2-3 Trees

- Although the most used implementation is that of Cormen et al., we'll present Sedgewick's because it is simpler
- And we'll only consider the search (trivial) and insertion operations
 - For the rest of the operations, refer to section 3.3 of Sedgewick's book
 - Or its associated [web page](#)
 - Or this paper [Left-leaning Red-Black Trees](#)
- But this presentation does not start with Red-Black trees but with another kind of (not binary) search trees: the **2-3-Search Trees**

2-3 Trees

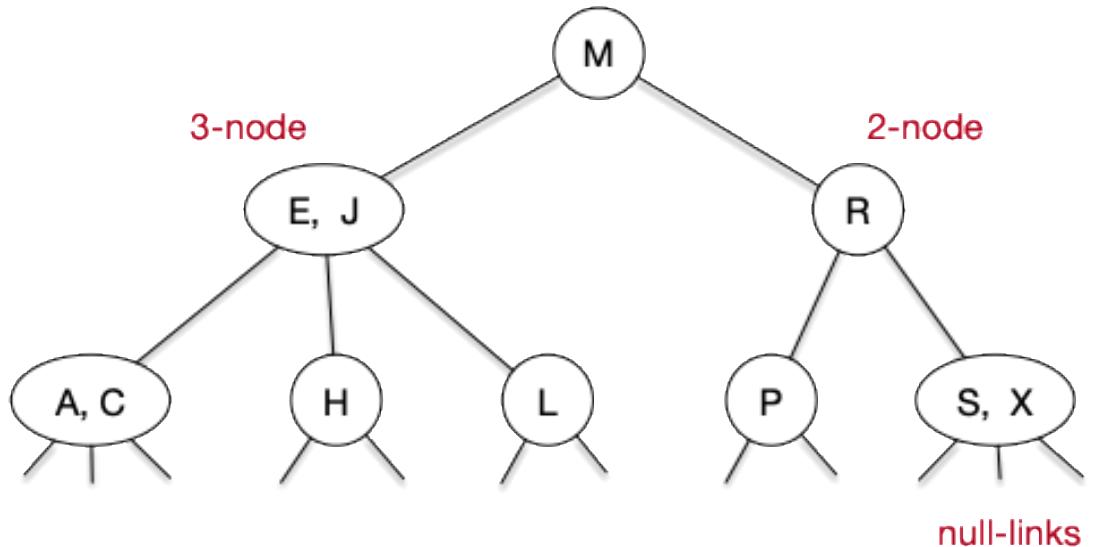
- A **2-3 Search Tree** is a tree that can be
 - **Empty**
 - A **2-Node** with a key (and associated value) and a left 23-subtree with smaller keys and a right 23-subtree with bigger keys
 - A **3-Node** with two keys (and associated values) and a left 23-subtree with smaller keys than those in the node, a middle 23-subtree with keys between those in the node, and a right 23-subtree with bigger keys than those in the node
- A **null-link** is a link to an empty tree.
- A **perfectly balanced 2-3-Search Tree** has all the **null-links** at the same distance from the root
 - We'll only consider those and simply call them **2-3-Trees**



2-3 Trees

Search for a key k

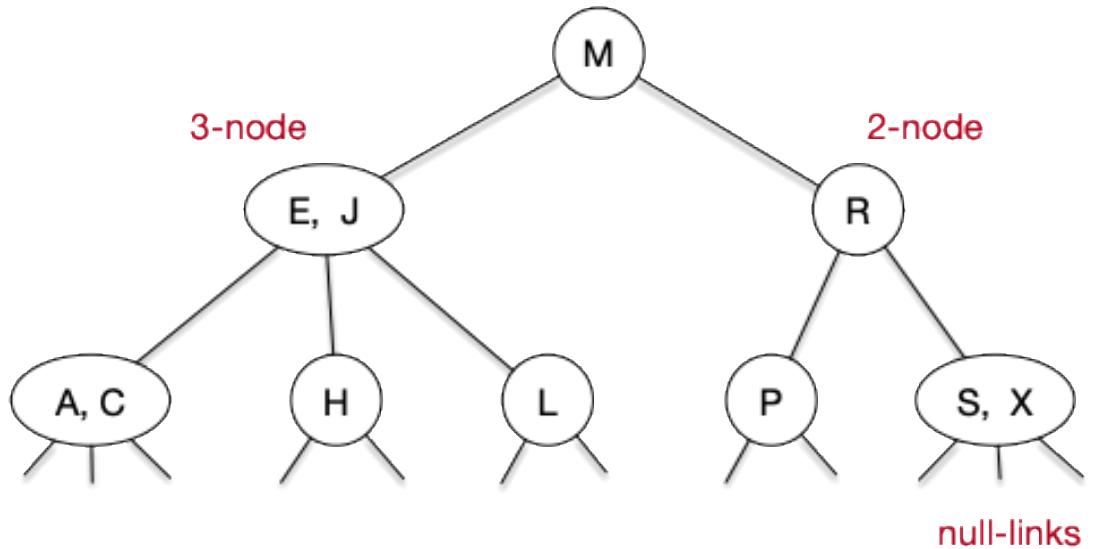
- A simple extension of the algorithm for BSTs
- As we only consider balanced 23-Trees, the cost of the search is logarithmic on the number of nodes
 - The height of the tree is between $\log_3 n$ and $\log_2 n$



2-3 Trees

Insertion of a key k (and associated value v):

- We proceed as in a BST
- We have several cases:
 - A. Insert into a 2-Node
 - B. Insert into a 3-Node
 - i. It's the single node of the tree
 - ii. His parent is a 2-Node
 - iii. His parent is a 3-Node
 - iv. It's the root

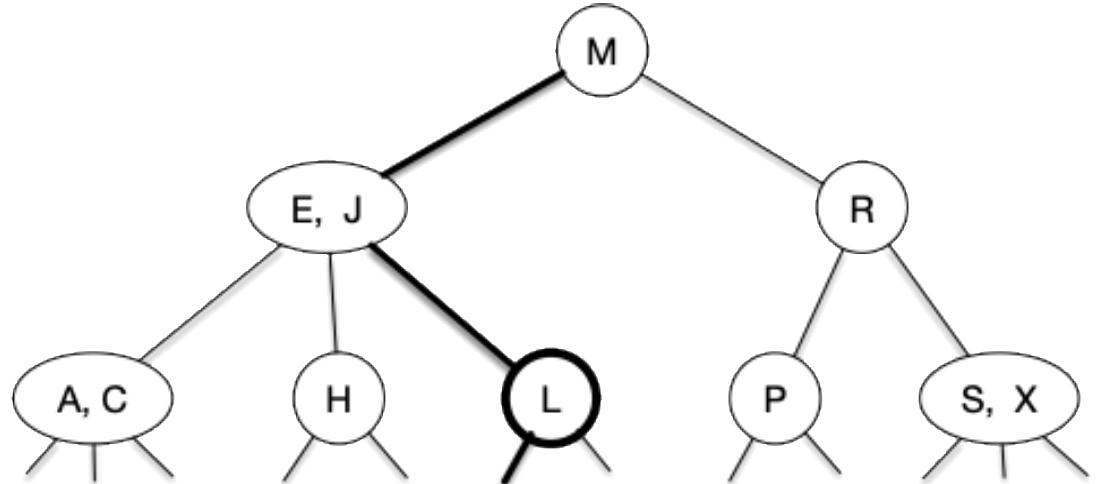


2-3 Trees

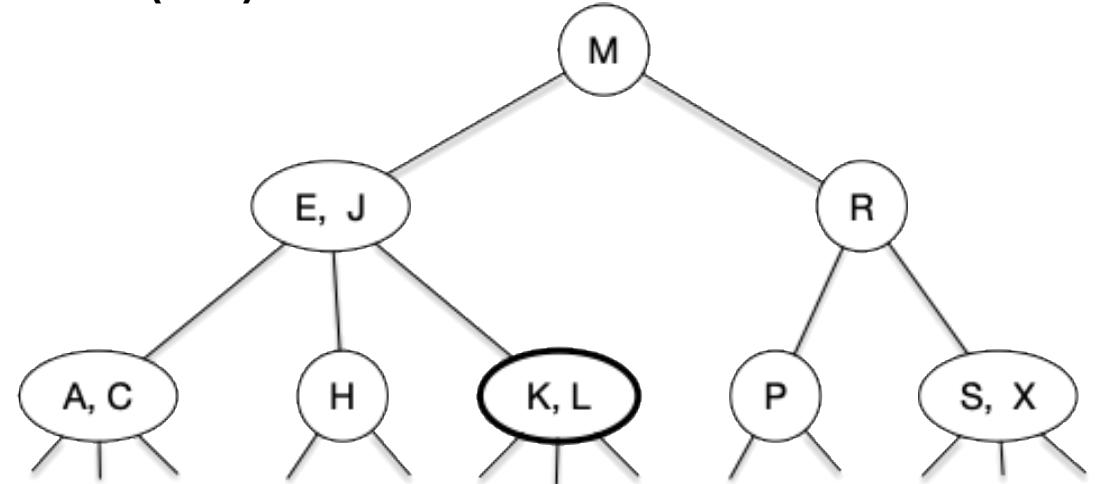
Insertion of a key k (and associated value v):

A. Insert into a 2-Node

- **Replace** the 2-Node where we fail with a 3-Node



Insert("K")



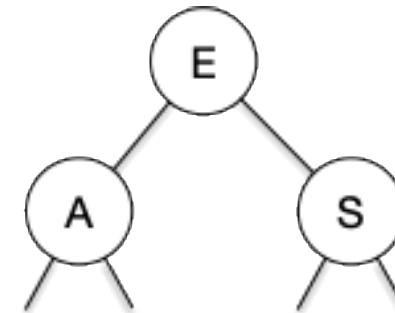
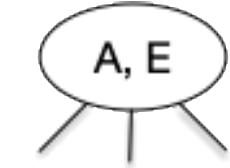
2-3 Trees

Insertion of a key k (and associated value v):

B. Insert into a 3-Node

- i. It's the single node of the tree
 - **Replace** the 3-Node with a **temporal** 4-Node
 - **Split** it into three 2-Nodes

Insert("S")



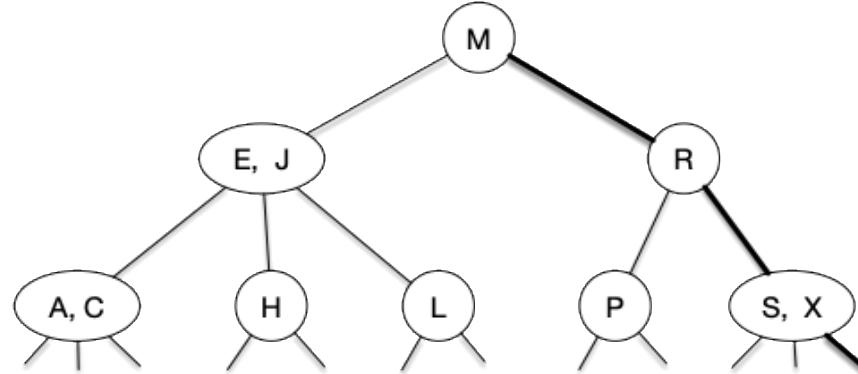
2-3 Trees

Insertion of a key k (and associated value v):

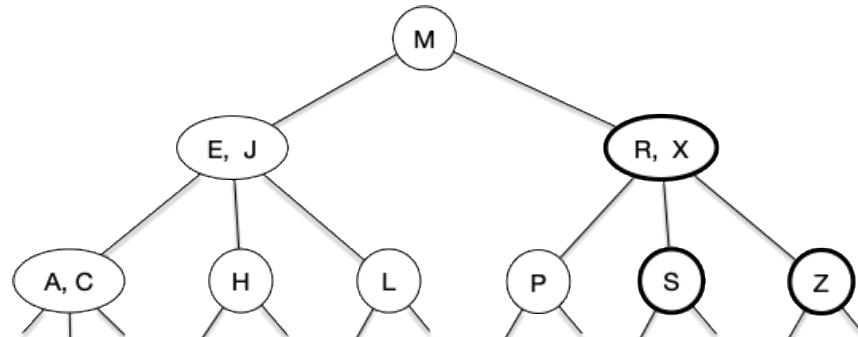
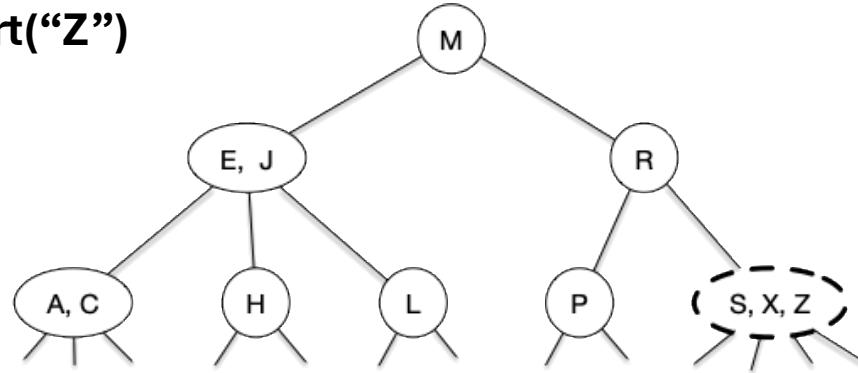
B. Insert into a 3-Node

ii. Its parent is a 2-Node

- **Replace** the 3-Node with a temporal 4-Node
- **Split** the 4-node into two 2-Nodes (for the left and right)
- **Insert** the middle key on the parent



Insert("Z")



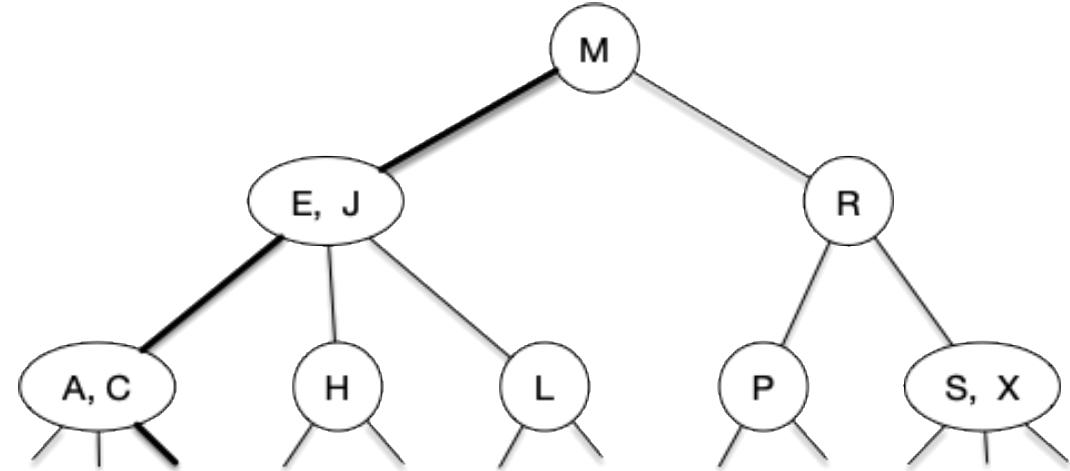
2-3 Trees

Insertion of a key k (and associated value v):

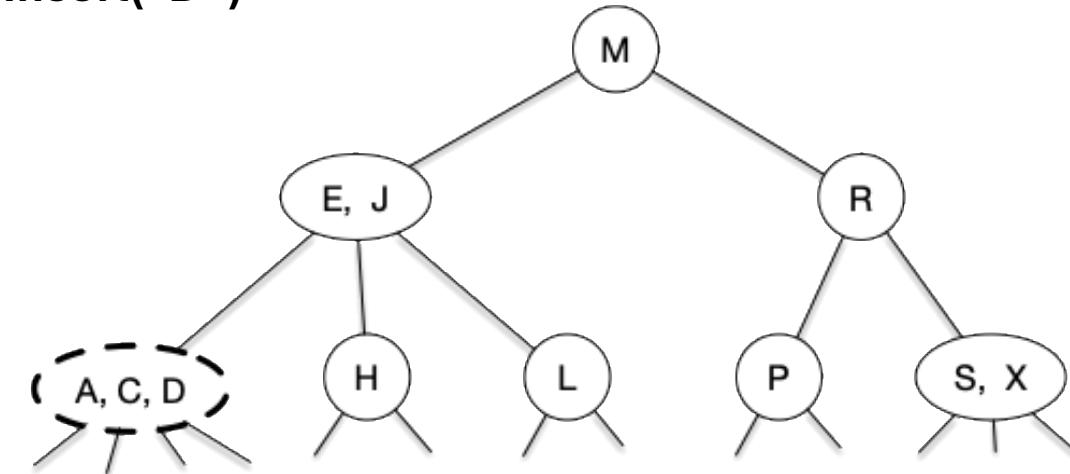
B. Insert into a 3-Node

iii. Its parent is a 3-Node

- **Replace** the 3-Node with a temporal 4-Node
- **Split** the 4-node into two 2-Nodes (for the left and right)
- **Insert** the middle key on the parent creating a new temporal 4-Node and ... (this can arrive up to the root)



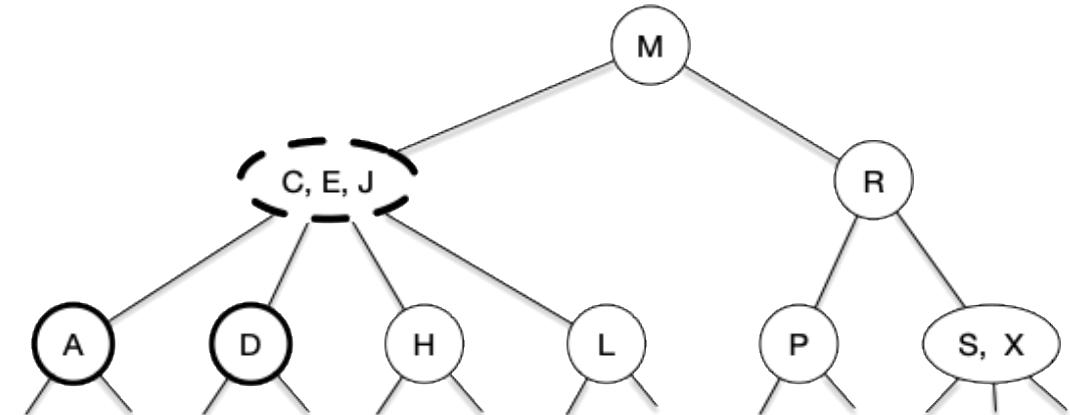
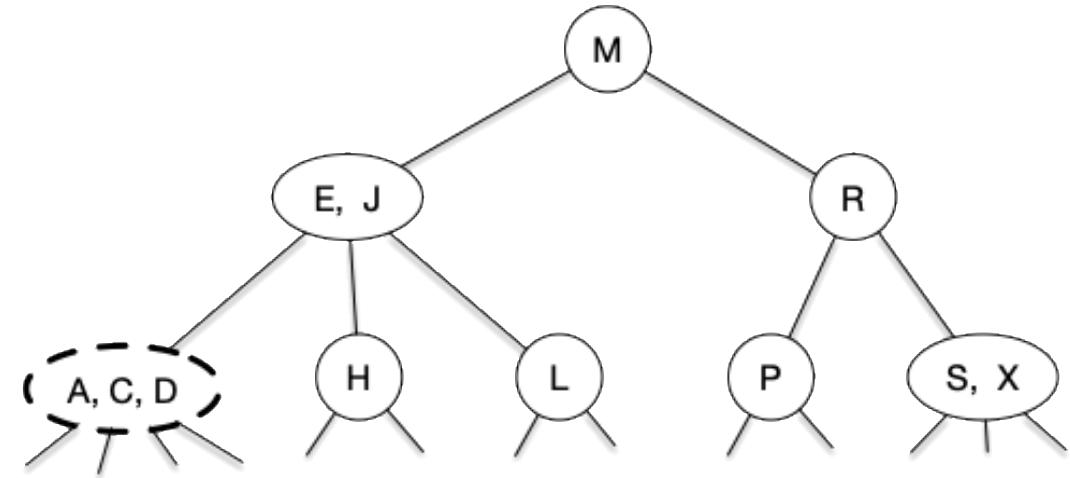
Insert("D")



2-3 Trees

Insertion of a key k (and associated value v):

- B. Insert into a 3-Node
 - iii. Its parent is a 3-Node
 - **Replace** the 3-Node with a temporal 4-Node
 - **Split** the 4-node into two 2-Nodes (for the left and right)
 - **Insert** the middle key on the parent creating a new temporal 4-Node and ... (this can arrive up to the root)



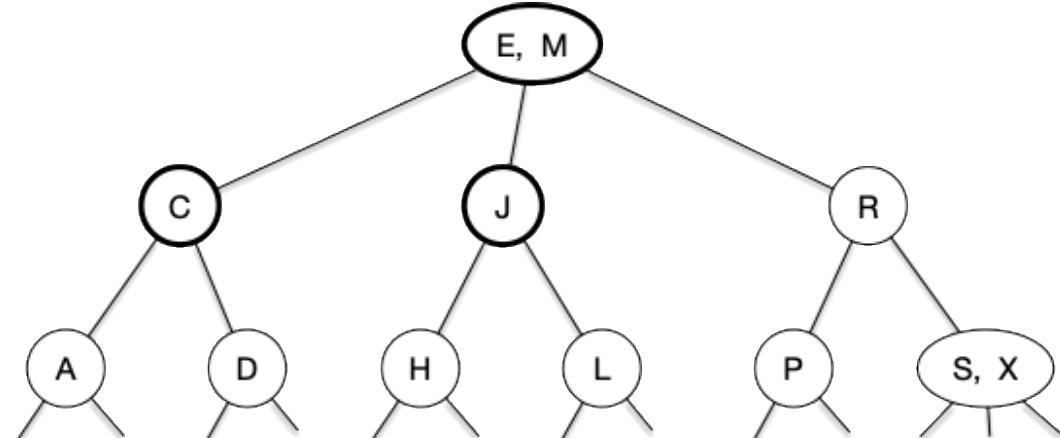
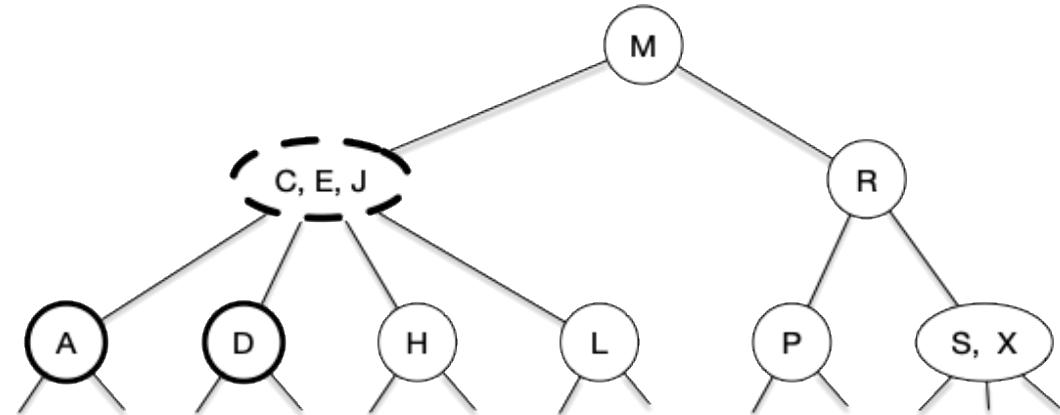
2-3 Trees

Insertion of a key k (and associated value v):

B. Insert into a 3-Node

iii. Its parent is a 3-Node

- **Replace** the 3-Node with a temporal 4-Node
- **Split** the 4-node into two 2-Nodes (for the left and right)
- **Insert** the middle key on the parent creating a new temporal 4-Node and ... (this can arrive up to the root)



2-3 Trees

Insertion of a key k (and associated value v):

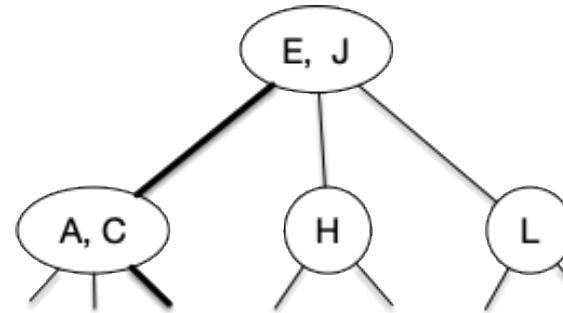
B. Insert into a 3-Node

iv. It's the root

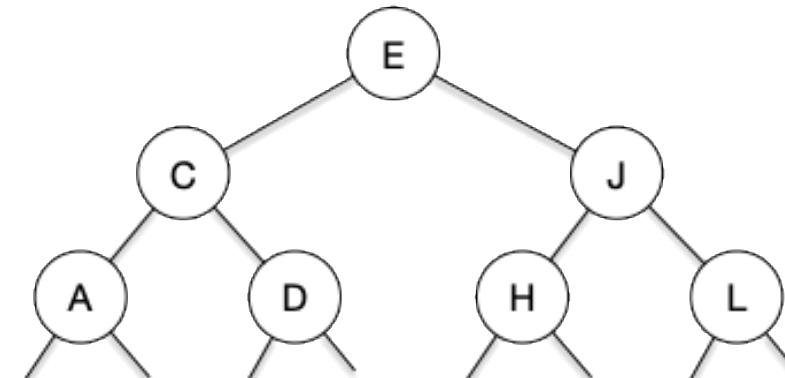
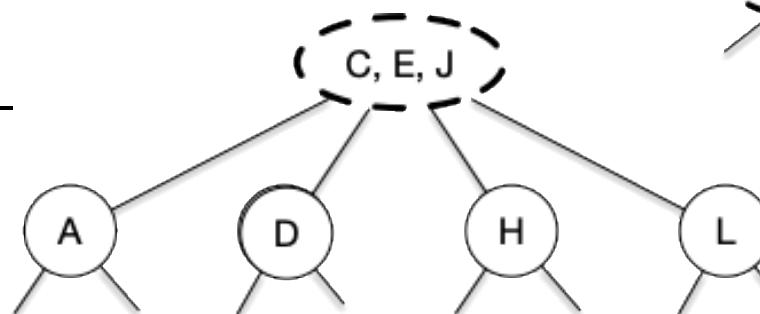
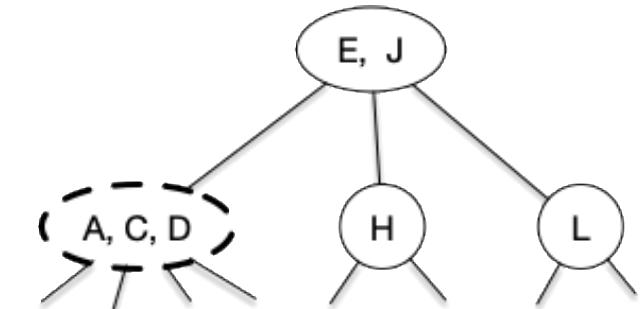
- **Replace** the root with a temporal 4-Node
- **Split** it into three 2-Nodes

NOTES:

- This is when there is a path formed by only 3-Nodes from the point of insertion to the root
- Actually, it's the same case as B.i



Insert("D")

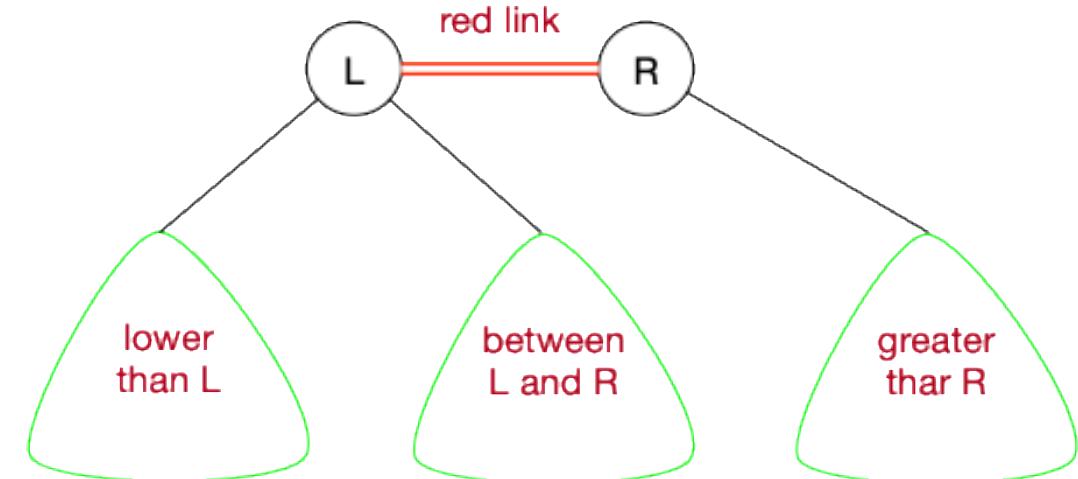
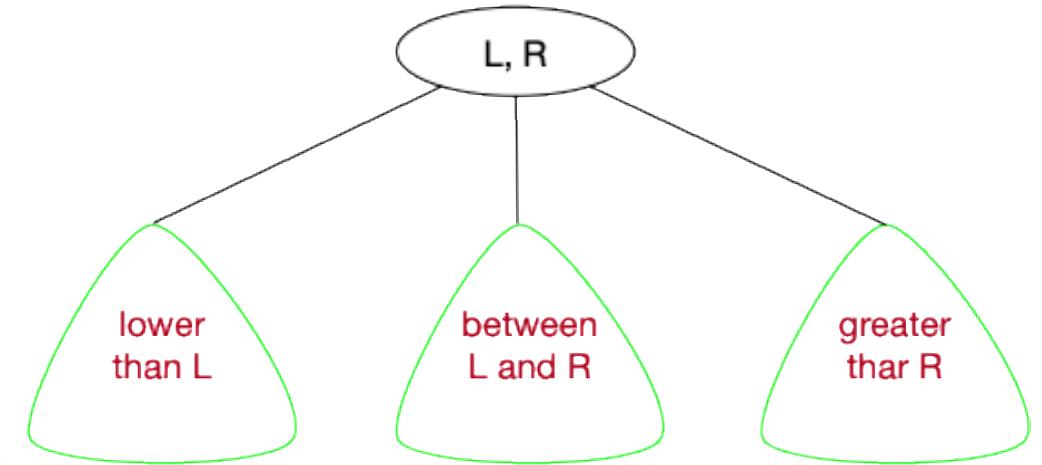


2-3 Trees

- **Splitting** a temporary 4-Node is a **local operation** of the tree ($\mathcal{O}(1)$)
 - No other part of the tree must be examined other than the specific nodes and links
 - And, at most, it propagates up to the root (**logarithmic** path)
- Besides, all transformations **preserve** the **order** and **perfect balance** of the tree
- **NOTE:**
 - Unlike BSTs, 23-Trees grow from the bottom to the top
- **Problem:** implementing it with different kinds of nodes is not easy.
- **Solution:** red-black trees are a way to **encode 2-3-trees** with only **one kind of node**.

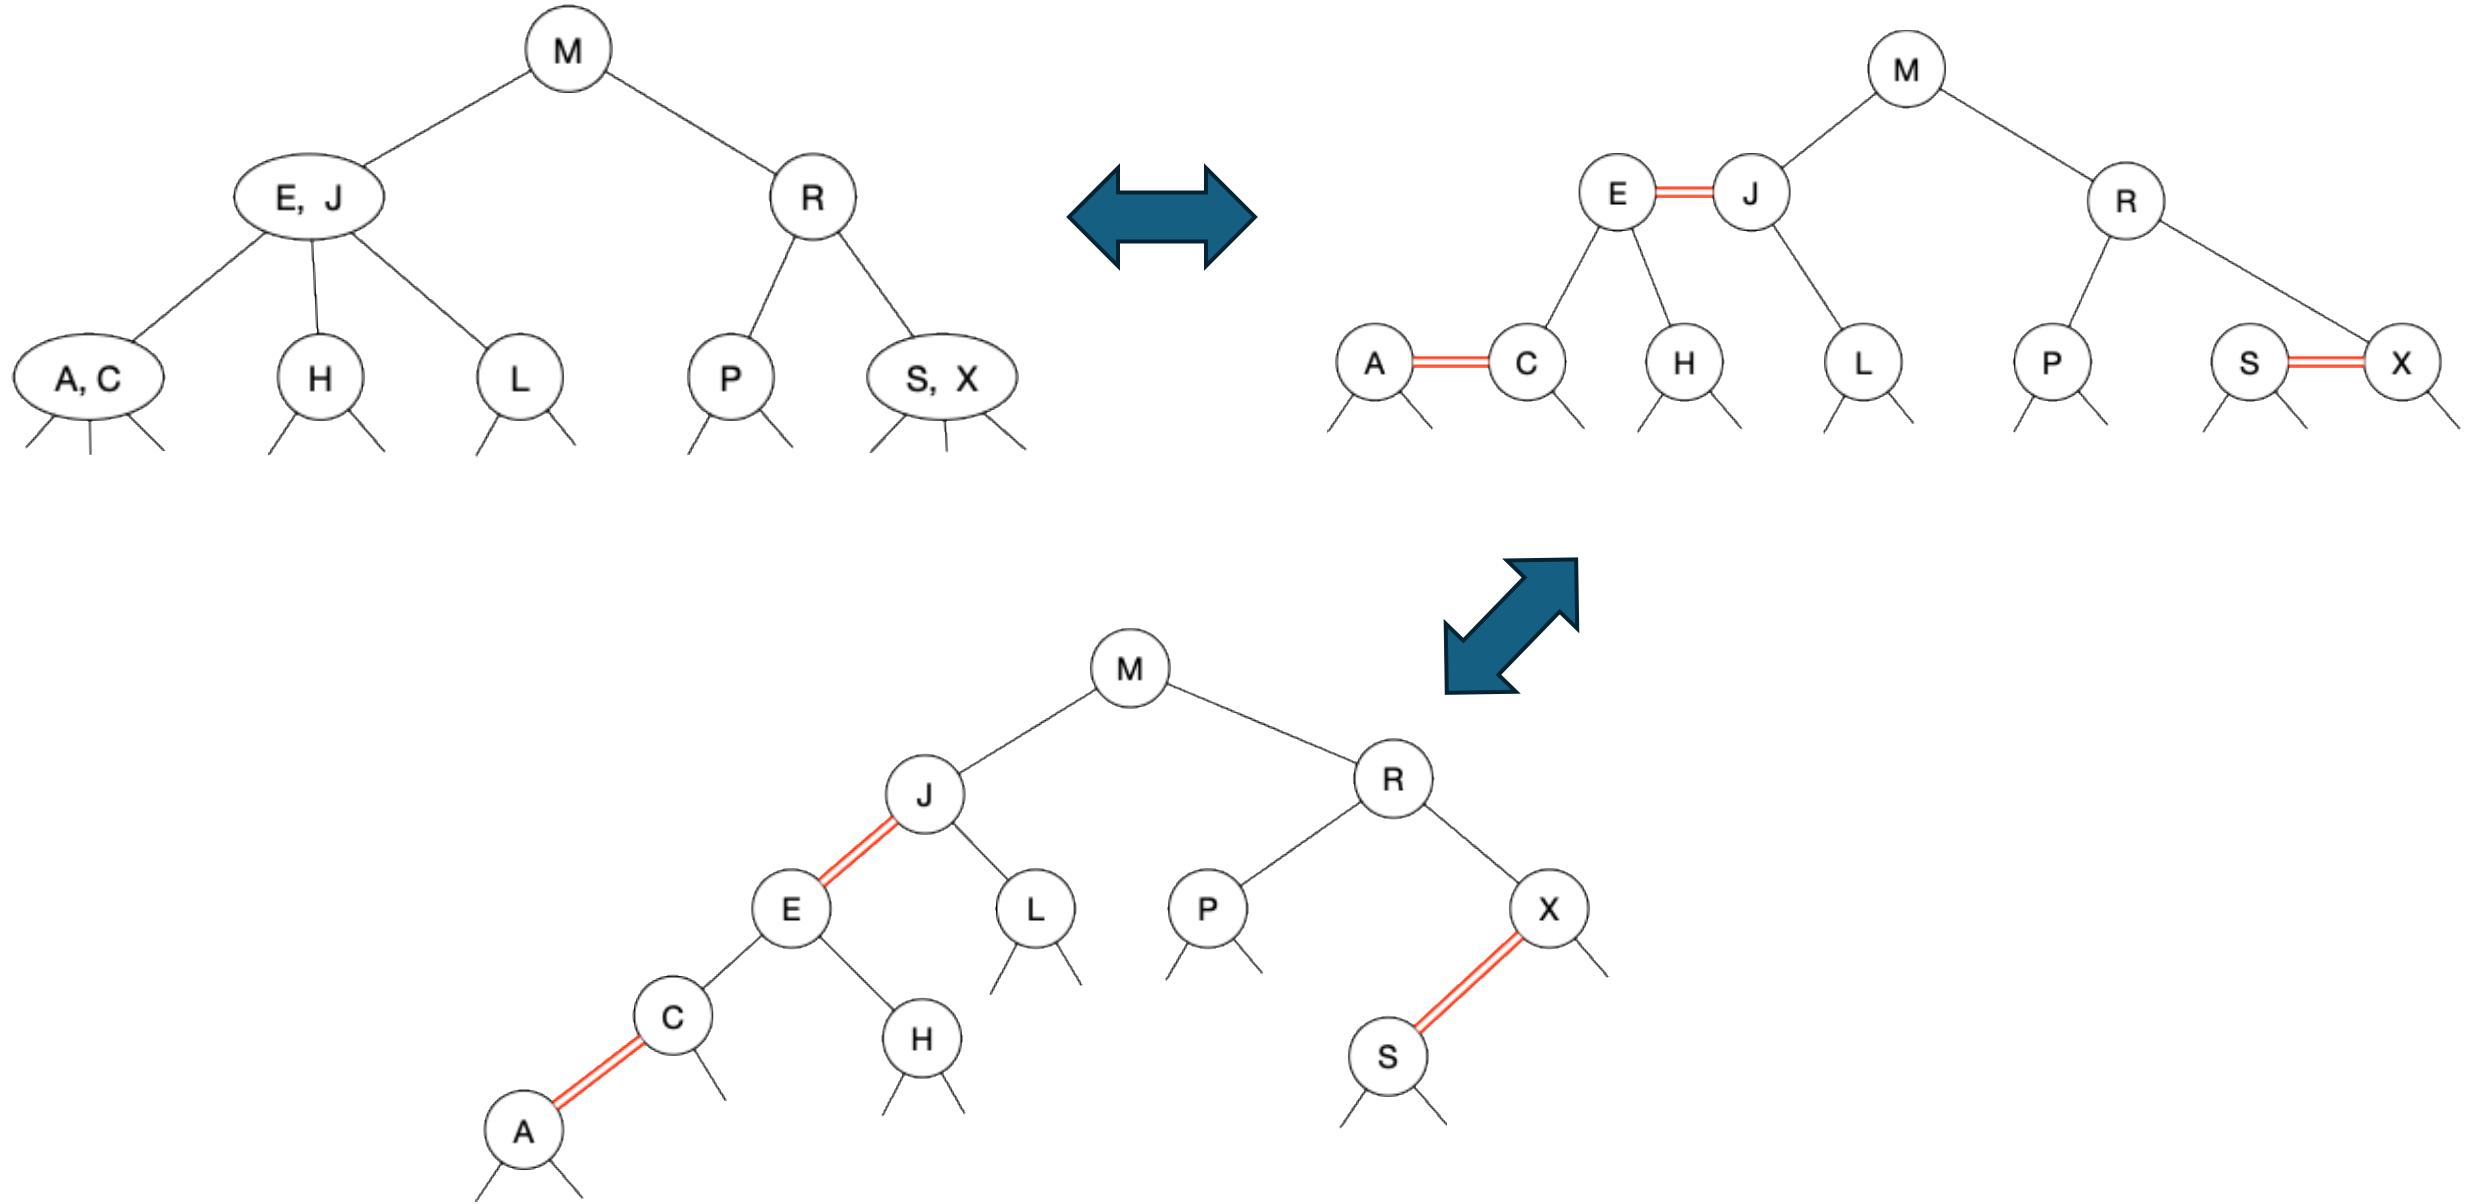
Red-Black Trees

- **NOTE:** There are slightly different versions of red-black trees, and we'll follow Sedgewick's
- The basic idea behind a red-black tree is to **encode 2-3 trees**
 - Starting with simple BSTs to encode 2-nodes
 - Adding additional information to encode 3-nodes
- Two different kind of links:
 - **Red links:** bind two 2-nodes to represent a 3-node
 - **Black links:** bind the 2-3 tree together
 - **Null-links** are always **black**



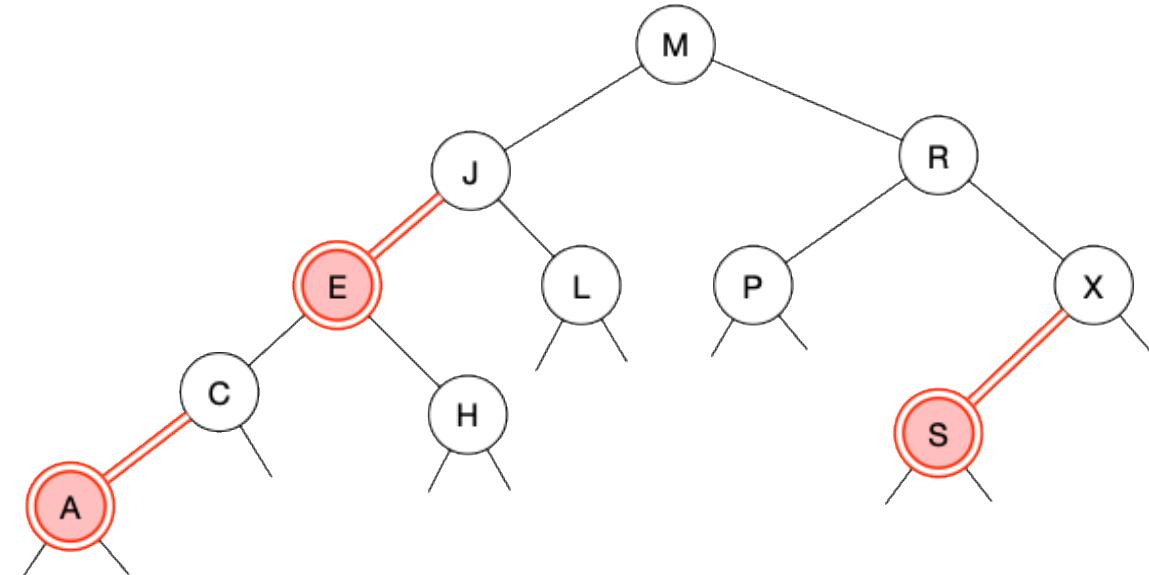
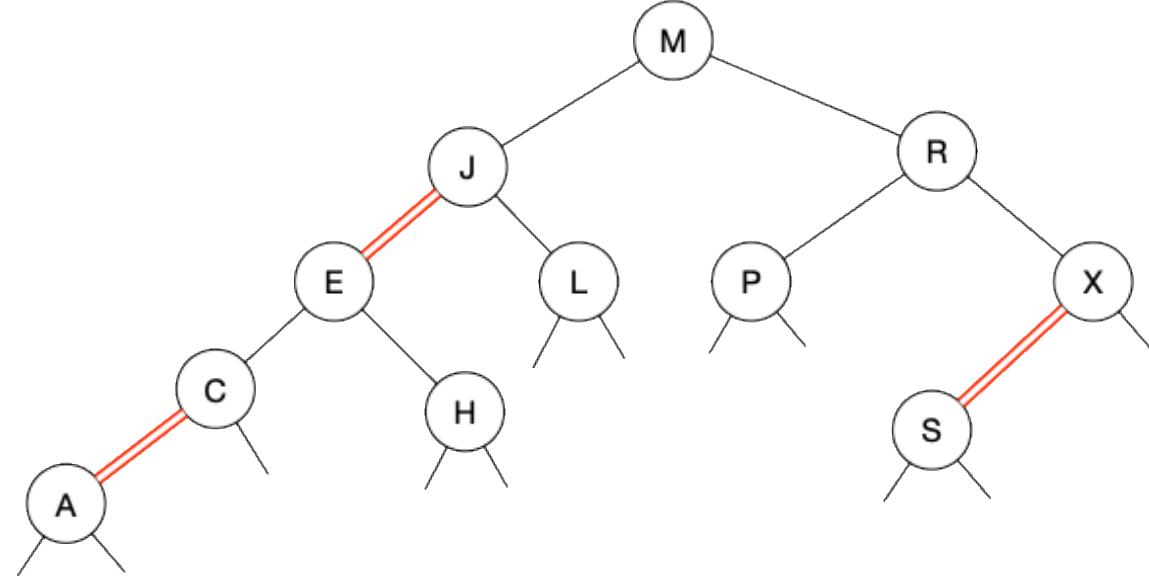
Red-Black Trees

- We can define this representation as BSTs having red and black links and satisfying the following restrictions:
 1. **Red links lean left**
 2. **No node has two red links** connected to it
 3. The tree has **perfect black balance**: every **path** to the **root** to a **null-link** has the **same number of blank links**
- This type of RB Trees are called **Left Leaning Red Black Trees (LLRBs)**
- There is a 1-to-1 relation (**isomorphism**) between LLRBs and 2-3 trees



Red-Black Trees

- As links are references, they cannot encode its colour
- But, as each node has only one link that comes from its parent, **we'll encode its colour in the child.**
- So, we'll have two kinds of nodes:
 - **Black nodes**: those in which their parent link is black
 - **Red nodes**: those in which their parent link is red
- The **root node will always be black**.



Red-Black Trees

- The full code of the implementation in Sedgewick's book can be found at [code](#)
- Sometimes the use of Java constructs is simplified:
 - Comparable
 - Non-static class Node
- **NOTE:** We have preferred to maintain the original implementation.

```
public class RedBlackBST<Key extends Comparable<Key>, Value> {  
  
    private static final boolean RED = true;  
    private static final boolean BLACK = false;  
  
    private Node root; // root of the BST  
  
    // BST helper node data type  
    private class Node {  
        private Key key; // key  
        private Value val; // associated data  
        private Node left, right; // links to left and right subtrees  
        private boolean color; // color of parent link  
        private int size; // subtree count  
  
        public Node(Key key, Value val, boolean color, int size) {...}  
    }  
  
    private boolean isRed(Node x) {  
        if (x == null) return false;  
        return x.color == RED;  
    }  
    ...  
}
```

Red-Black Trees

- The **search** algorithm is the **same** as in **regular BST**
- That is, a top-down search following the ordering imposed by the keys

```
public Value get(Key key) {  
    if (key == null)  
        throw new IllegalArgumentException("key is null");  
    return get(root, key);  
}  
  
// value associated with the given key in subtree rooted at x;  
// null if no such key  
private Value get(Node x, Key key) {  
    while (x != null) {  
        int cmp = key.compareTo(x.key);  
        if      (cmp < 0) x = x.left;  
        else if (cmp > 0) x = x.right;  
        else return x.val;  
    }  
    return null;  
}
```

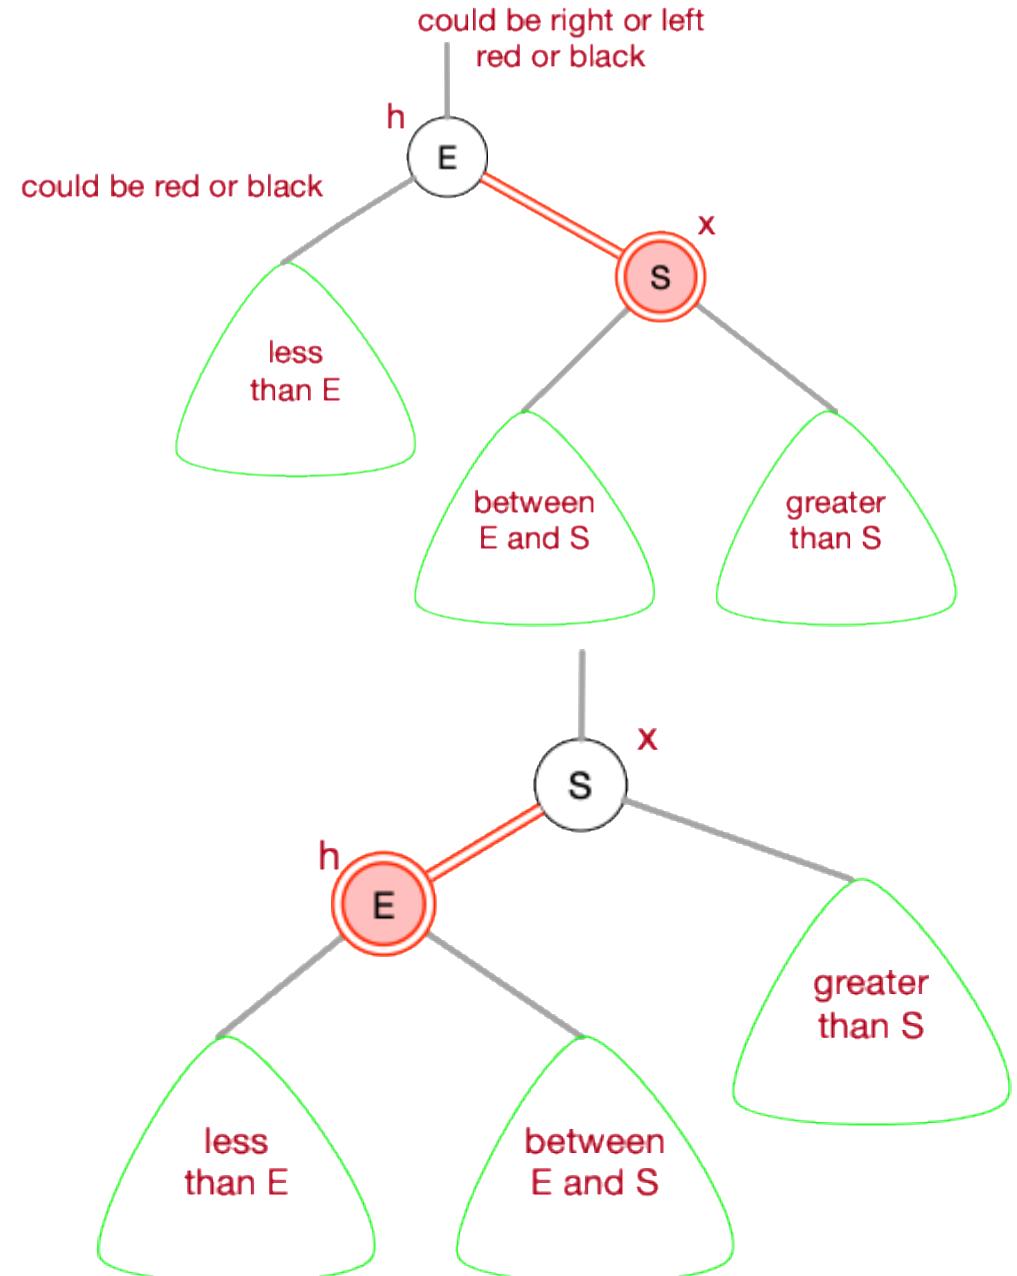
Red-Black Trees

- All the modifying operations (e.g. insertion) must respect the **invariant** properties of the tree, that is
 1. **Red links lean left**
 2. **No node has two red links** connected to it
 3. **The tree has perfect black balance**: every path to the root to a null-link has the same number of blank links
- But, sometimes, in the middle of them, they will allow
 - Right leaning red links
 - Two red links in a row
 - A node with two red links to both children
- There are two operations that correct this: **rotations** (two versions) and **colour flipping**.
 - All these operations are local ($\mathcal{O}(1)$) and can propagate up the tree, which has **logarithmic height**

Red-Black Trees

```
// make a right-leaning link lean to the left
private Node rotateLeft(Node h) {
    assert (h != null) && isRed(h.right);
    // assert (h != null) && isRed(h.right) && !isRed(h.left);
    // for insertion only
    Node x = h.right;
    h.right = x.left;
    x.left = h;
    x.color = h.color;
    h.color = RED;
    x.size = h.size;
    h.size = size(h.left) + size(h.right) + 1;
    return x;
}
```

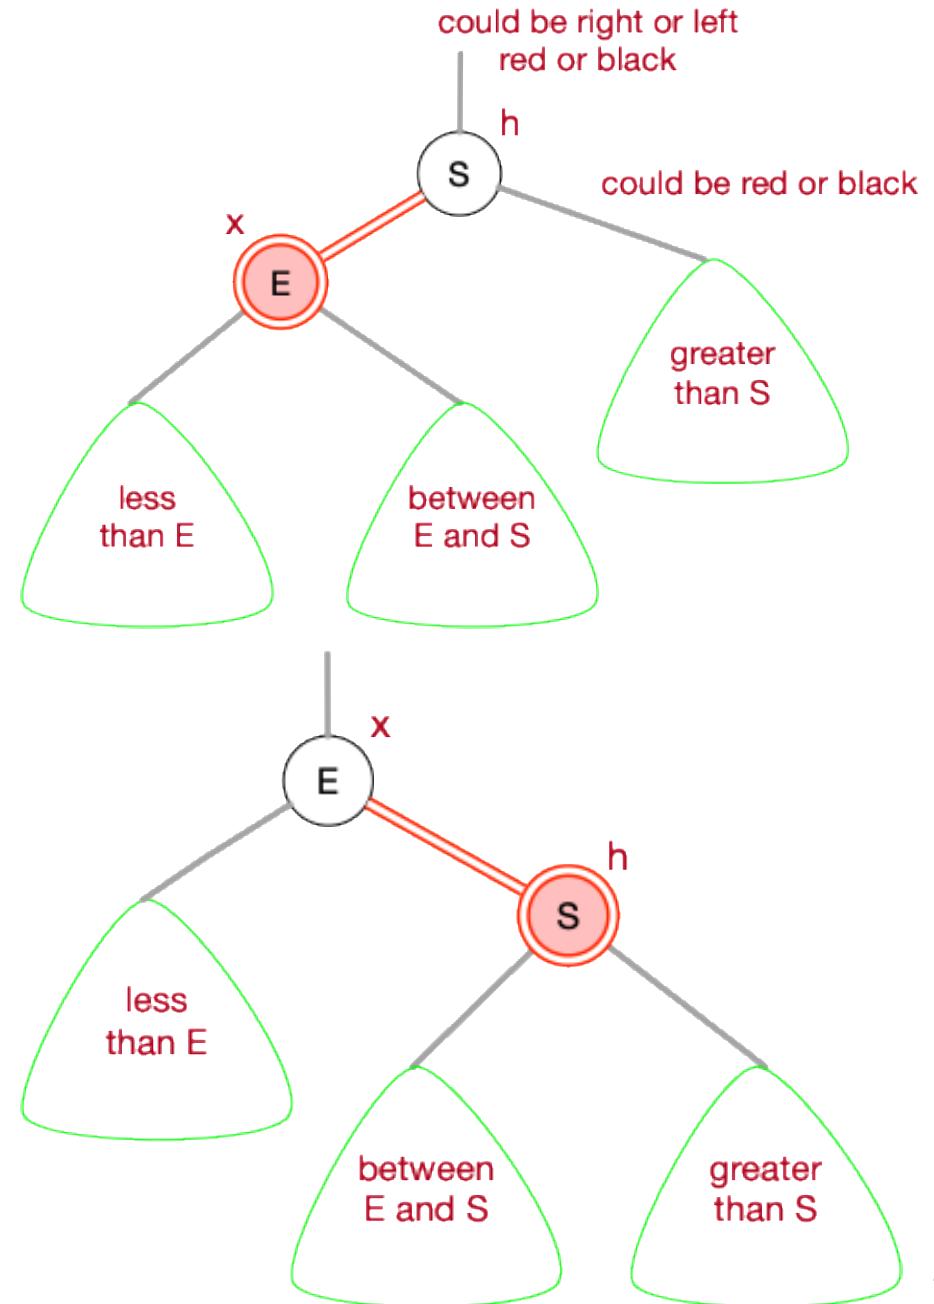
returns the new root



Red-Black Trees

```
// make a left-leaning link lean to the right
private Node rotateRight(Node h) {
    assert (h != null) && isRed(h.left);
    // assert (h != null) && isRed(h.left) && !isRed(h.right);
    // for insertion only
    Node x = h.left;
    h.left = x.right;
    x.right = h;
    x.color = h.color;
    h.color = RED;
    x.size = h.size;
    h.size = size(h.left) + size(h.right) + 1;
    return x;
}
```

returns the new root

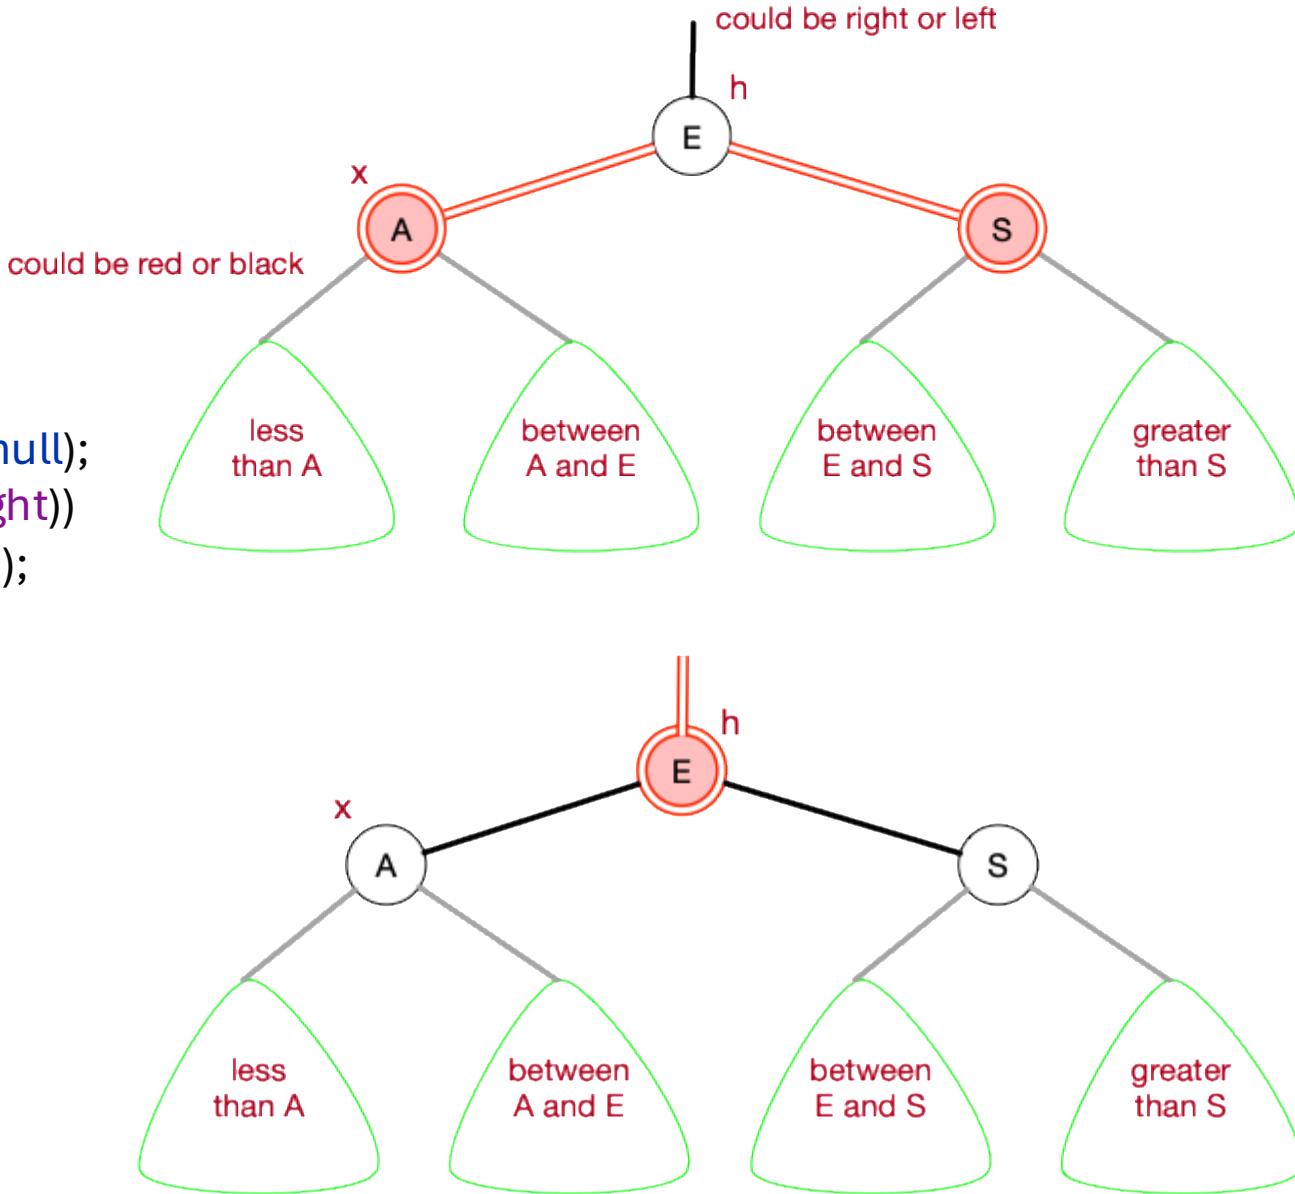


Red-Black Trees

```
// flip the colors of a node and its two children
private void flipColors(Node h) {
    // h must have opposite color of its two children
    assert (h != null) && (h.left != null) && (h.right != null);
    assert (!isRed(h) && isRed(h.left) && isRed(h.right))
        || (isRed(h) && !isRed(h.left) && !isRed(h.right));
    h.color = !h.color;
    h.left.color = !h.left.color;
    h.right.color = !h.right.color;
}
```

We can interpret flip as splitting a 4-node into two 2-nodes and inserting on the parent

- red link attaches middle node to parent
- black links are the resulting left and right 2-nodes after the split



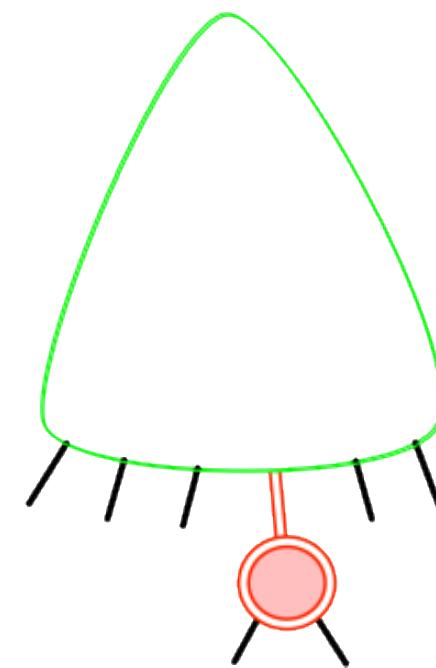
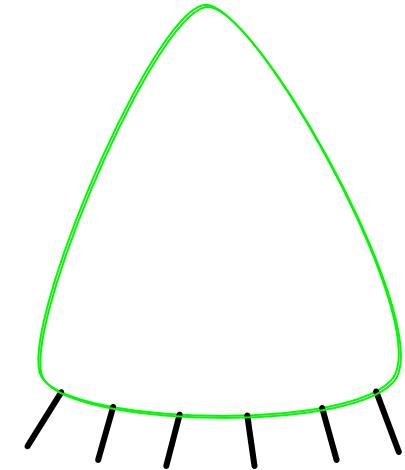
Red-Black Trees

- The insertion algorithm is almost the same as BST's
 - Doing a **top-down** search for finding the **insertion node**
 - If a new node must be created, it is inserted as a **leaf**
- The difference is that, **bottom-up**, when the recursive calls return, the algorithm will **restore the invariant properties** if they're violated

```
public void put(Key key, Value val) {  
    if (key == null)  
        throw new IllegalArgumentException("key is null");  
  
    root = put(root, key, val);  
    root.color = BLACK;  
}  
  
// insert the key-value pair in the subtree rooted at h  
private Node put(Node h, Key key, Value val) {  
    if (h == null) return new Node(key, val, RED, 1);  
  
    int cmp = key.compareTo(h.key);  
    if      (cmp < 0) h.left  = put(h.left,  key, val);  
    else if (cmp > 0) h.right = put(h.right, key, val);  
    else            h.val   = val;  
  
// TODO: restore invariant  
  
    return h;  
}
```

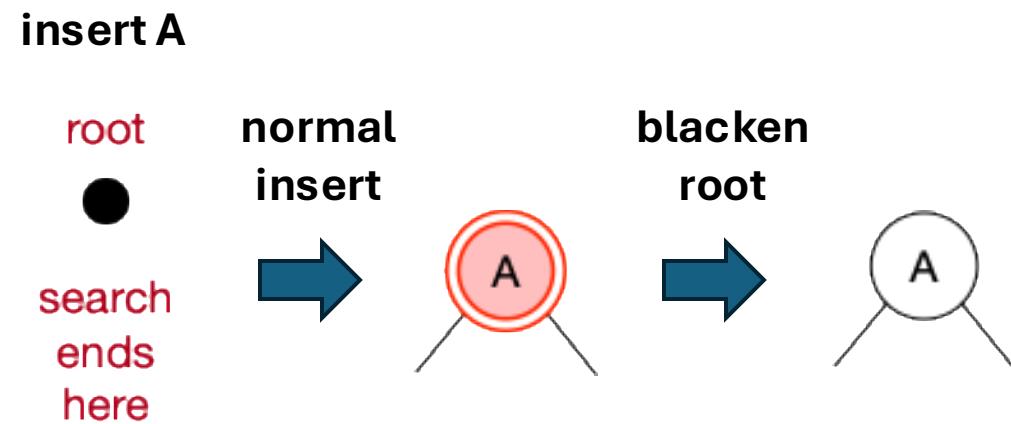
Red-Black Trees

- Inserting the **new node** as **RED**, ensured that the **perfect black balance** property is not violated
 - All null-links continue to be at the same black distance to root
- After the insertion, the **root node** is always coloured as **BLACK**
 - A red root node makes no sense



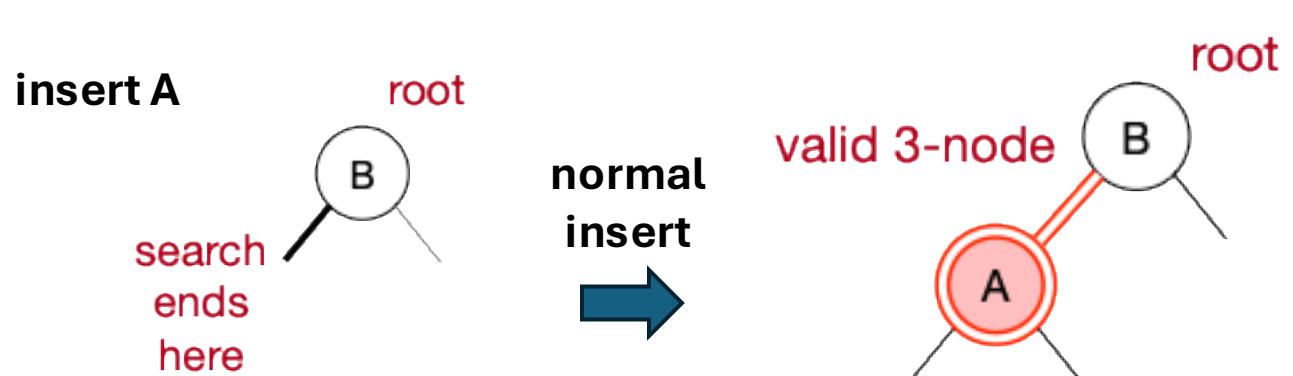
Red-Black Trees

- Let's begin with three easy cases:
 - Inserting to an **empty tree**:
 - simply create the new root and make it black



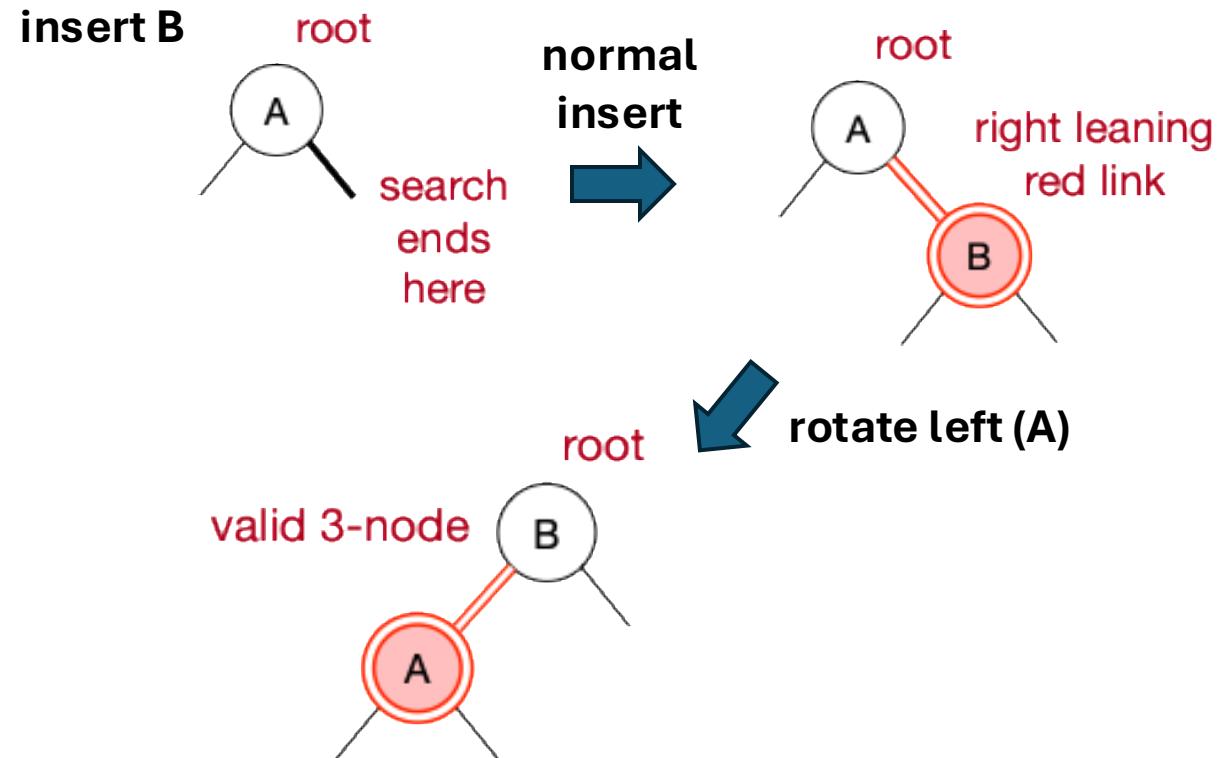
Red-Black Trees

- Let's begin with three easy cases:
 - Inserting to an **empty tree**:
 - simply create the new root and make it black
 - Inserting to a tree with a **single 2-node**
 - Insert as **left child**:
 - the new RED node makes the root node a 3-node



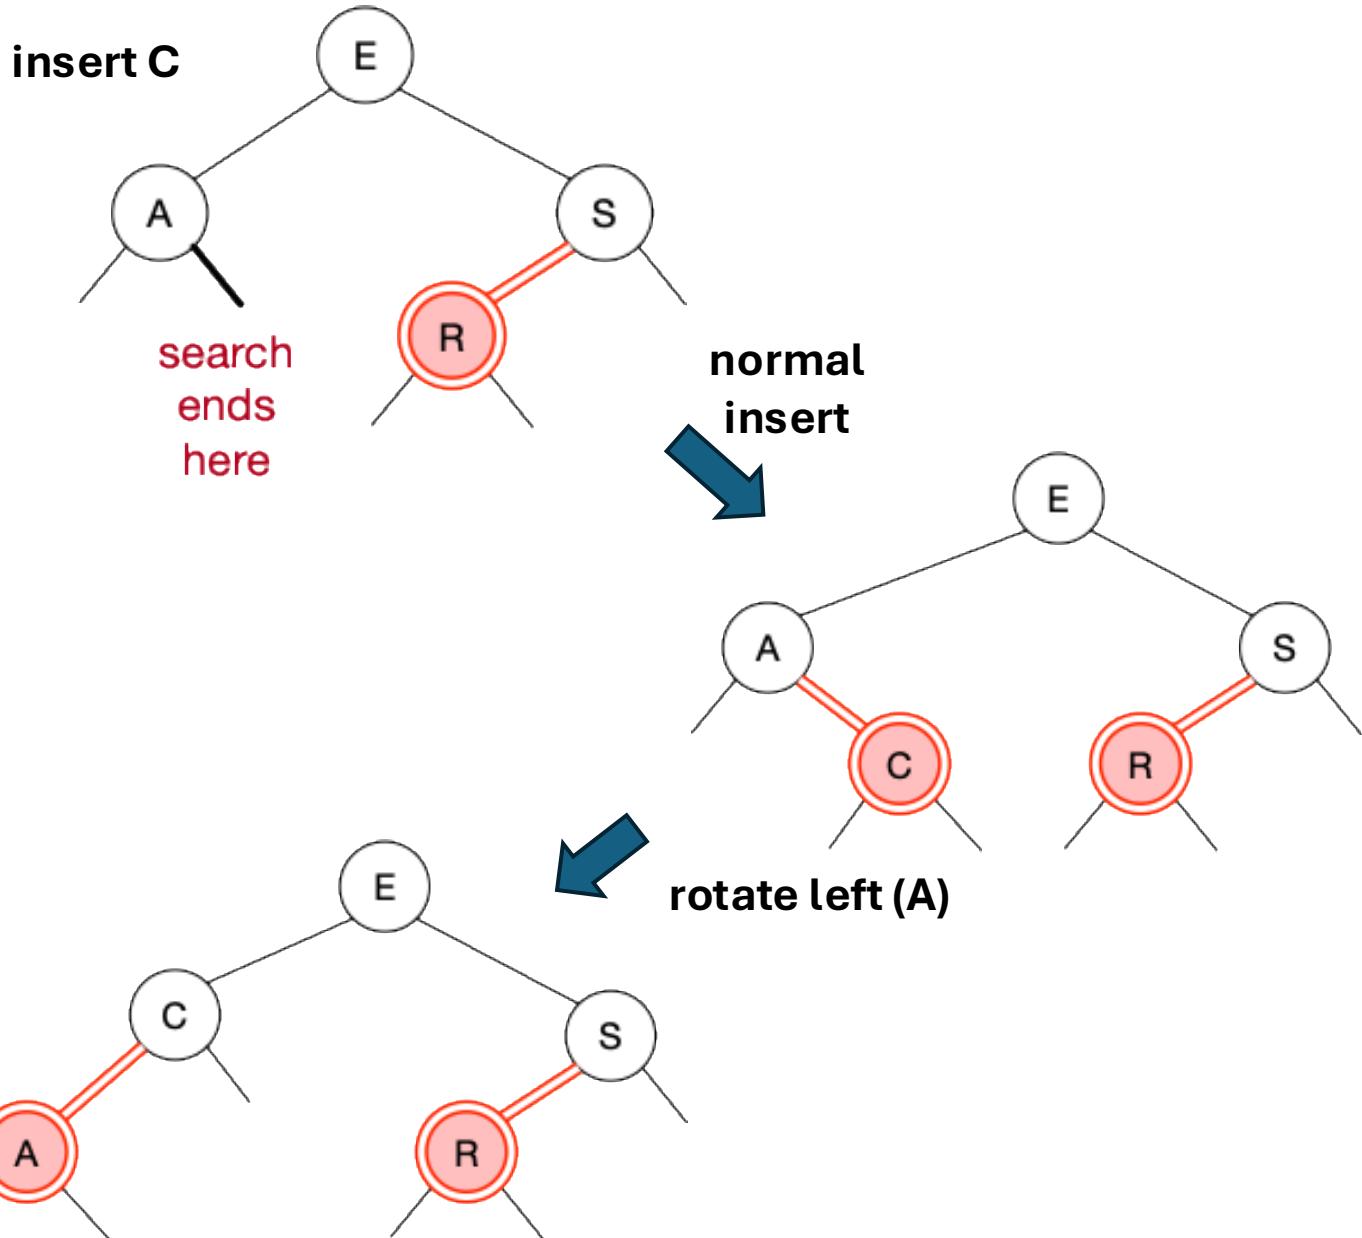
Red-Black Trees

- Let's begin with three easy cases:
 - Inserting to an **empty tree**:
 - simply create the new root and make it black
 - Inserting to a tree with a **single 2-node**
 - Insert as **left child**:
 - the new RED node makes the root node a 3-node
 - Insert as **right child**:
 - the new RED node makes a right leaning link
 - so, we need to do a **left rotation** (of the parent of the new node)



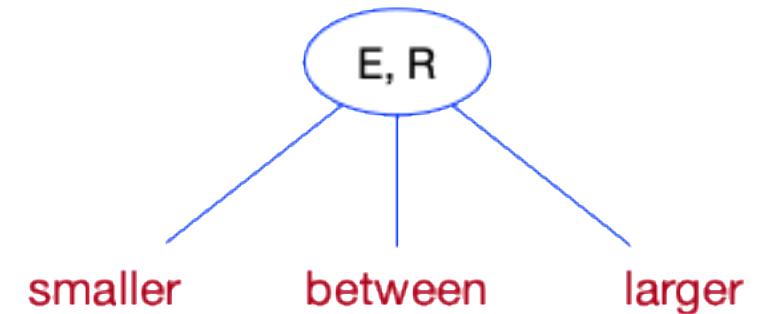
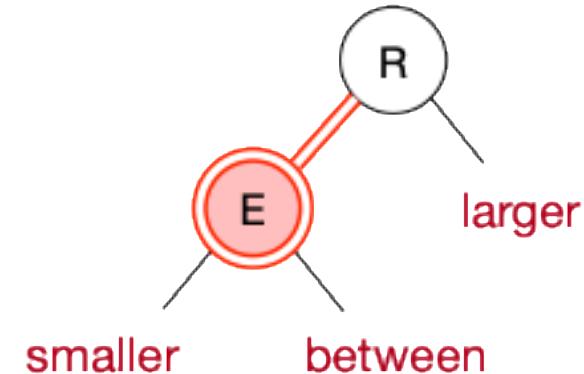
Red-Black Trees

- The same situation happens when adding to 2-nodes at the bottom of the tree.
- Let's show only the case when a left rotation of the parent of the new node is needed



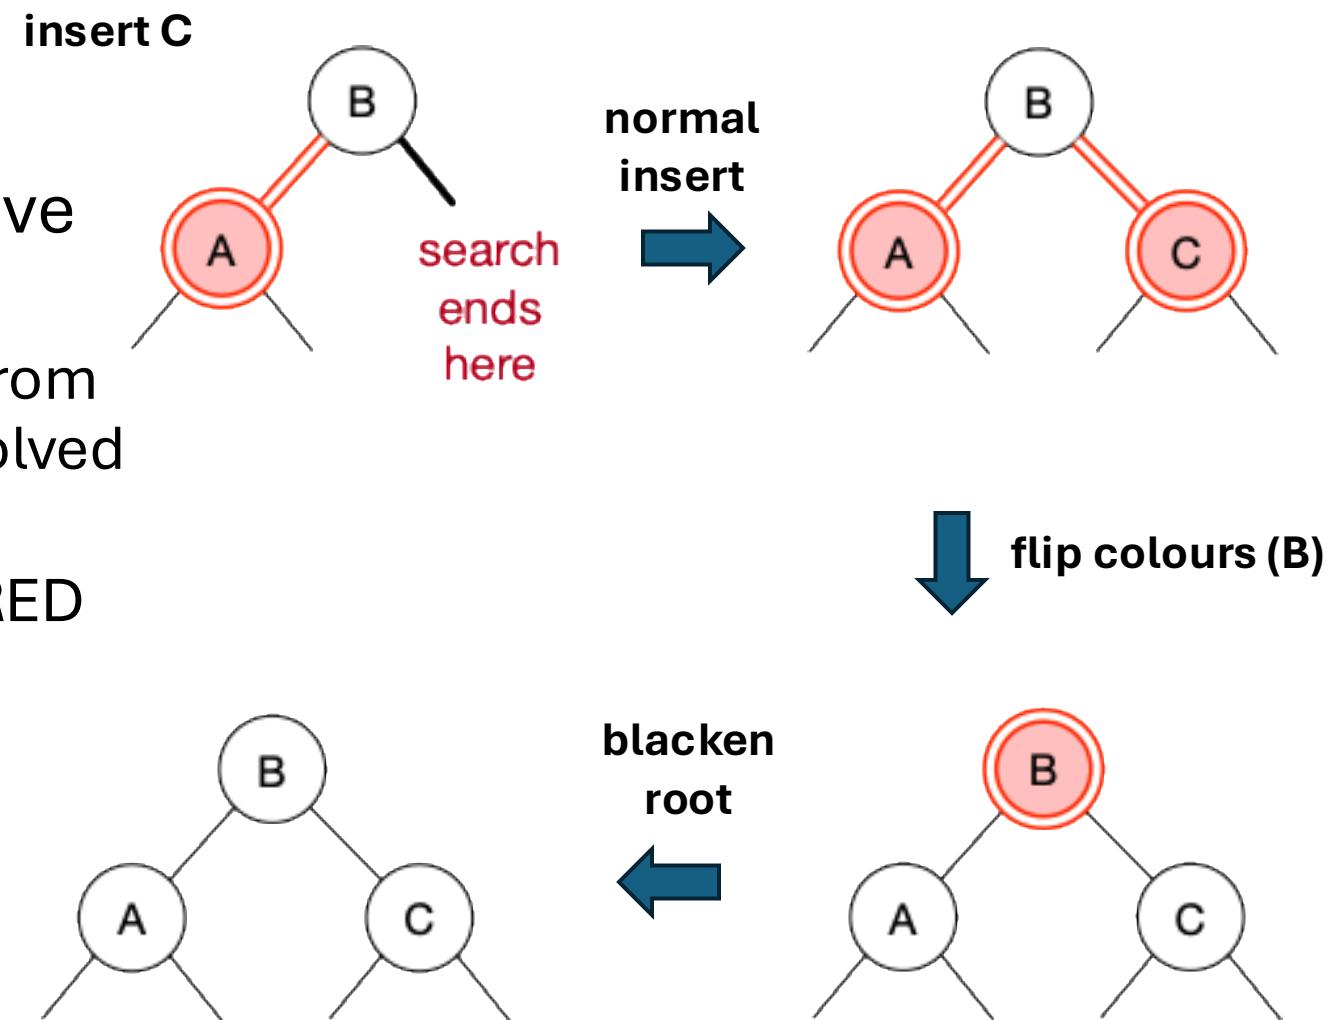
Red-Black Trees

- Let's consider when the tree is a single 3-node
- There are three possibilities in this case:
 - The new key is **larger** than those in the tree
 - The new key is **smaller** than those in the tree
 - The new key is **between** those in the tree
- NOTE: These are the same cases we had when analysing insertion in 2-3 Trees



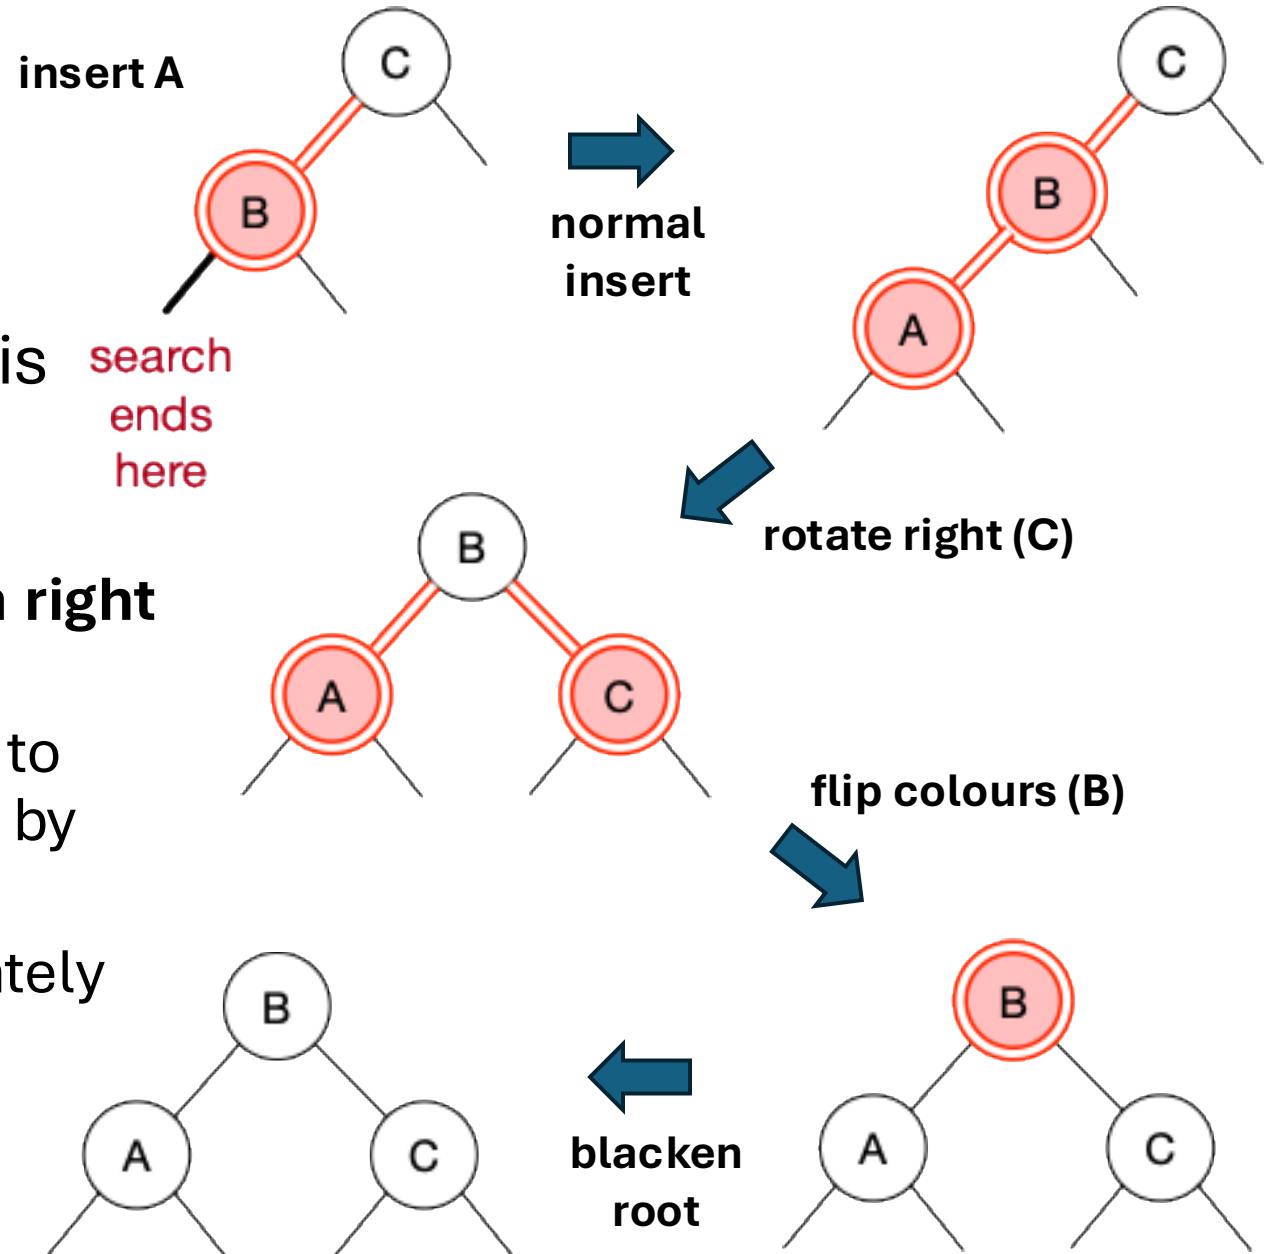
Red-Black Trees

- When the key is **larger**, we have the easiest case
 - We have to RED links coming from the same parent that can be solved by **flipping the colours**
 - The **root** is temporarily made RED
 - But it's immediately restored restored to **BLACK**



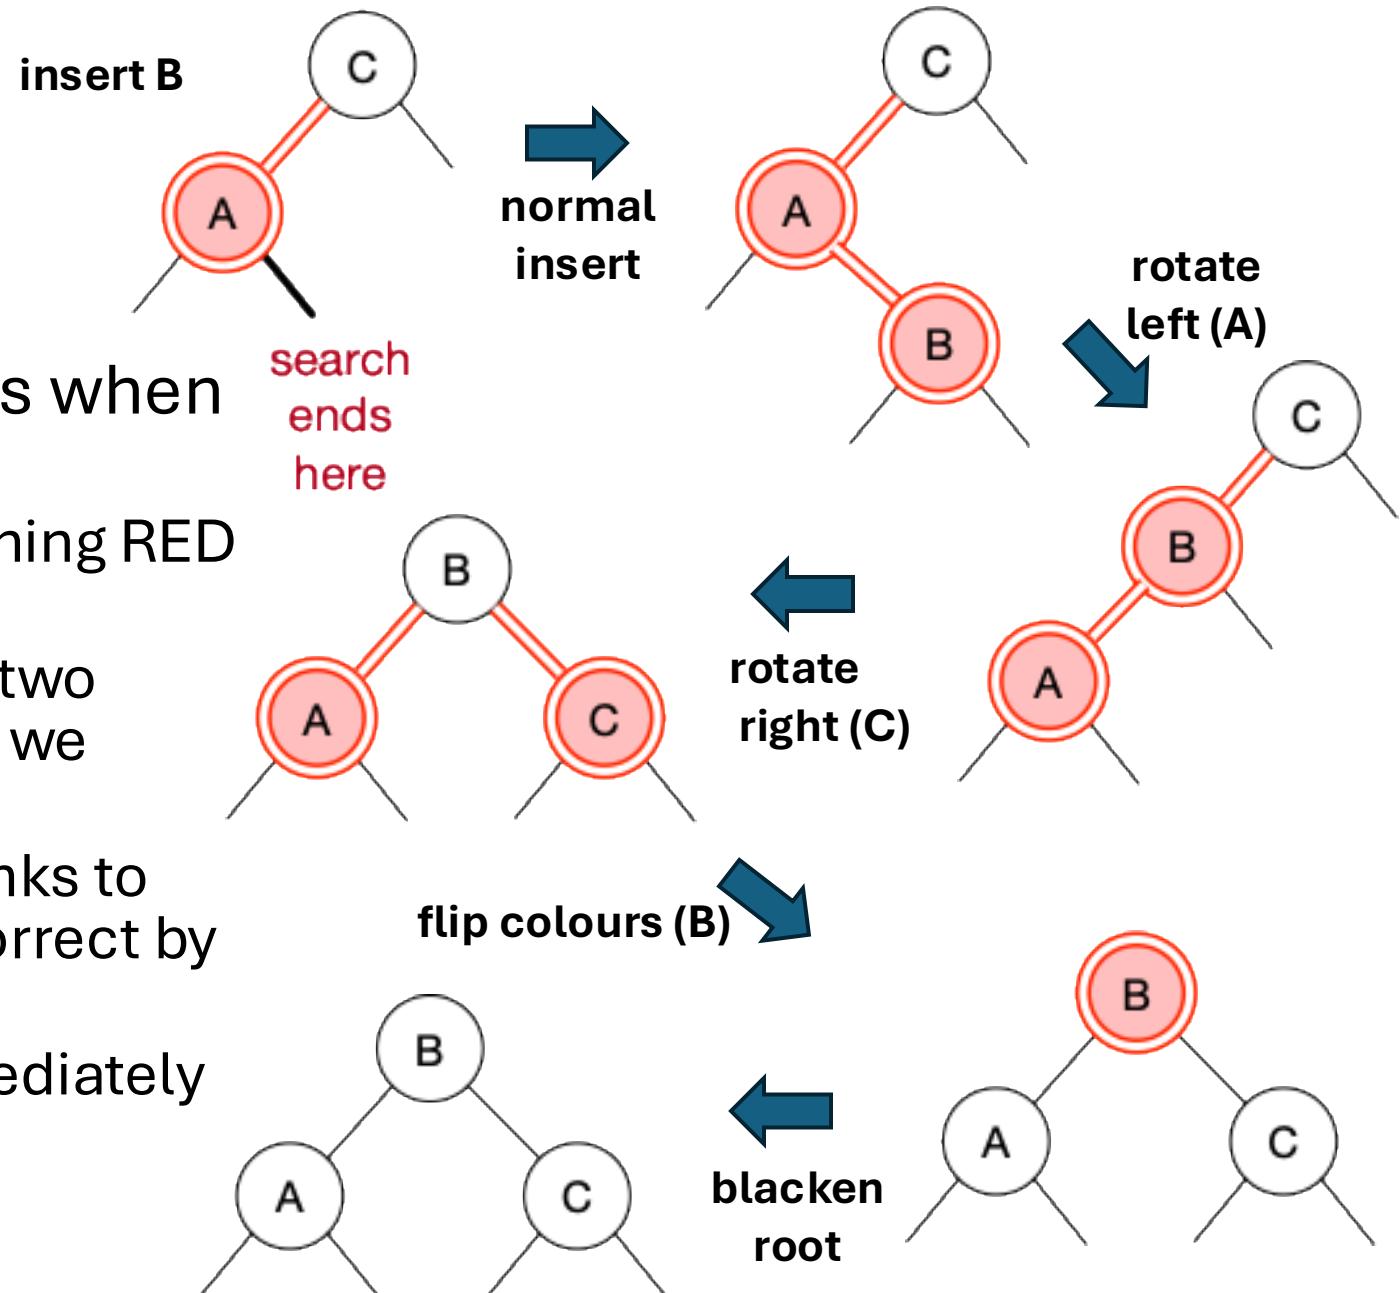
Red-Black Trees

- The next case we'll consider is when the key is **smaller**
 - First, to avoid a chain of two consecutive RED links we do a **right rotation**
 - But this creates two RED links to the same parent that we solve by **flipping colours**
 - As before, the **root** is immediately changed to **BLACK**



Red-Black Trees

- The last case to consider is when the key is in **between**
 - First, we solve the right leaning RED link with a **left rotation**
 - This causes two a chain of two consecutive RED links that we solve by a **right rotation**
 - But this creates two RED links to the same parent that we correct by **flipping the colours**
 - As always, the **root** is immediately changed to **BLACK**

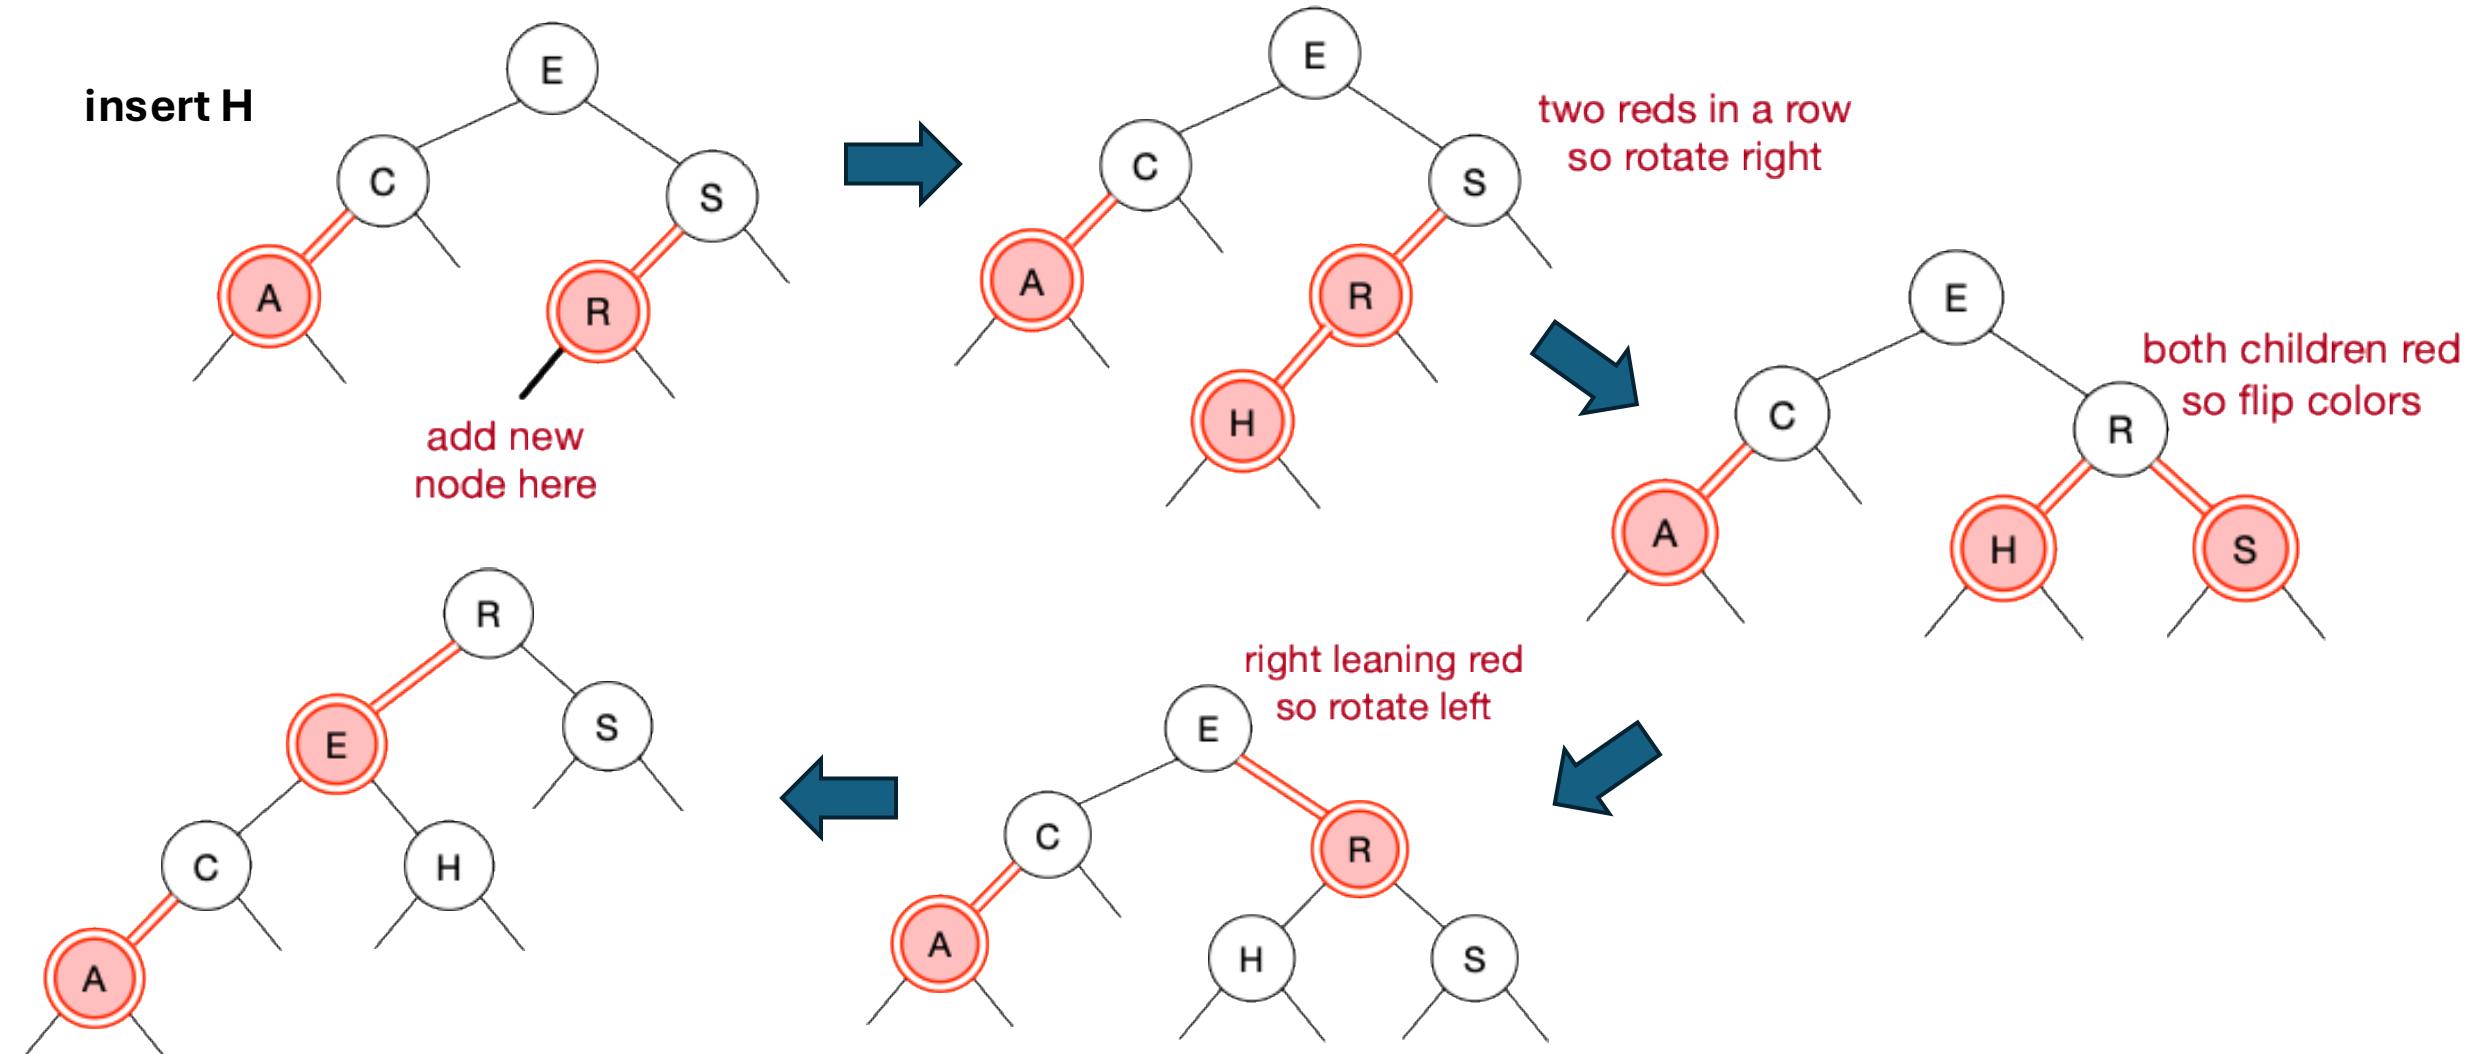


Red-Black Trees

- The only case we must consider is **inserting in a 3-Node** that is at the **bottom** of the tree
- Locally, we'll have the **same three cases**, and we'll proceed as when the 3-Node was the only node in the tree (the root)
 - NOTE: **All the three cases ended by a flipping colours operation**
- So, what we'll be the **difference** if any?
 - That now, as **the node is no longer the root**, it **won't be changed to BLACK**
 - So, a **RED link is propagated up the tree**
- How will it be treated?
 - By the **beauty of recursion**: from the point of view of the parent, it will be as it was inserted as a new node (which are always RED)
 - So, the very structure of the recursion will take care of that !!

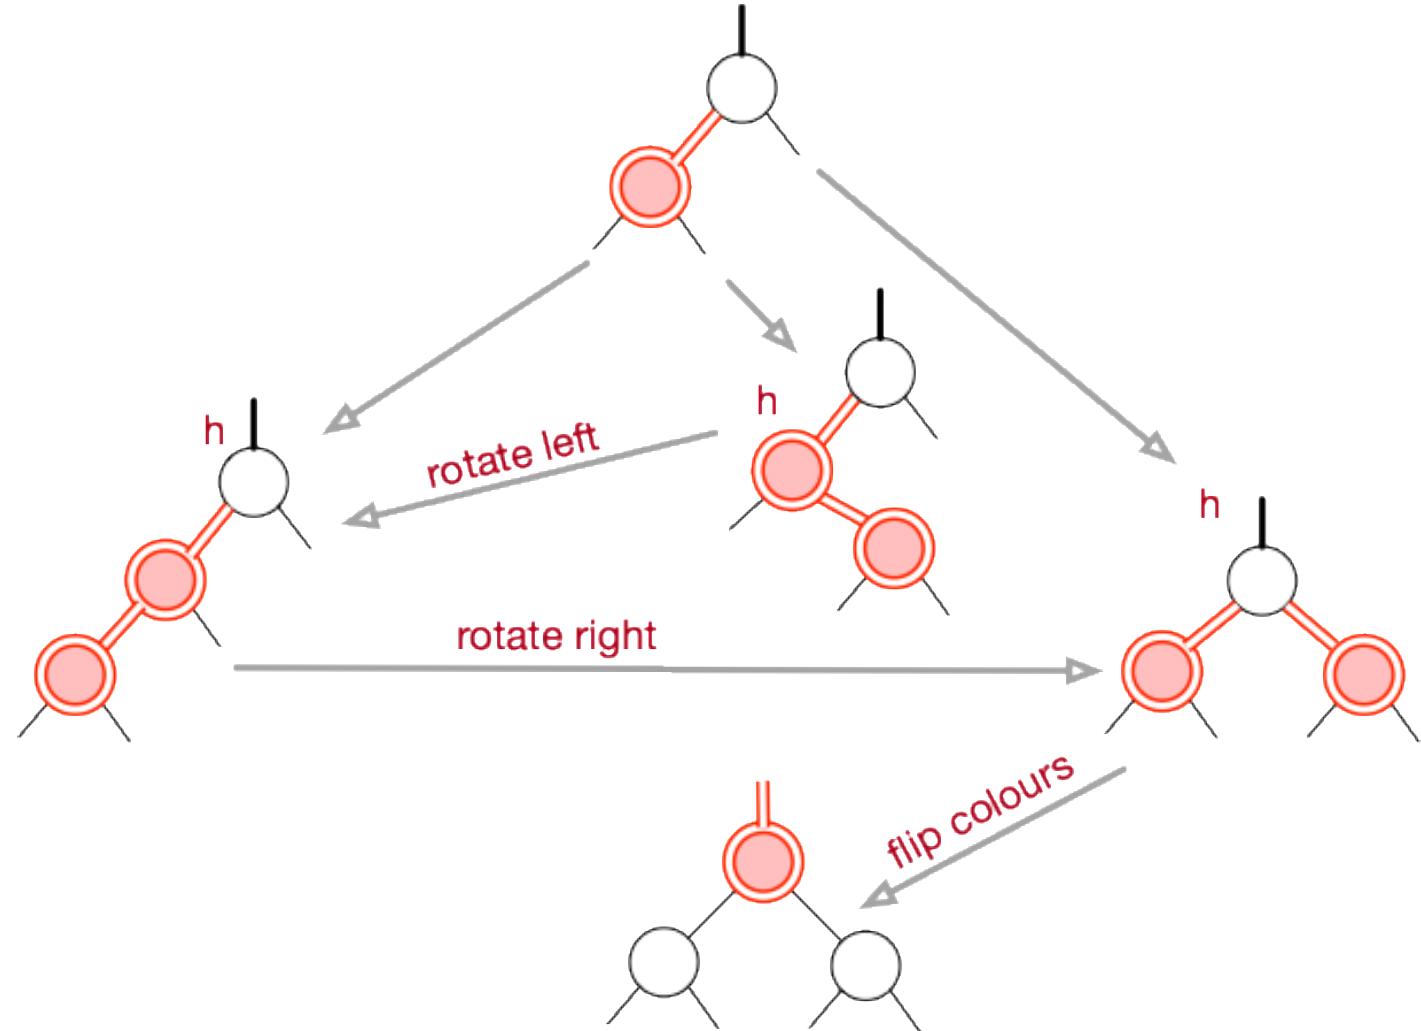
Red-Black Trees

insert H



Red-Black Trees

We can sum-up all these three cases, the three possible insertion points, the transformations and the passing up of the RED link in a single diagram



NOTE: The returns to callers (parent node) are not shown and understanding them is key to mastering the algorithm.

Red-Back Trees

- After doing the normal top-down insertion in a BST, we arrange bottom-up the invariant violations that we may have created
- And thanks to recursion,
 - when we produce a problem up to the tree ...
 - ... it will be taken care of when the current recursive call is ended ...
 - ... in the call corresponding to its parent

```
private Node put(Node h, Key key, Value val) {  
    if (h == null) return new Node(key, val, RED, 1);  
  
    int cmp = key.compareTo(h.key);  
    if      (cmp < 0) h.left  = put(h.left, key, val);  
    else if (cmp > 0) h.right = put(h.right, key, val);  
    else           h.val   = val;  
  
    // fix-up any right-leaning links  
    if (isRed(h.right) && !isRed(h.left))    h = rotateLeft(h);  
    if (isRed(h.left)  && isRed(h.left.left)) h = rotateRight(h);  
    if (isRed(h.left)  && isRed(h.right))   flipColors(h);  
  
    h.size = size(h.left) + size(h.right) + 1;  
  
    return h;  
}
```

B-Trees

B-Trees

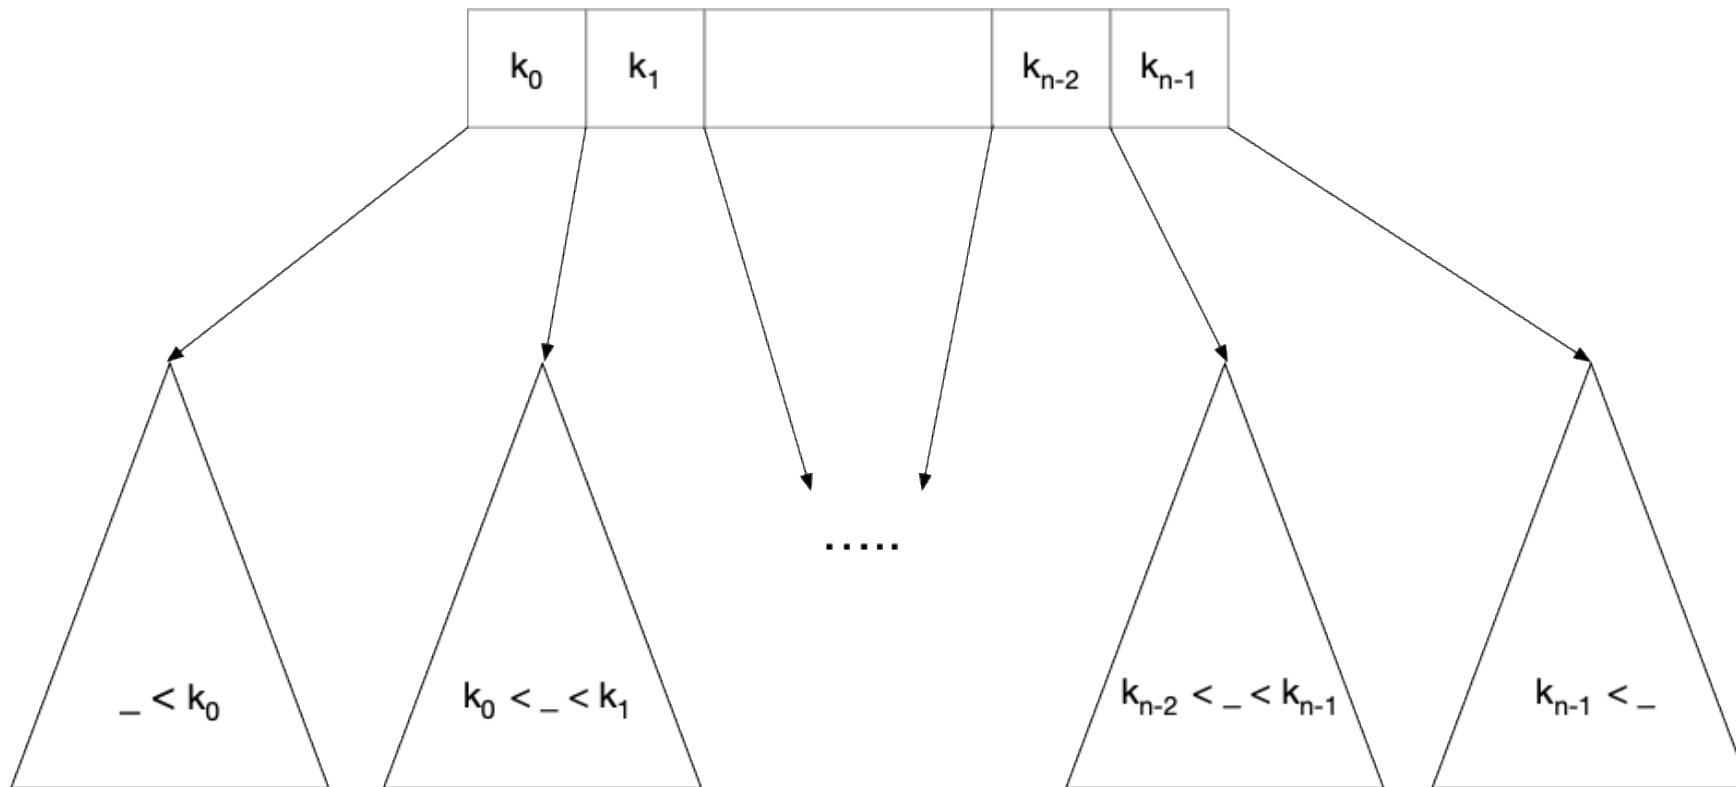
- B-Trees are balanced search trees designed to work well on **disk drives** or other **direct-access secondary storage devices**
 - They're like red-black trees, but they are better at minimizing the number of data access operations
 - They differ from red-black trees in that nodes can have many children (from a few to hundreds), i.e. **the branching factor can be quite large**
- So, B-Trees
 - Generalize binary search trees
 - And their insert and remove operations leave the B-Tree balanced
 - So, the height of the tree is $\mathcal{O}(\log n)$
- As we did in the case of binary search trees, we'll only show the keys in the tree because the value associated to it is simply a payload.

B-Trees

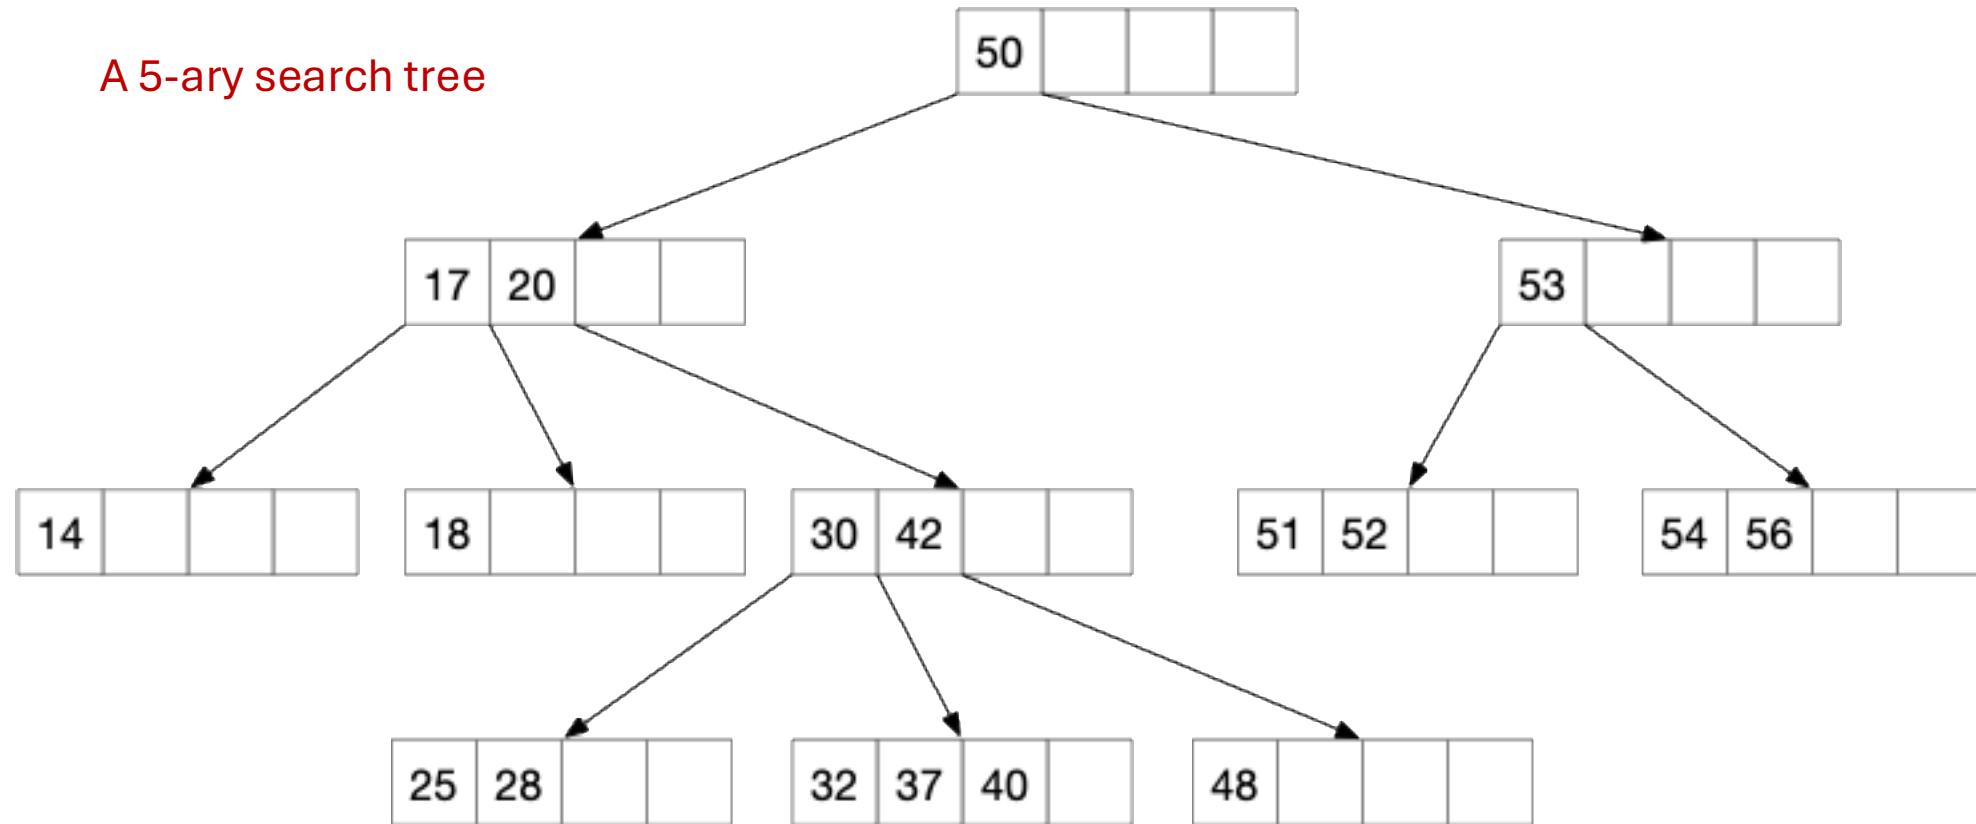
- An **m-ary search tree** generalizes binary search trees by
 - **m** is the order of the tree, that is, the maximum degree of any node
 - **n** is the number of keys associated with any node ($n < m$) and ($n + 1$) is the number of children of this node
 - k_0, k_1, \dots, k_{n-1} are the keys associated with a node
 - The keys at each node are sorted increasingly so $k_i < k_{i+1}$
 - The properties of a search tree are respected, so that for a key k_i
 - All keys in the **first i subtrees** are **smaller**
 - All keys in the **last $n - i$ subtrees** are **bigger**

B-Trees

$$k_0 < k_1 < \dots < k_{n-2} < k_{n-1}$$



B-Trees



B-Trees

- m-ary search trees have the same problem as simple binary search trees
 - They can be unbalanced
 - So, insert / delete / search operations are not logarithmic
- **B-Trees are balanced m-ary search trees**
- They were developed by Rudolf Bayer & Edward W, McCreight in 1970
 - [Organization and Maintenance of Large Ordered Indices](#), SIGFIDET Workshop 1970: 107-141
- **NOTE:** There are many variations of them, and we'll follow Ribó's presentation

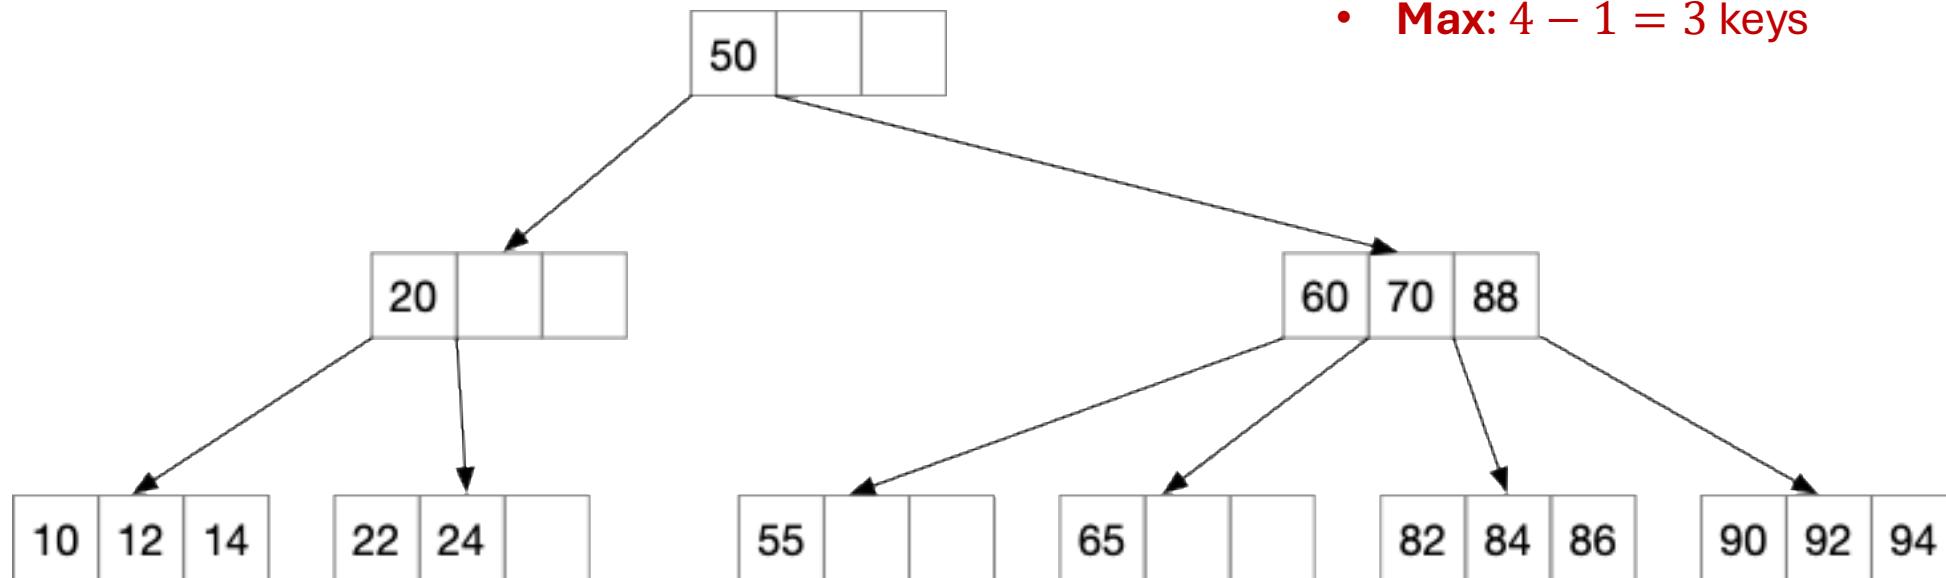
B-Trees

- A **B-Tree** of **order m** (with $m \geq 3$) is a **m -ary search tree** such that
 - **The root** must have **at least 2** children (and 1 key)
 - Unless it is a leaf, or the b-tree is empty
 - **All non-root nodes** must have **at least $\left\lceil \frac{m}{2} \right\rceil - 1$ keys**
 - So, if it's an **internal node**, it'll have **at least $\left\lceil \frac{m}{2} \right\rceil$ children**
 - **All nodes** will have **at most $m - 1$ keys**
 - **All leaves** are at the **same level**
- So, the cost of searching for a key in a B-Tree of order m and size n is $O(\log_{\left\lceil \frac{m}{2} \right\rceil} n)$

B-Trees

B-Tree of order 4

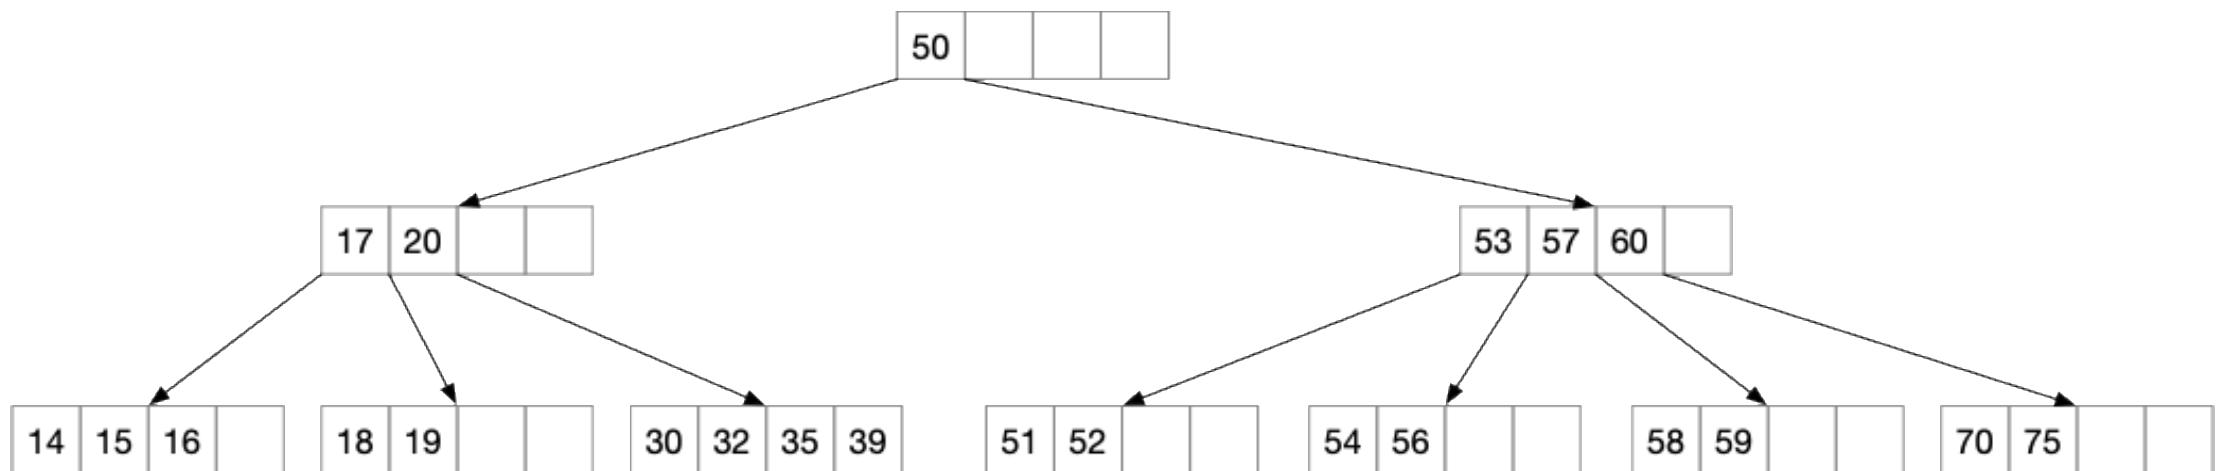
- Min: $\left\lceil \frac{4}{2} \right\rceil - 1 = 2 - 1 = 1$ key
- Max: $4 - 1 = 3$ keys



B-Trees

B-Tree of order 5

- Min: $\left\lceil \frac{5}{2} \right\rceil - 1 = 3 - 1 = 2$ keys
- Max: $5 - 1 = 4$ keys



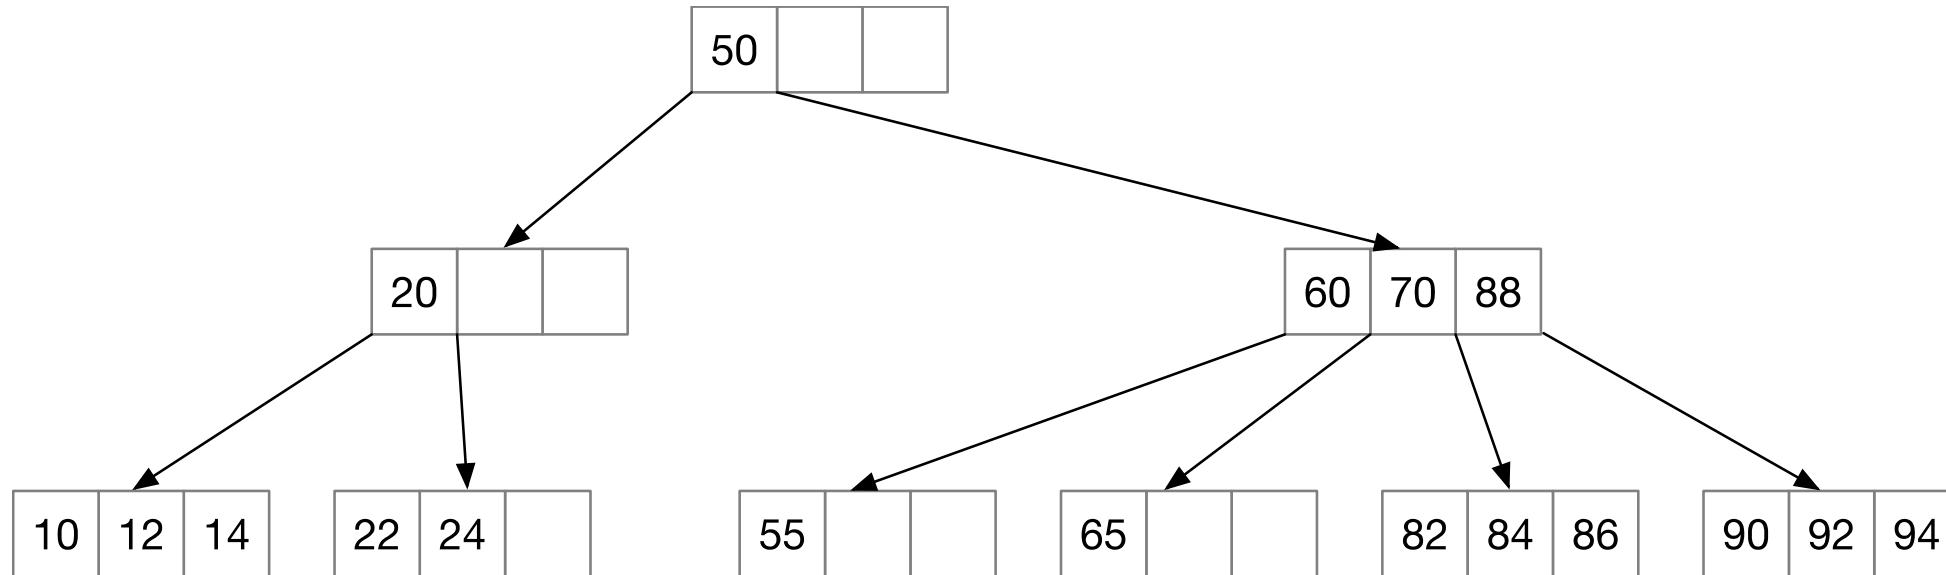
B-Trees

- **Insertion of the pair: $k \rightarrow v$**
 1. Search the key k in the tree
 2. If the key is found, substitute the current value associated to it by v
 3. If not, consider the node (a leaf) h in which to add the pair (k, v)
 4. If the node h is not full, we add the pair to it
 5. If it is,
 - i. Split the leaf in two, considering also the new key to be added.
The split consist in taking the median value, leaving the smaller keys in the node and creating a new node with the bigger keys.
 - ii. The median value (with the new node attached to it as right child) is inserted in the parent node, and the steps 3 & 4 are repeated with the parent as the h node (which now it's not a leaf)
NOTE: When moving up the key, if the key had already a right child, as its right child must be the new node, we must perform a *little arrangement* (a kind of *rotation*)

B-Trees

Let's consider the following B-Tree (**order $m = 4$**)

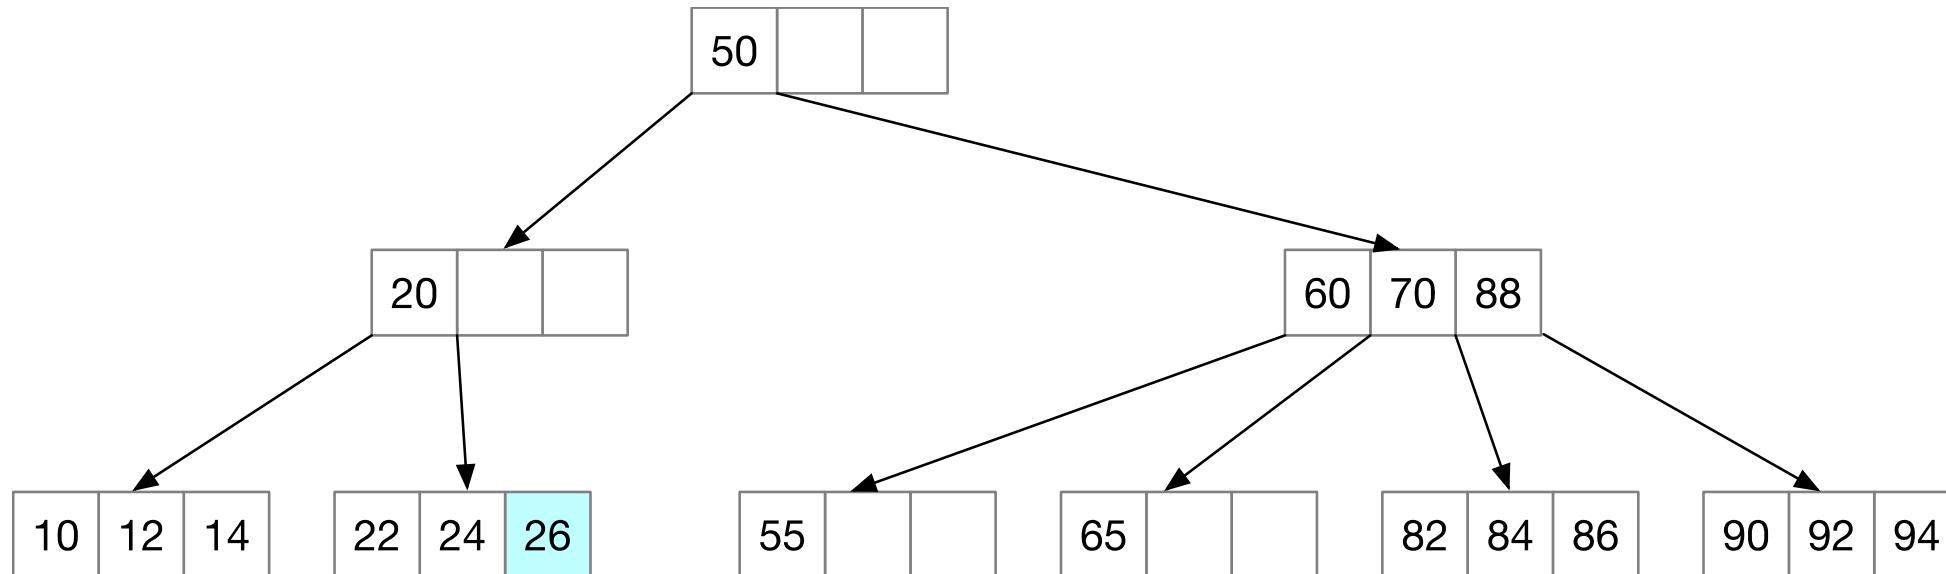
- **Minimum number of keys** = $\left\lceil \frac{4}{2} \right\rceil - 1 = 2 - 1 = 1$



B-Trees

Let's **insert key 26**

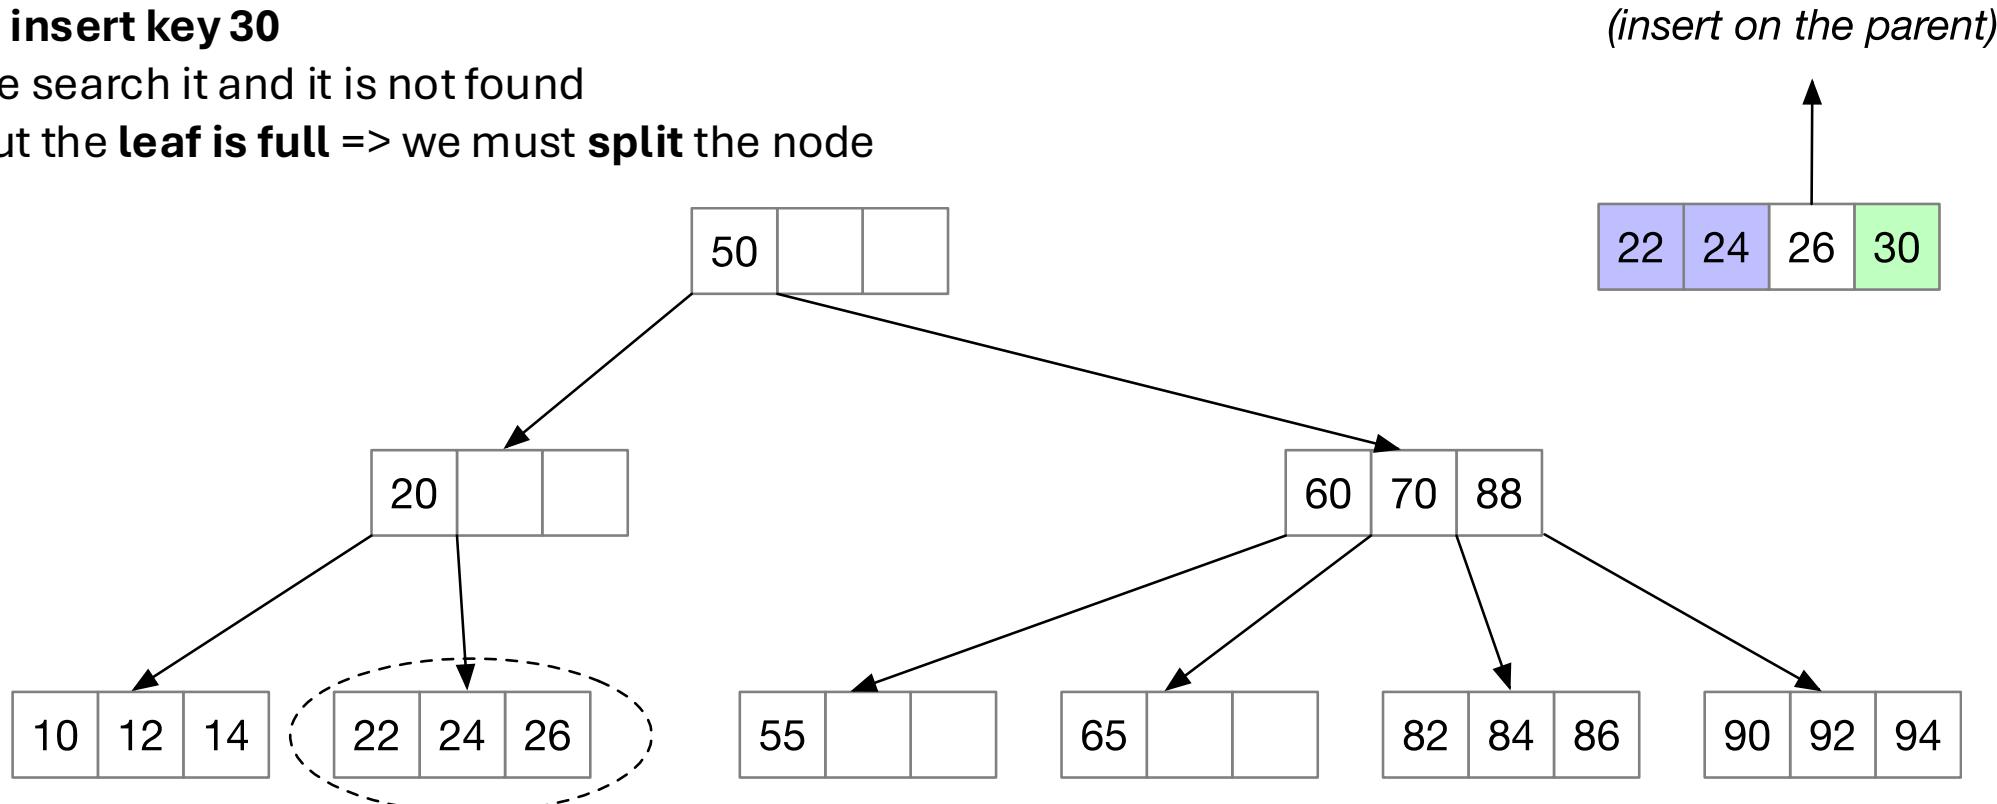
- we search it and it is not found
- **leaf is not full**, so we **add 26** to it



B-Trees

Let's **insert key 30**

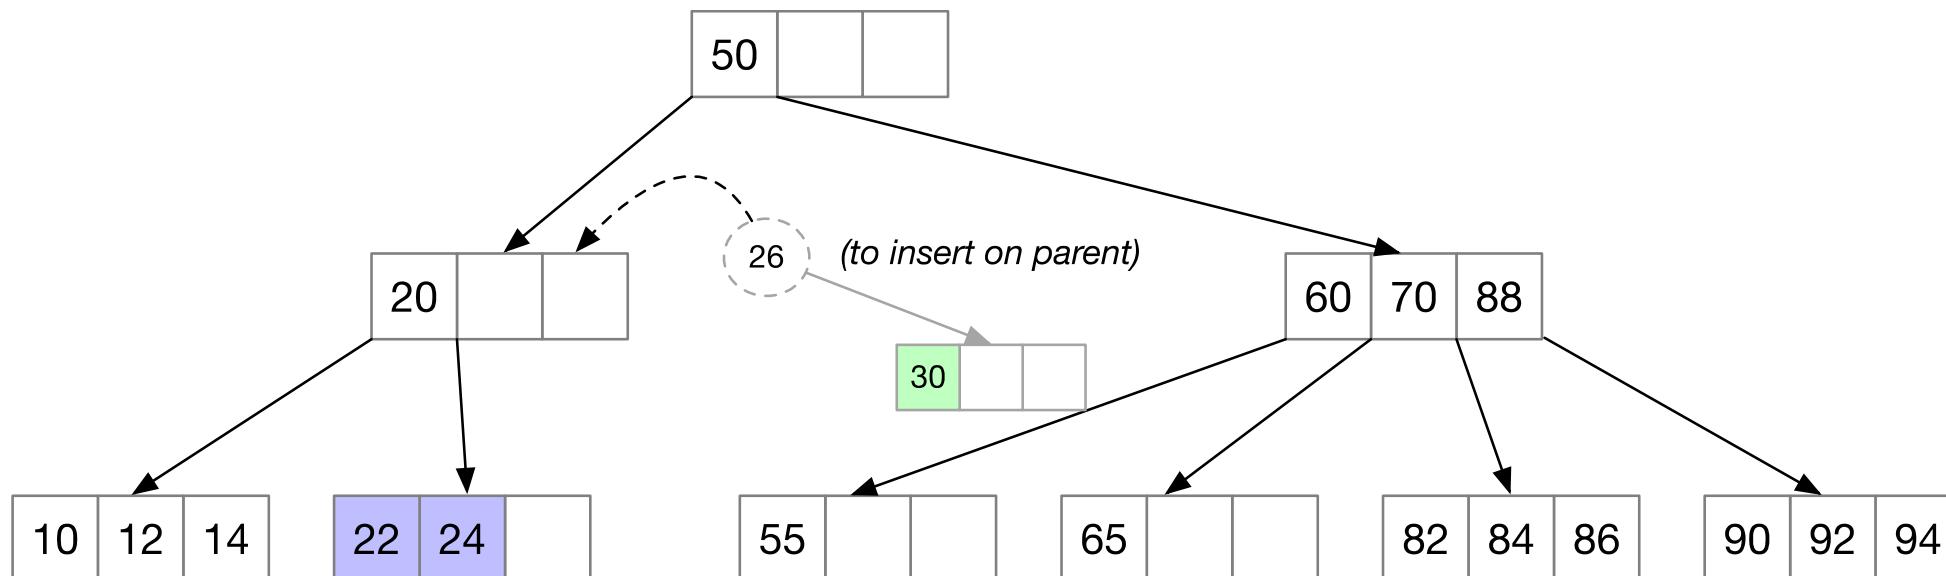
- we search it and it is not found
- but the **leaf is full** => we must **split** the node



B-Trees

Let's **insert key 30**

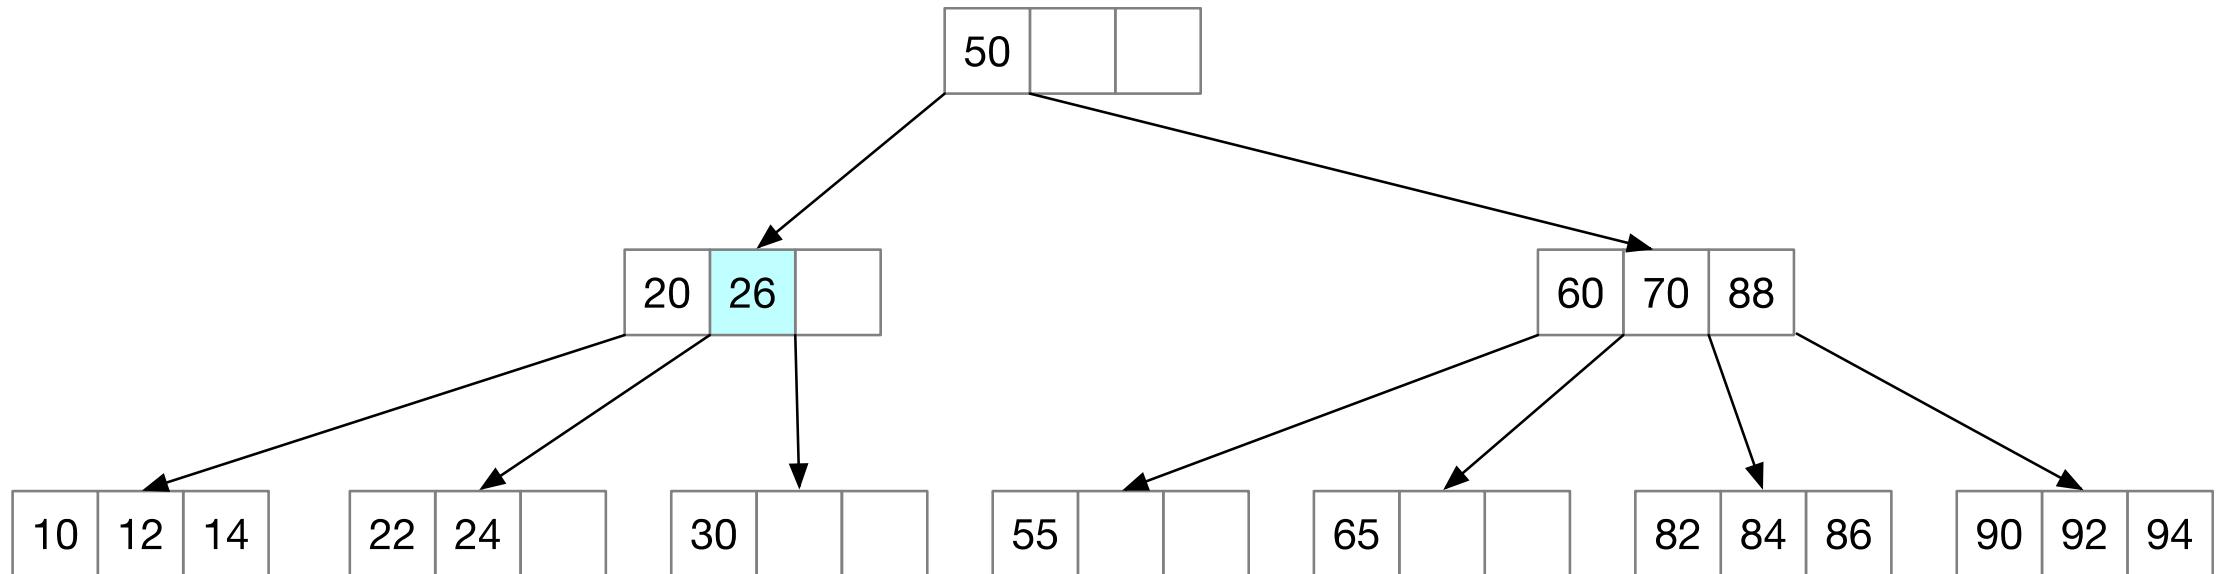
- we search it and it is not found
- but the **leaf is full** => we must **split** the node
- we **insert 26 on parent**
 - as 26 was on a left, it had no right child, so we don't need any rearrangement



B-Trees

Let's **insert key 30**

- Let's **insert key 26** on parent
 - The node is **not full**, so we **add 26** to it

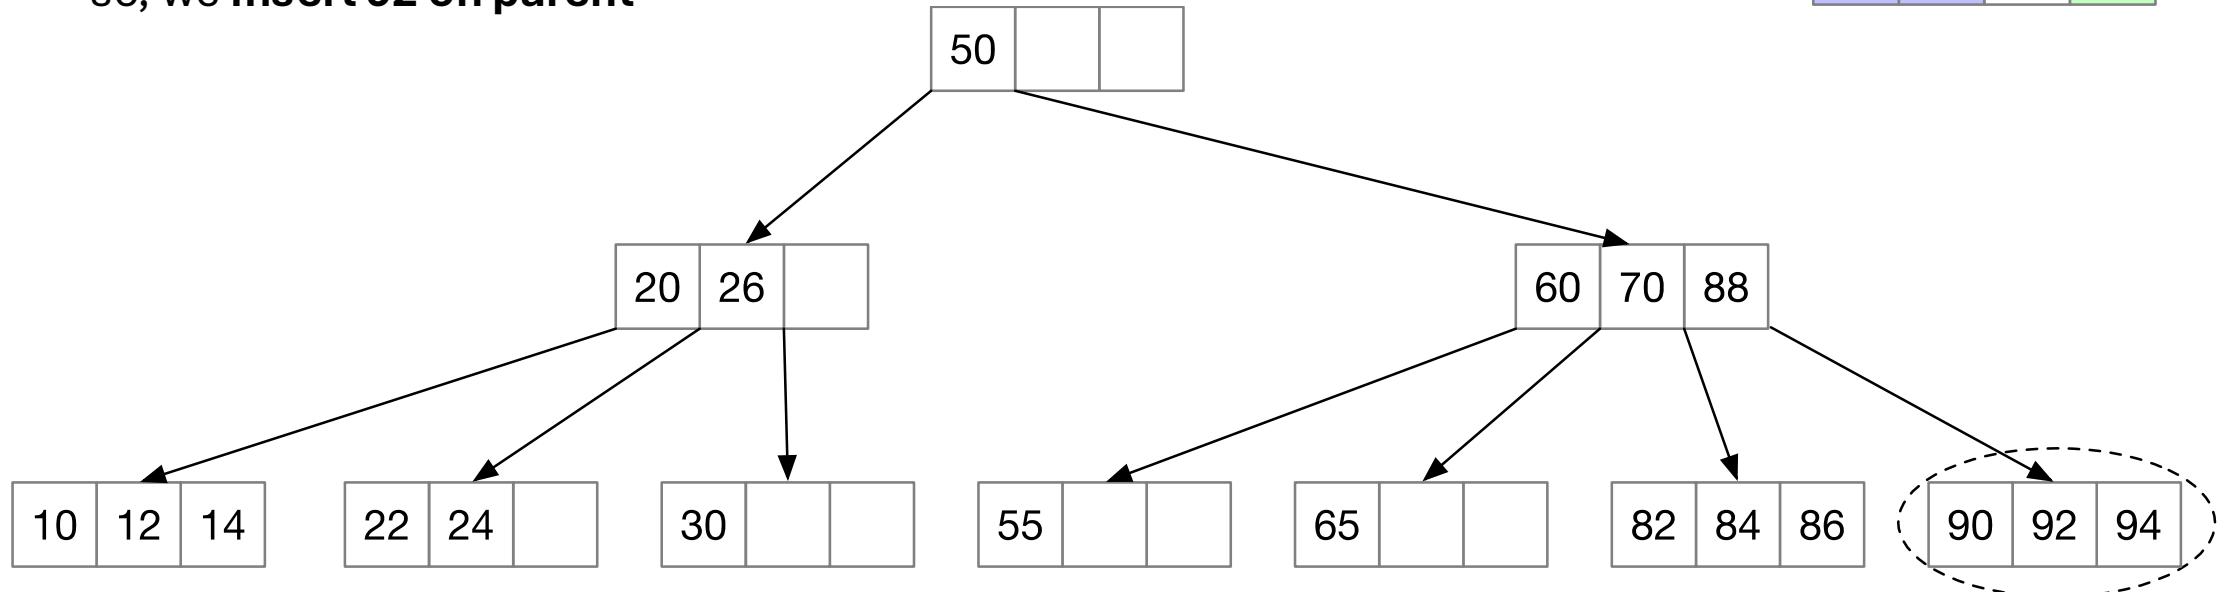


B-Trees

(to insert on parent)

Now, let's **insert key 91**

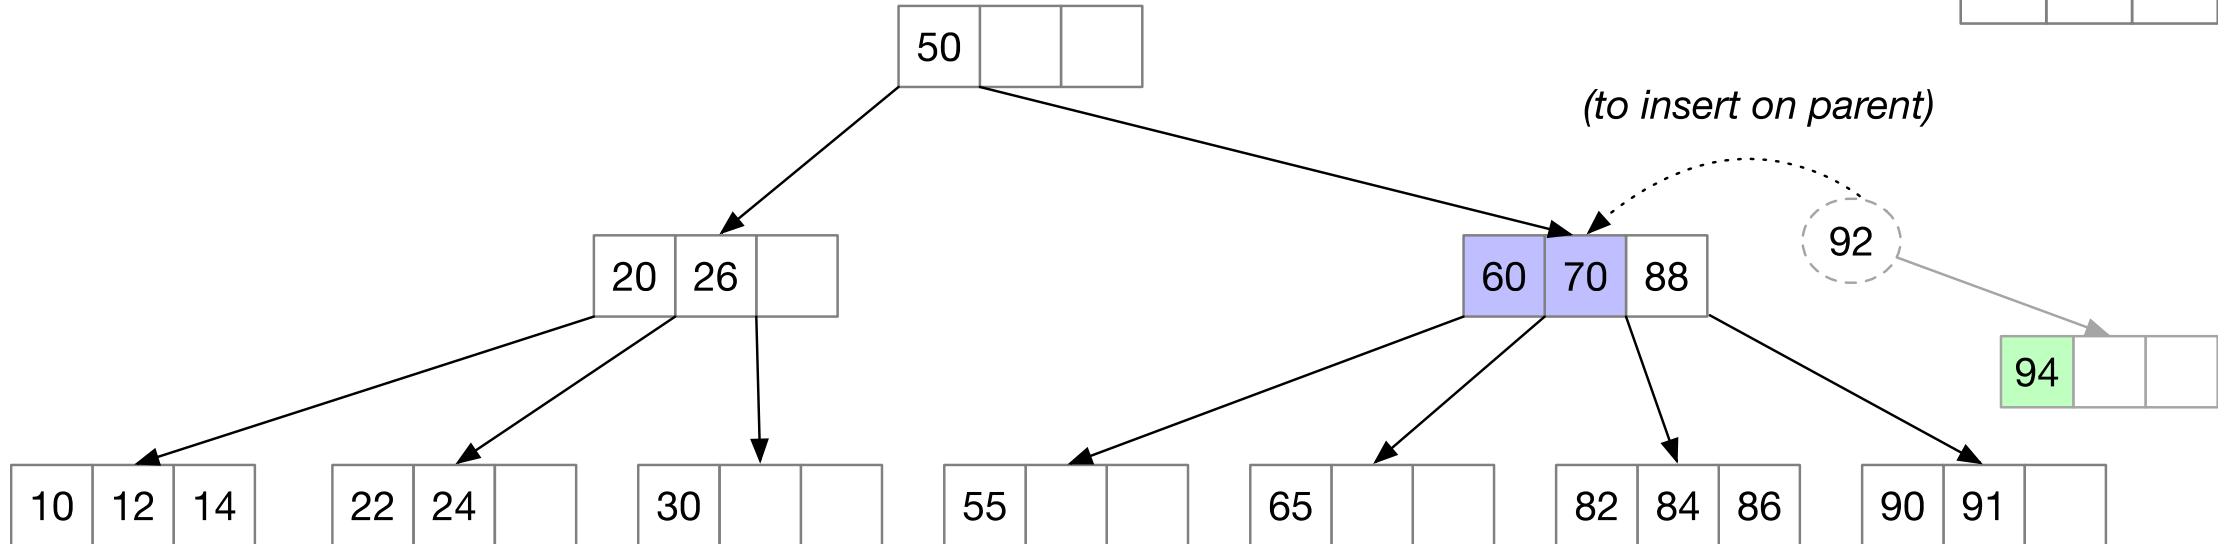
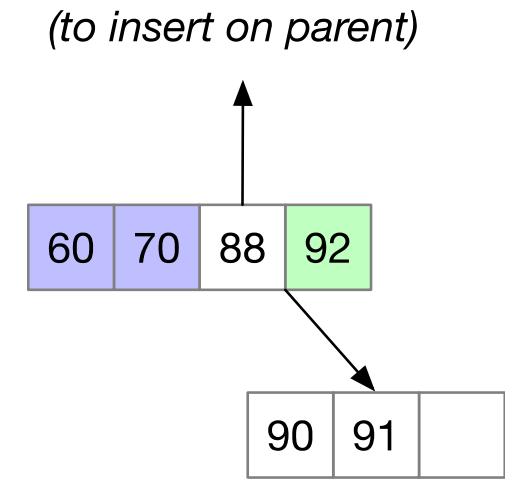
- we search for it, but it's **not found**
- the leaf is **full** => we must **split** it
- so, we **insert 92 on parent**



B-Trees

Now, let's **insert key 91**

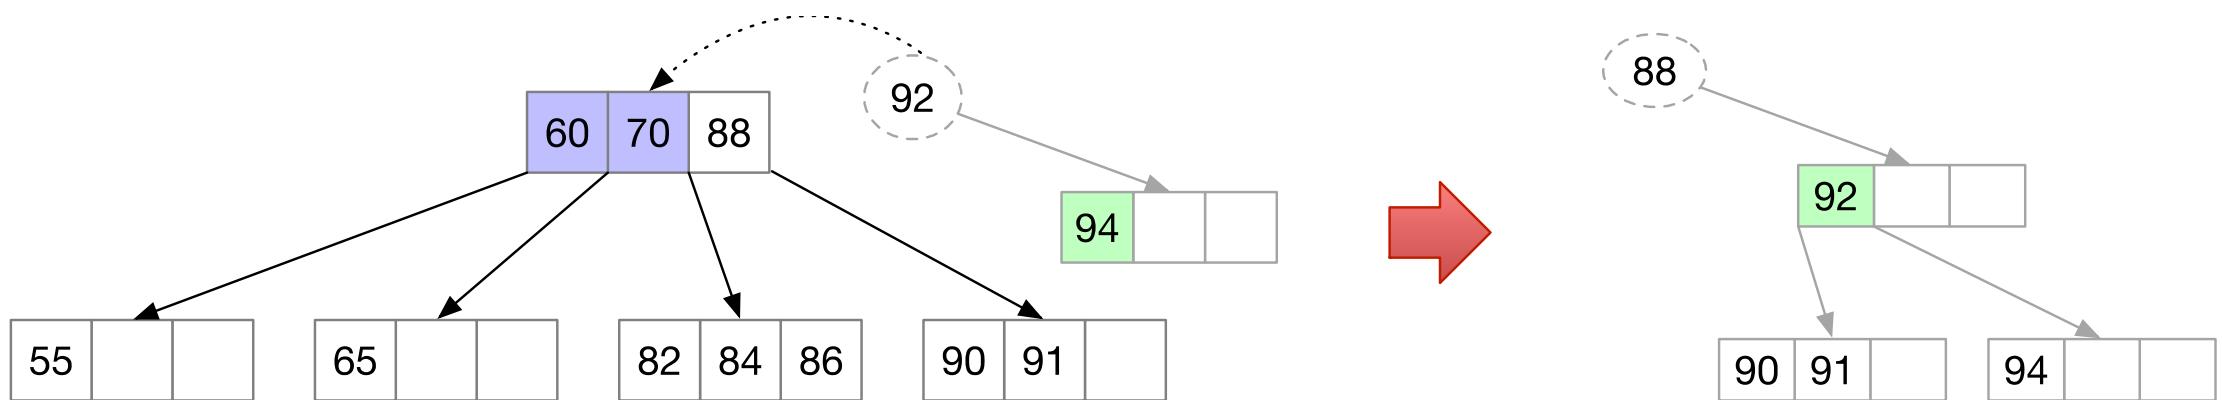
- Let's insert **92 on parent**
 - But the **parent is full** => we **split the node**
 - So, we must **insert 88 in its parent'**



B-Trees

Now, let's **insert key 91**

- Let's insert **92 on parent**
 - But the **parent is full** => we **split** the node
 - So, we must **insert 88 in its parent'**
 - As the right child is the new node, we must do a rearrangement if 88 already has a right node

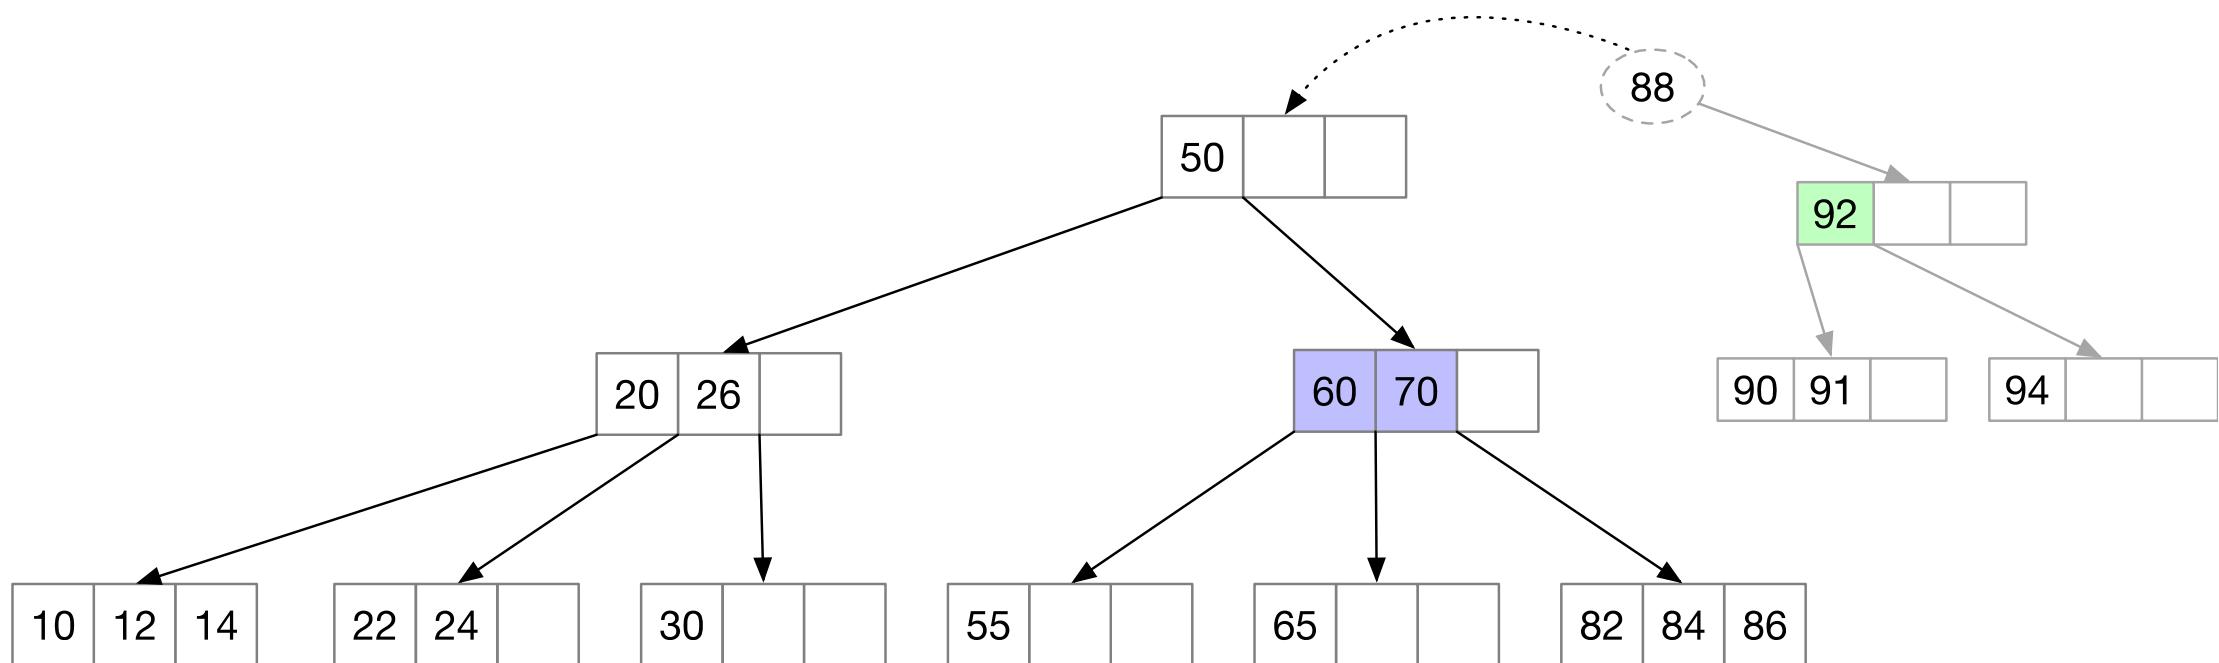


B-Trees

Now, let's **insert key 91**

- Let's insert **92** on parent
 - Let's **insert 88 in its parent'**

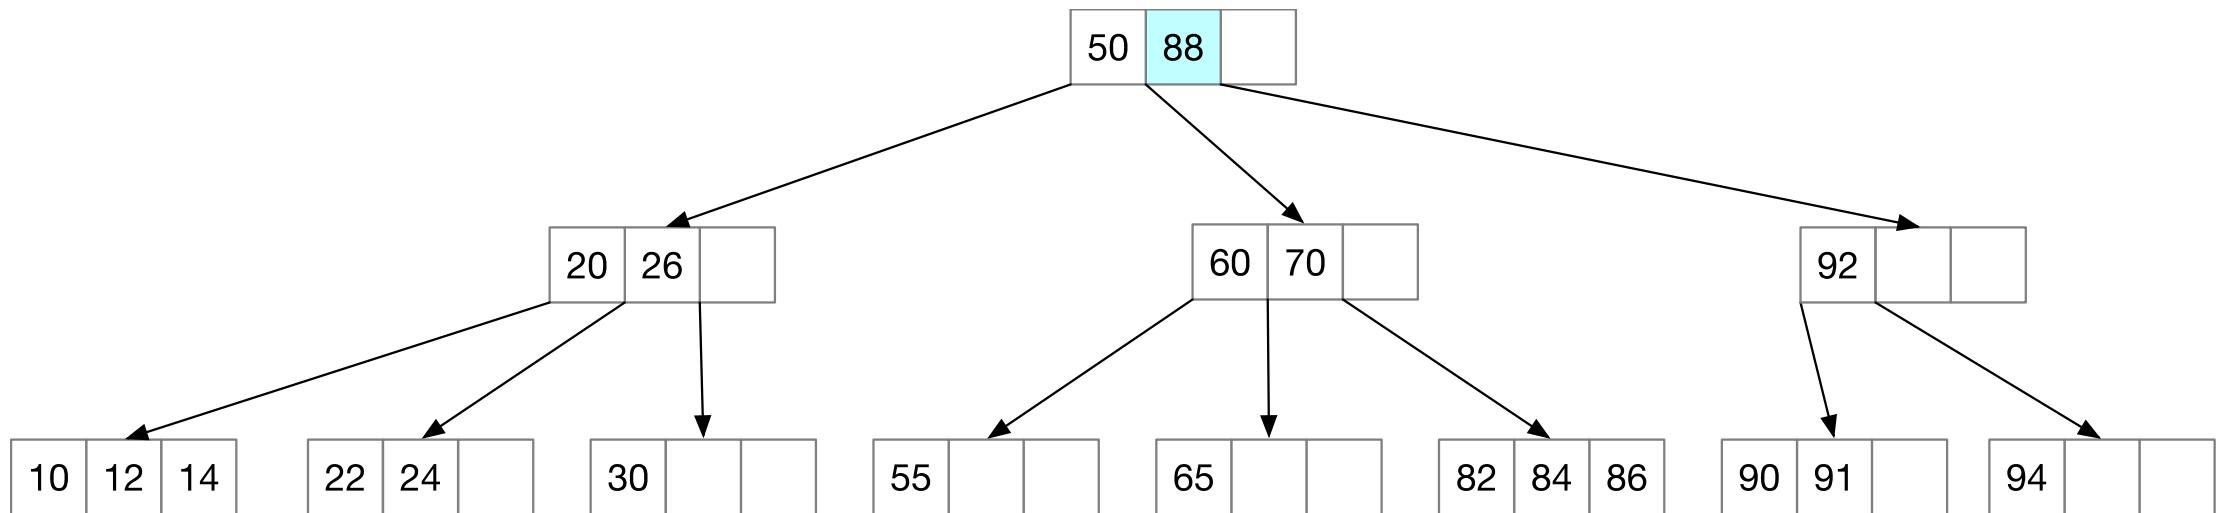
(to insert on parent)



B-Trees

Now, let's **insert key 91**

- Let's insert **92** on parent
 - Let's **insert 88 in its parent'**
 - The node is **not full**, so we **add** to it



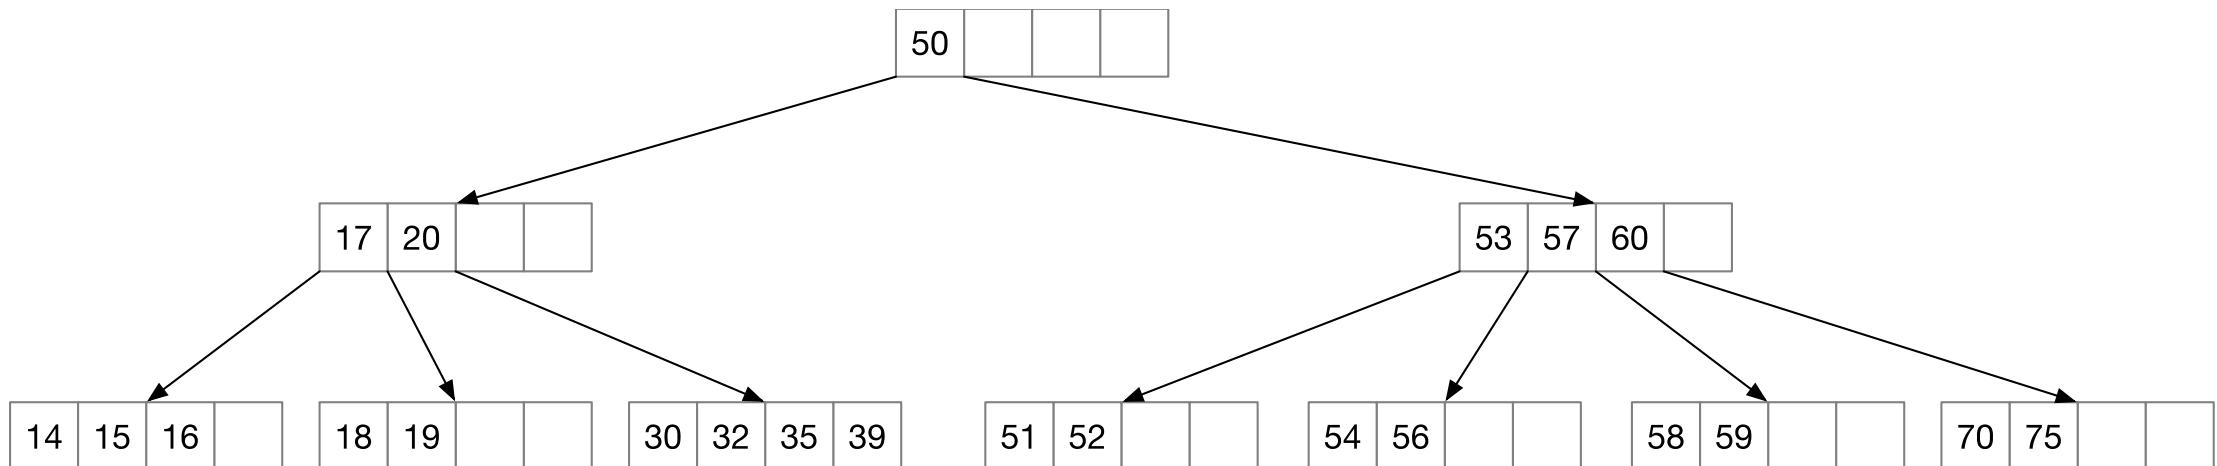
B-Trees

- **Deletion of the key k**
 1. Search the key k in the tree
 2. If the key is found in an **internal node**
 - i. **Substitute** it by the pair (k', v') , where k' is the **next of key k** in ascending order
 - ii. **Delete** (recursively) key k' (from the **leaf**)
 3. If the key is found in a **leaf node**, **remove** it from the leaf
 - i. If $node$ has **at least** $\left\lceil \frac{m}{2} \right\rceil - 1$ keys remaining, we're **finished**
 - ii. If not, but there is an **adjacent sibling** with an **excess ok keys** (i.e. more than $\left\lceil \frac{m}{2} \right\rceil - 1$), then **redistribute keys with sibling and parent**
 - iii. If not, **merge node** with an **adjacent sibling**. This will make the parent lose a child and a key. Return to 3.1 using the parent as $node$

B-Trees

Let's consider the following B-Tree (**order $m = 5$**)

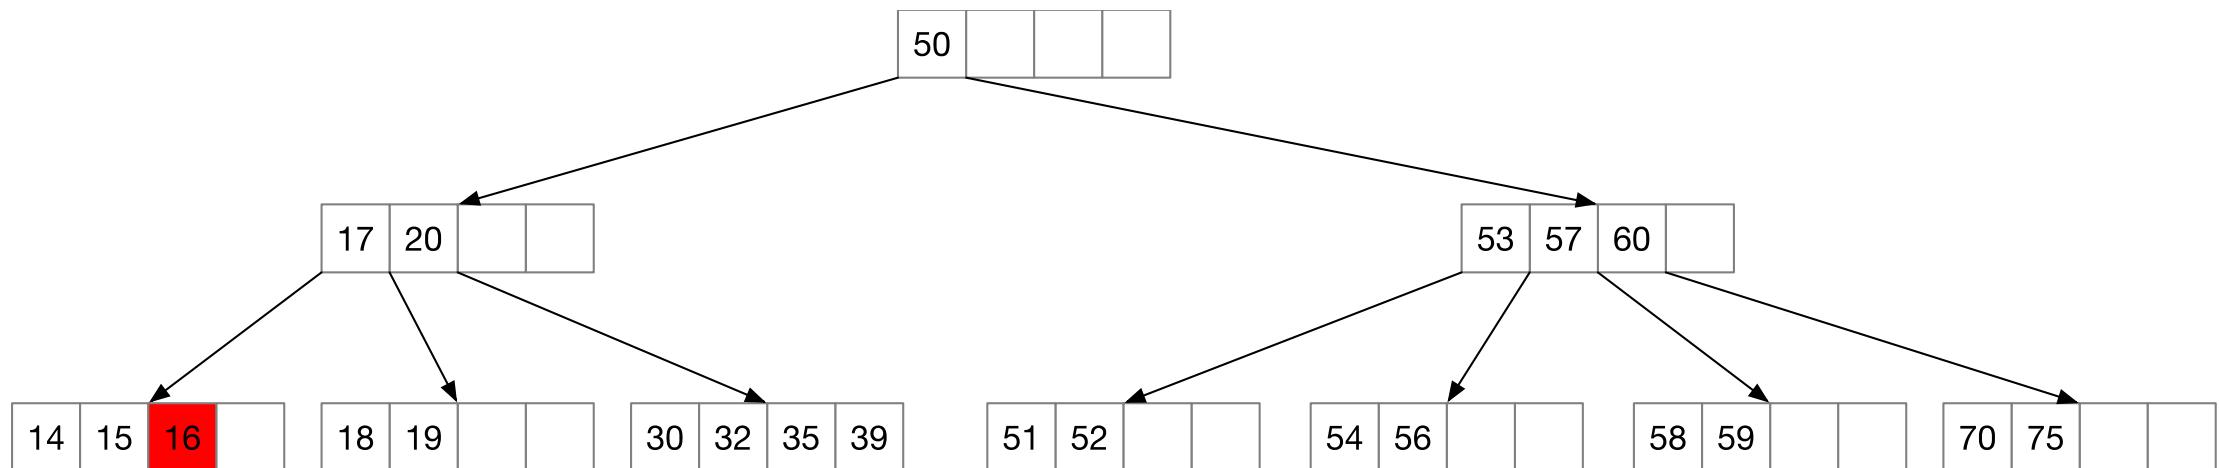
- **Minimum number of keys** = $\left\lceil \frac{5}{2} \right\rceil - 1 = 3 - 1 = 2$
- node with **3 or 4 keys** -> **excess**
- node with **1 key** -> **shortage**



B-Trees

Let's **delete key 16**

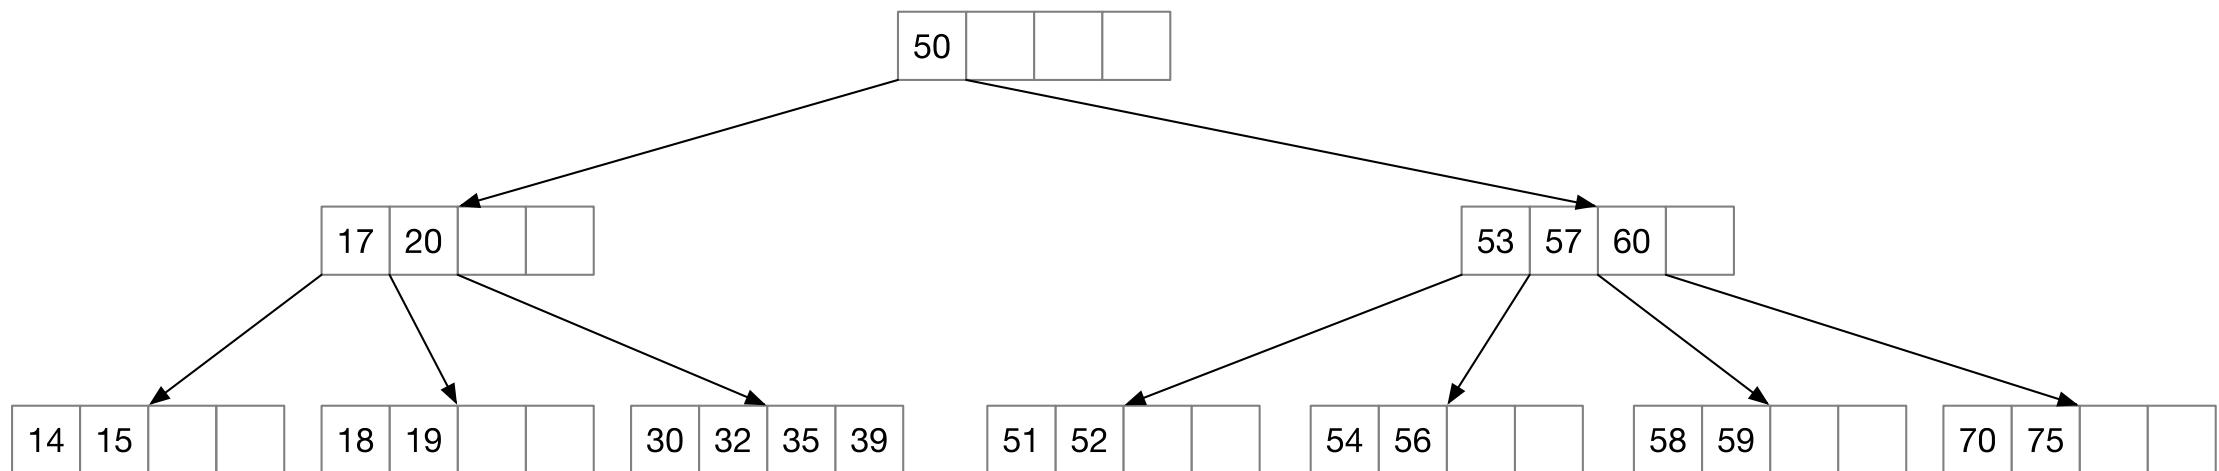
- we find it on a leaf



B-Trees

Let's **delete key 16**

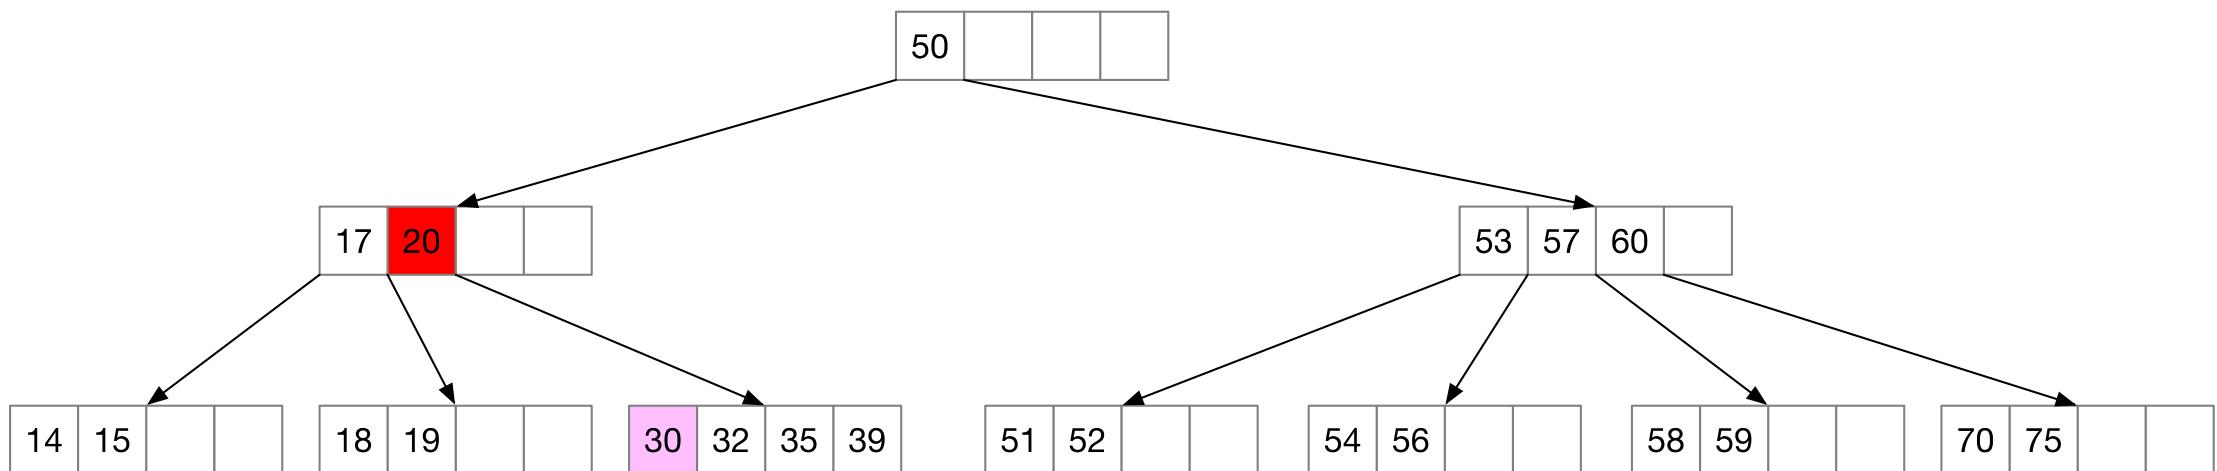
- we find it on a **leaf**
- we **remove** it, but **no shortage**, so we're **finished**



B-Trees

Let's **delete key 20**

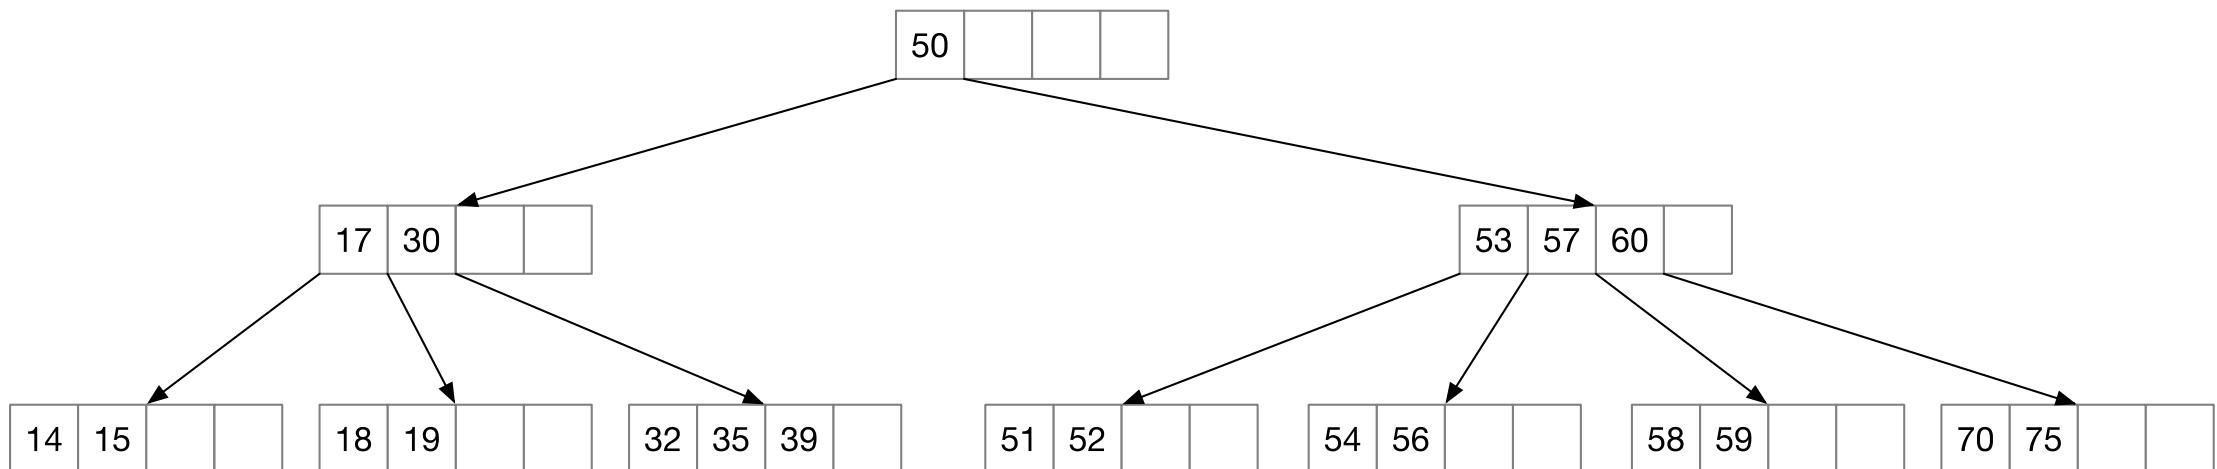
- we find it on an **internal node**
- we **find the next** key, which is 30
 - we **substitute** 20 with 30
 - we **delete** 30 from the **leaf**



B-Trees

Let's **delete key 20**

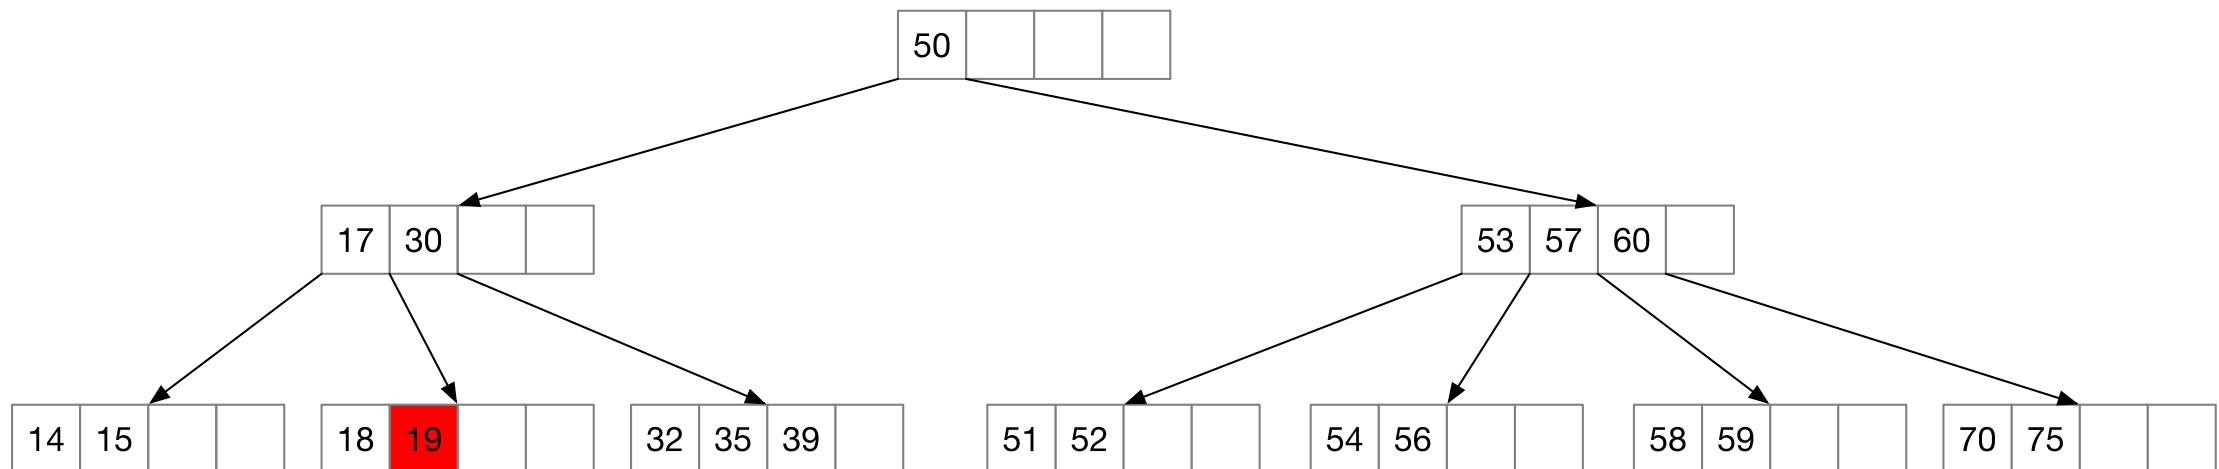
- Let's **remove 30** from the leaf
 - we have **no shortage**, so we're **finished**



B-Trees

Let's **delete key 19**

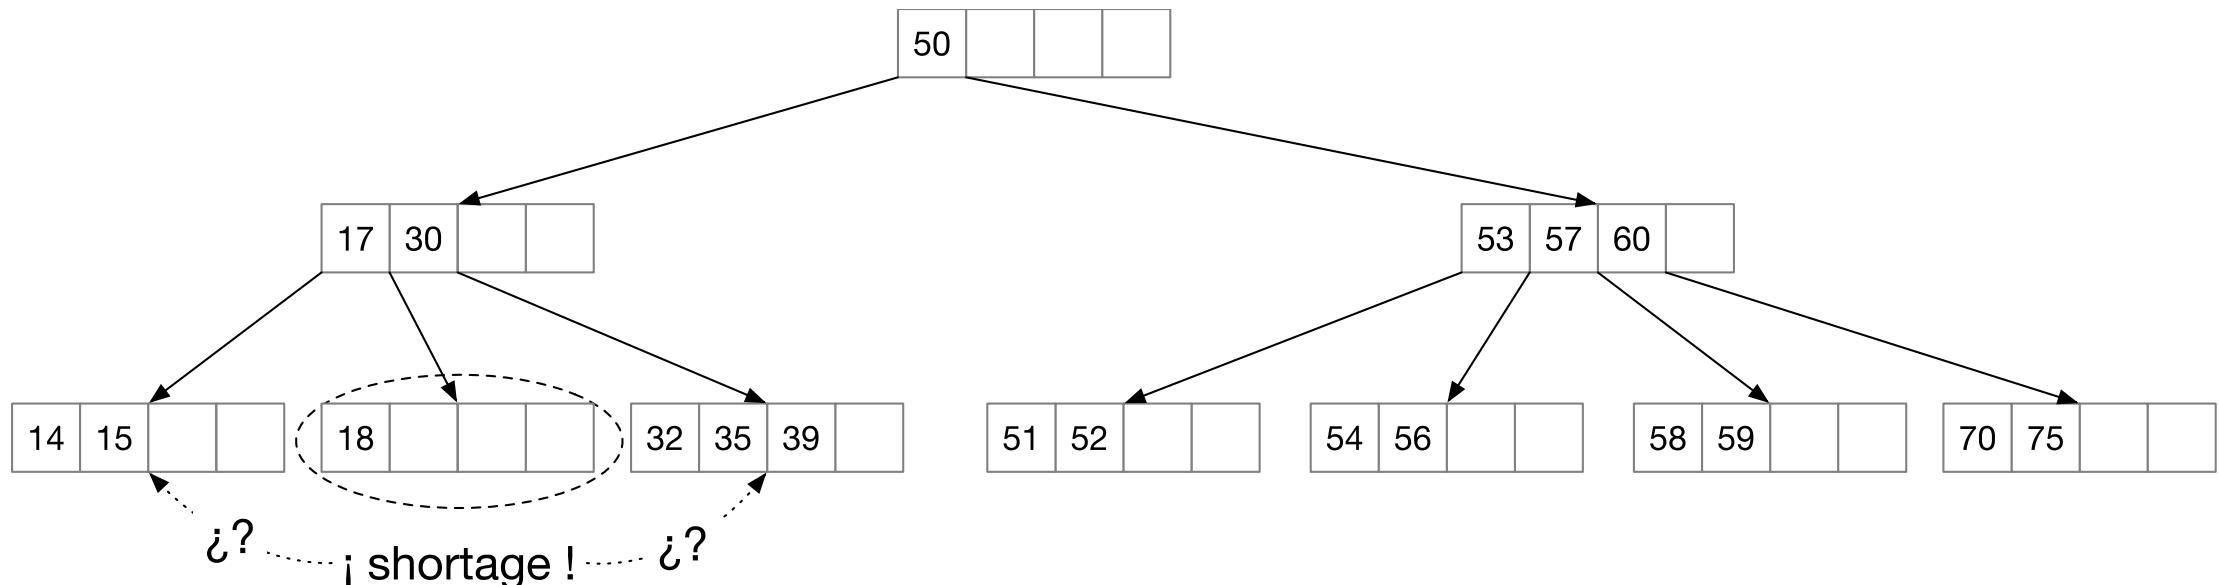
- we find it on a **leaf**, so we **remove** it



B-Trees

Let's **delete key 19**

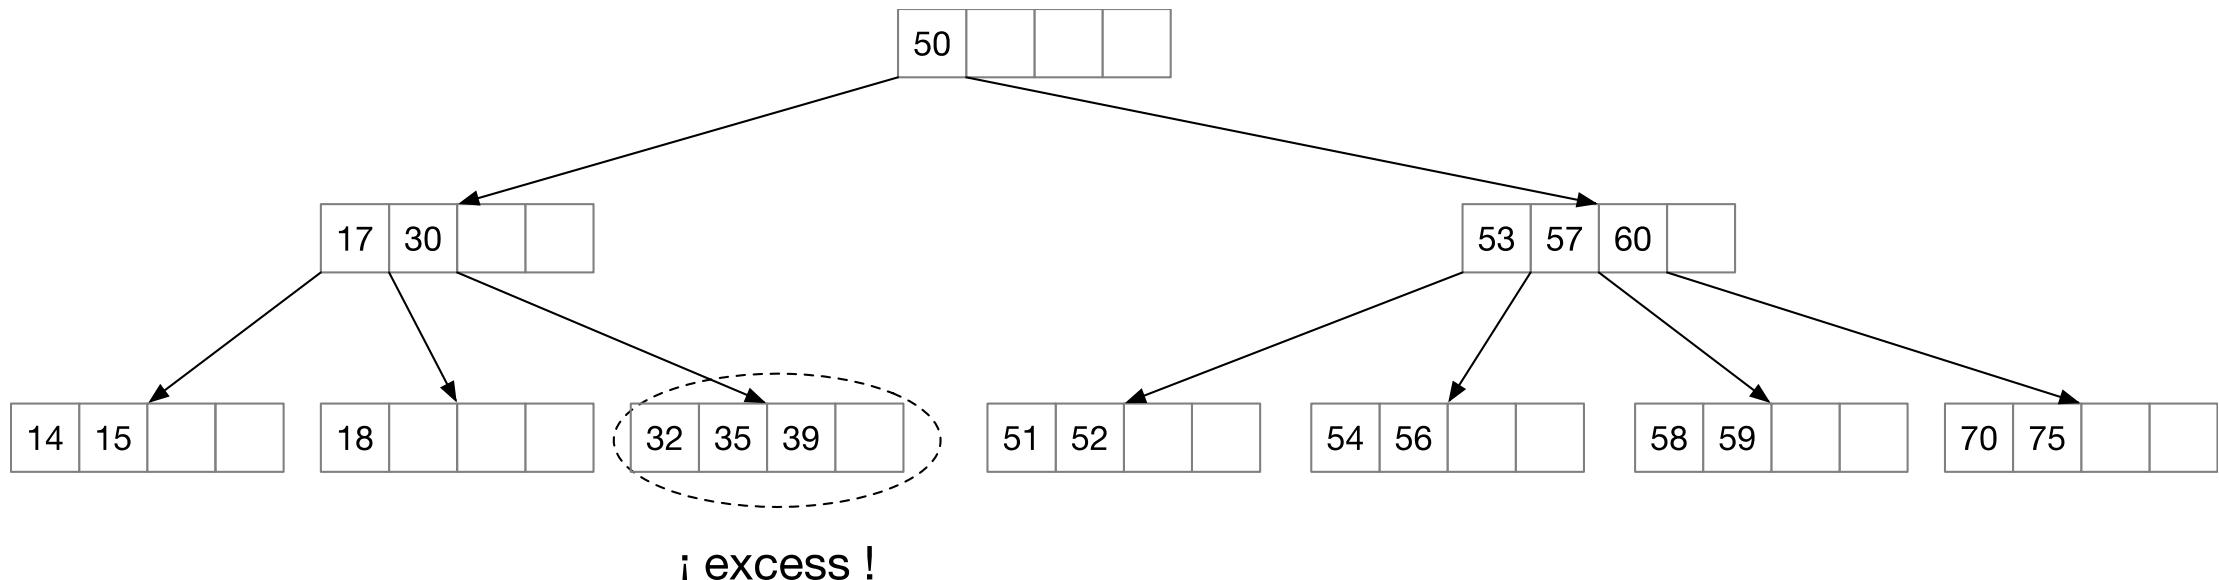
- we find it on a **leaf**, so we **remove it**
- but we incur in **shortage**
 - is there **any adjacent sibling with excess**?



B-Trees

Let's **delete key 19**

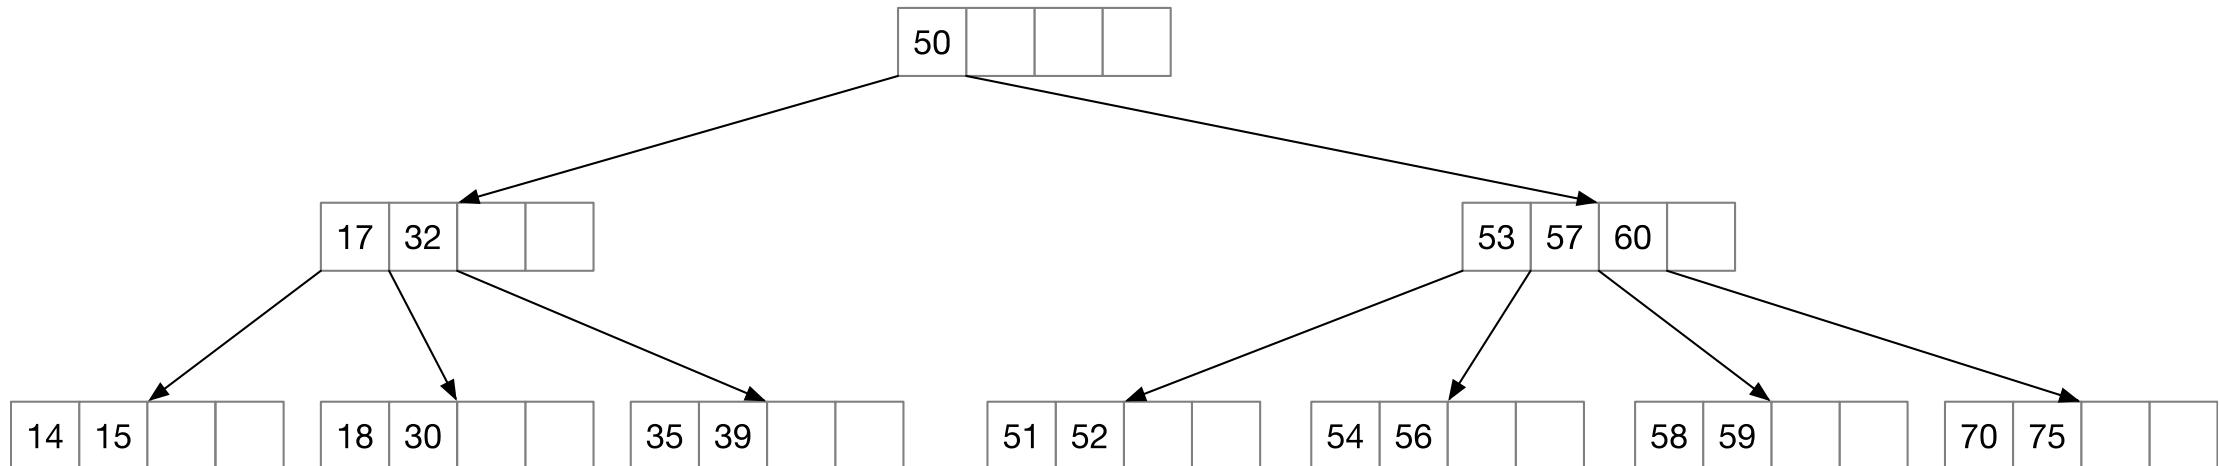
- we find it on a **leaf**, so we remove it
- but we incur in **shortage**
 - is there **any adjacent sibling with excess?**
 - **yes**



B-Trees

Let's **delete key 19**

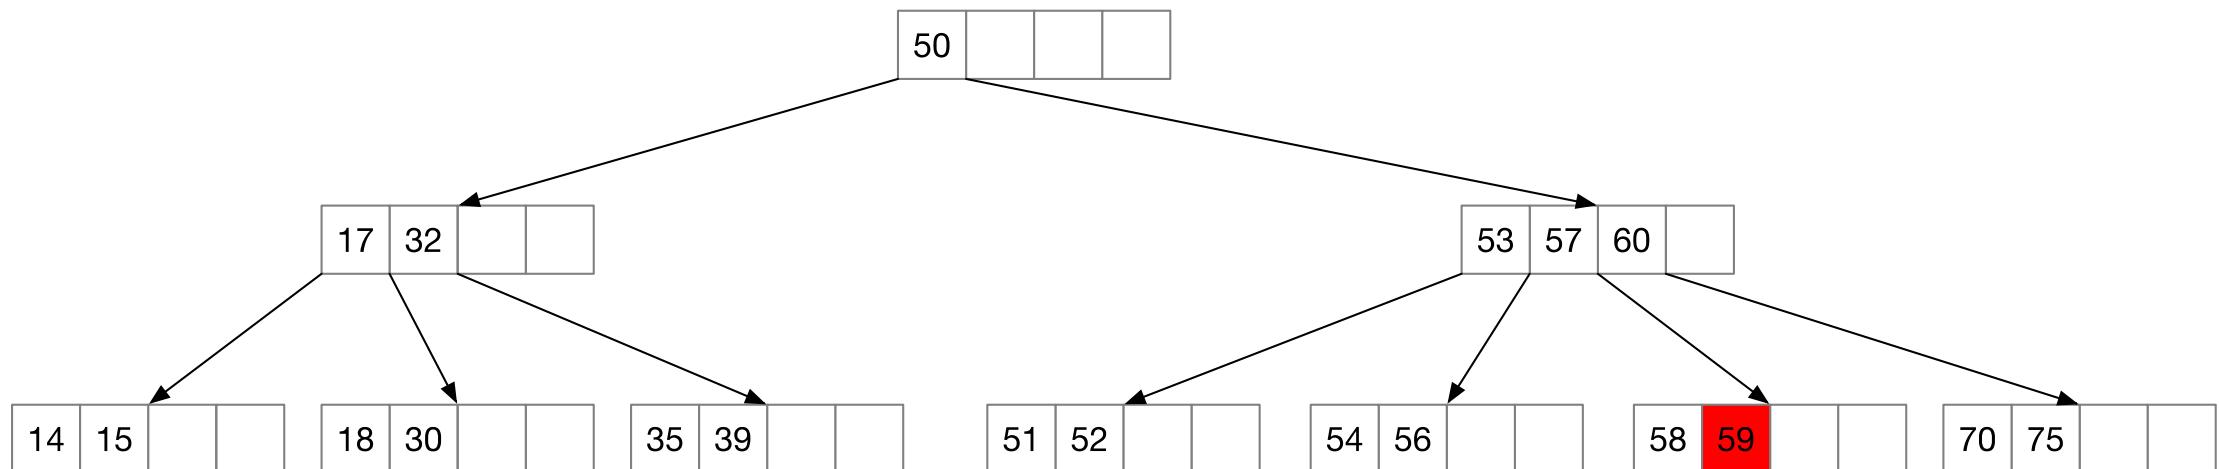
- we find it on a **leaf**, so we remove it
- but we incur in **shortage**
 - is there **any adjacent sibling with excess?**
 - **yes => we redistribute the keys**



B-Trees

Let's **delete key 59**

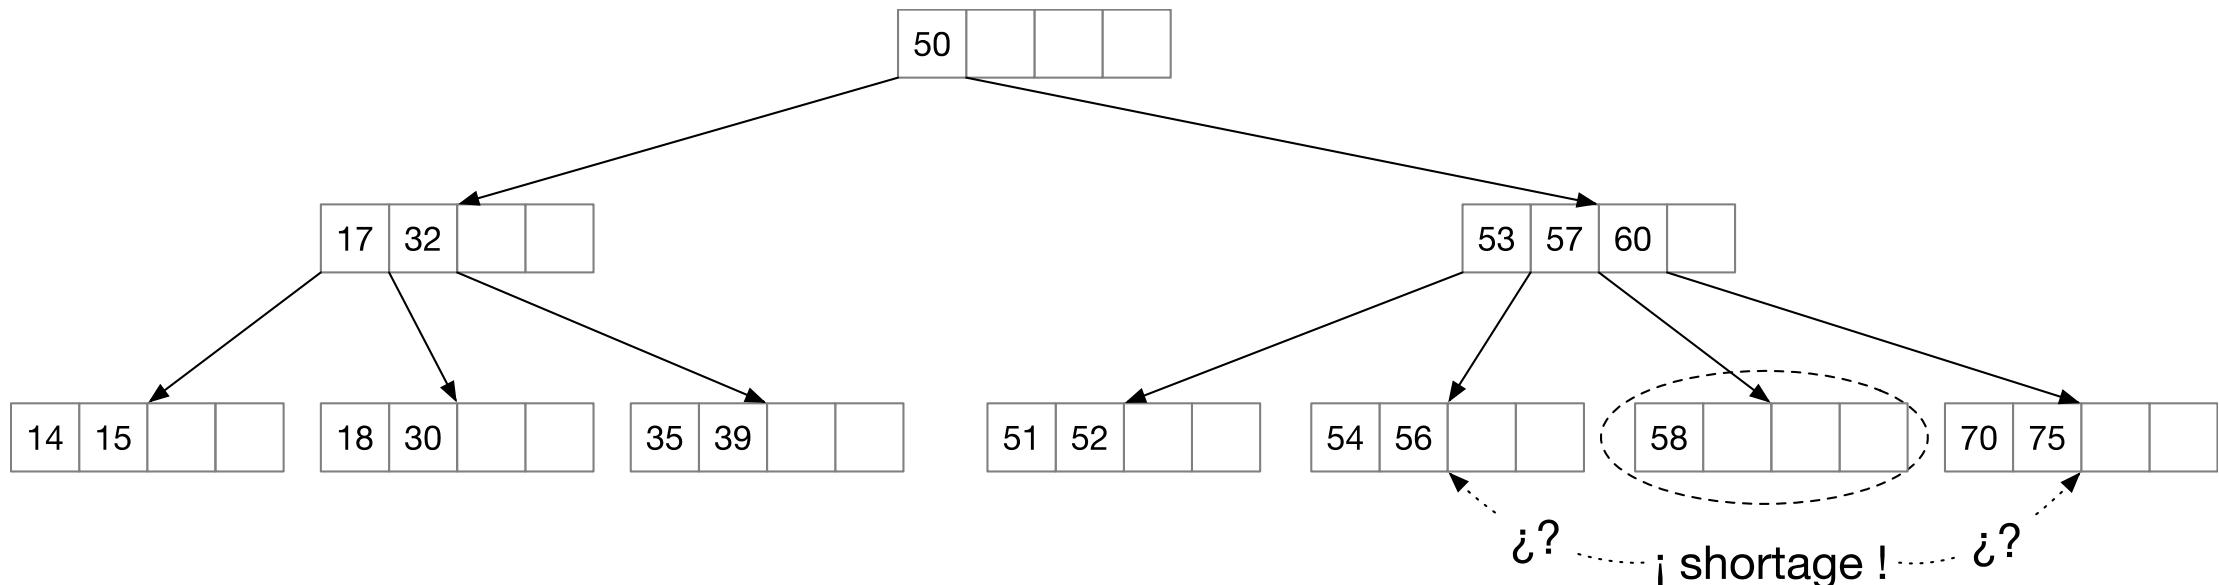
- we search for it, and we find it in a **leaf**



B-Trees

Let's **delete key 59**

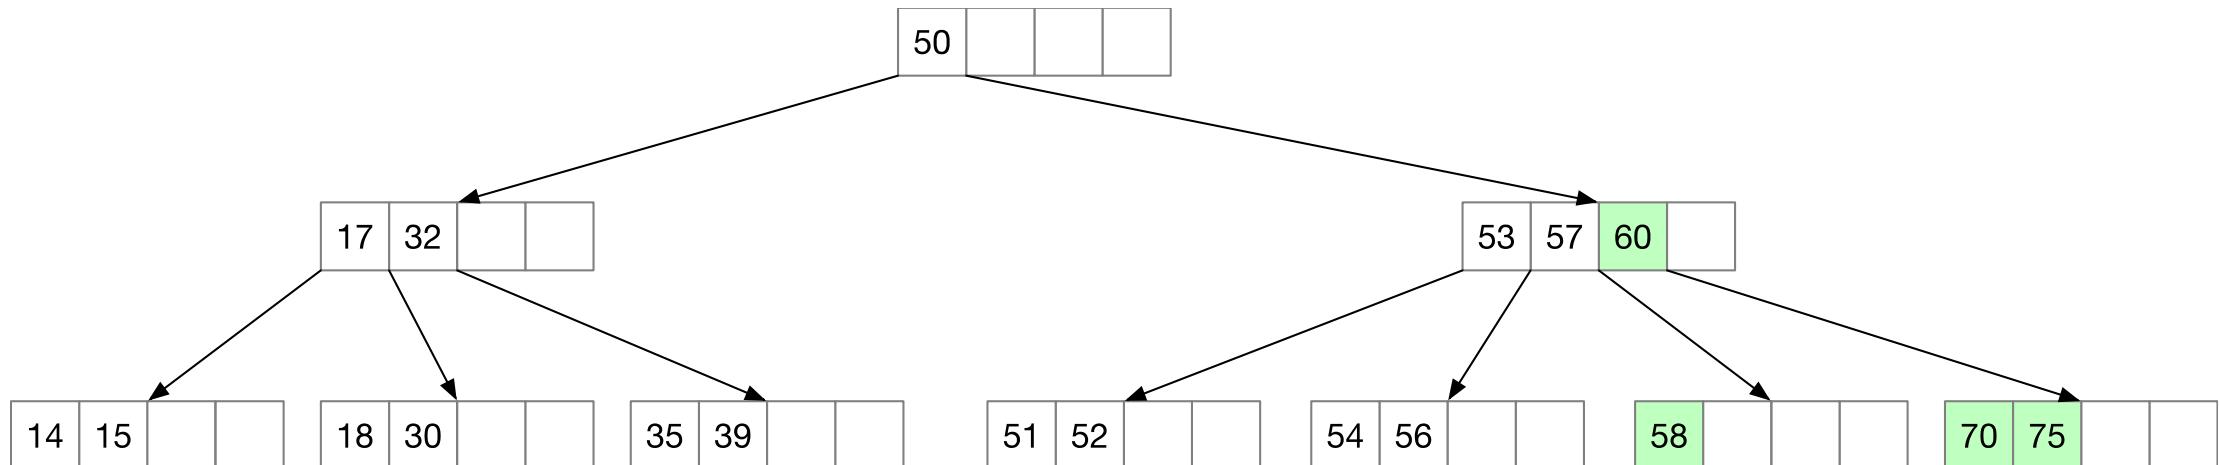
- we search for it, and we find it in a **leaf**
- we **remove** it but we incur in **shortage**
- is there any **adjacent sibling with excess**?
- **no**



B-Trees

Let's **delete key 59**

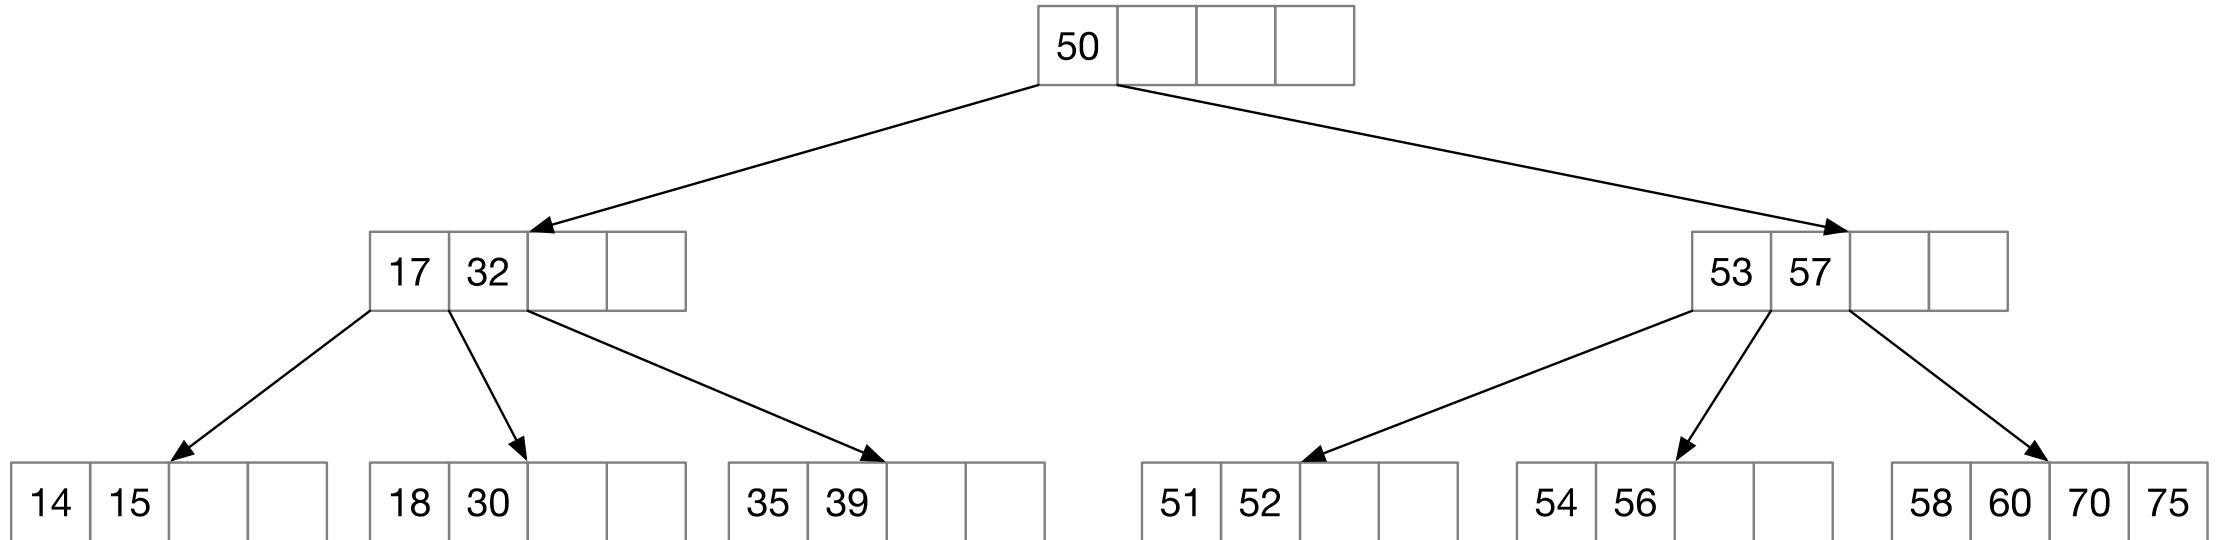
- we search for it, and we find it in a **leaf**
- we **remove** it but we incur in **shortage**
- is there **any adjacent sibling with excess?**
- **no =>** we must do a **merge**



B-Trees

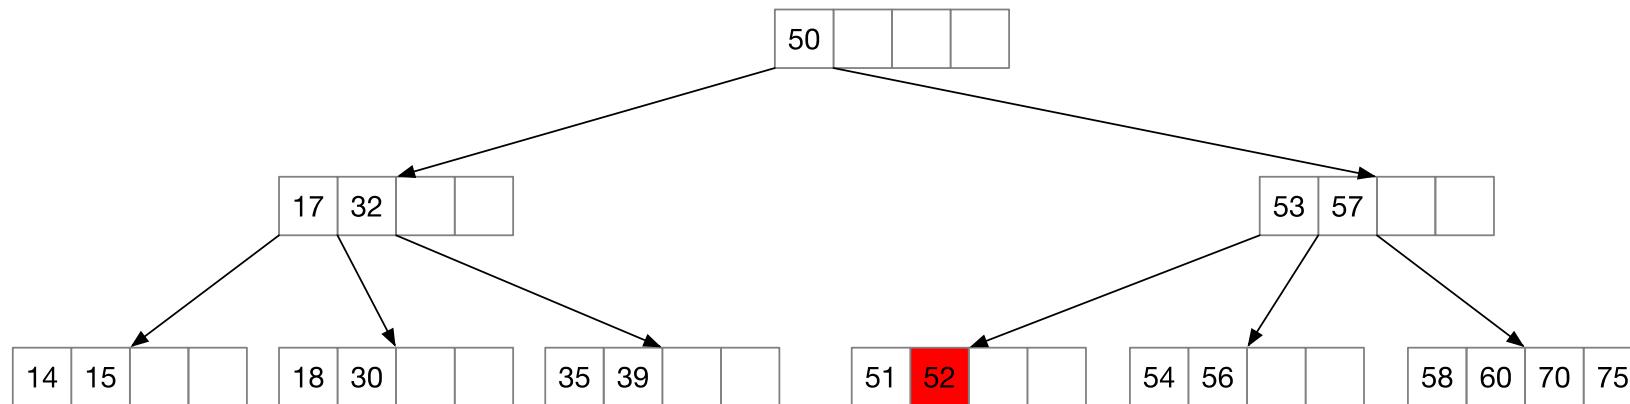
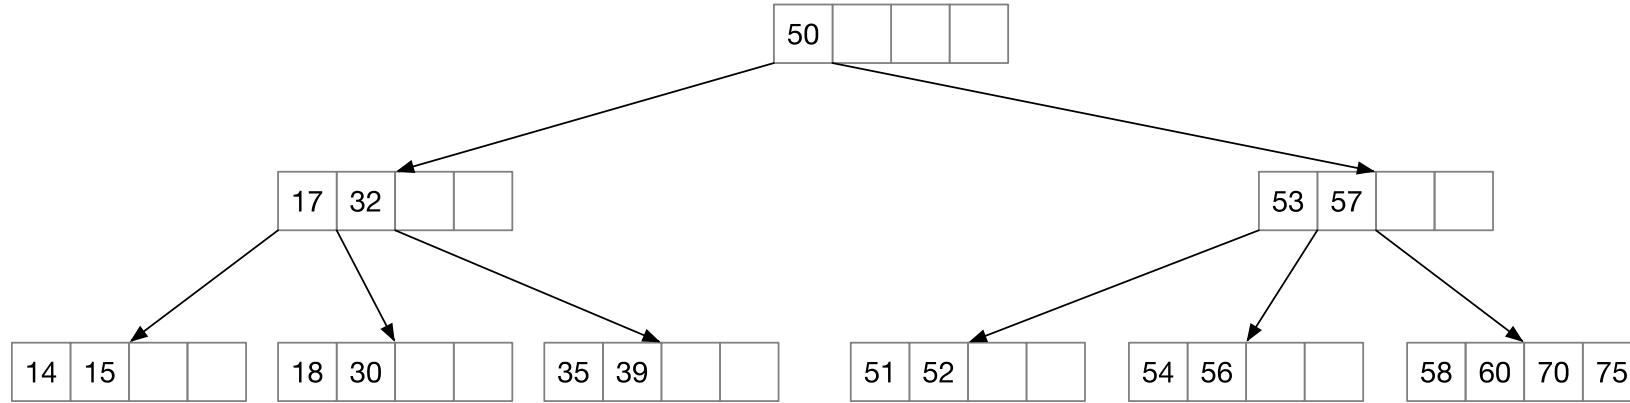
Let's **delete key 59**

- we search for it, and we find it in a **leaf**
- we **remove** it but we incur in **shortage**
- is there **any adjacent sibling with excess**?
- **no** => we must do a **merge**
- the **parent loses a key** but has **no shortage**, we're **finished**



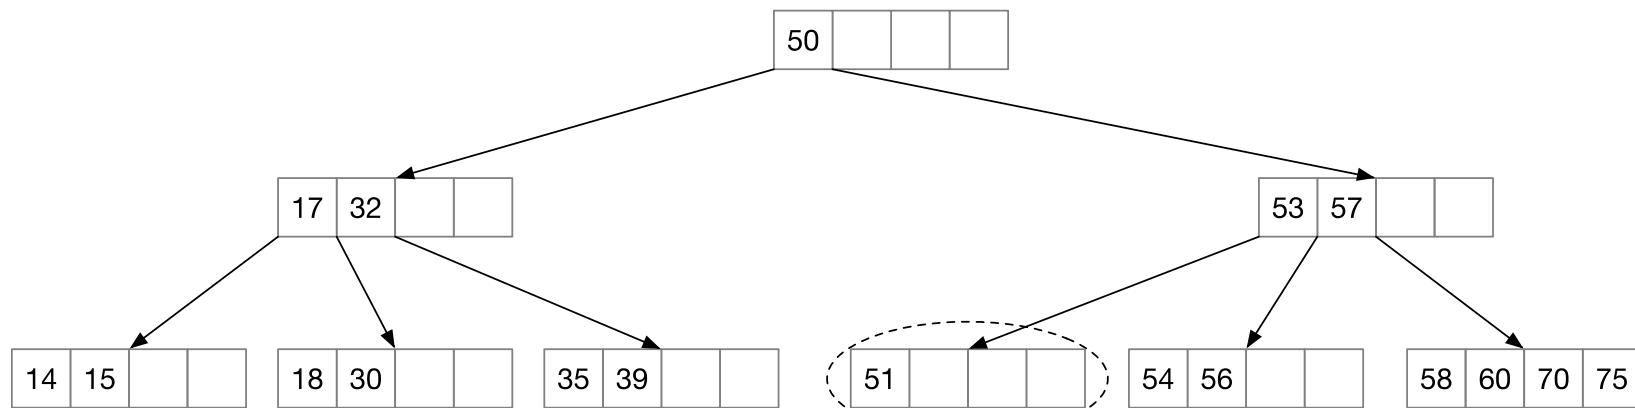
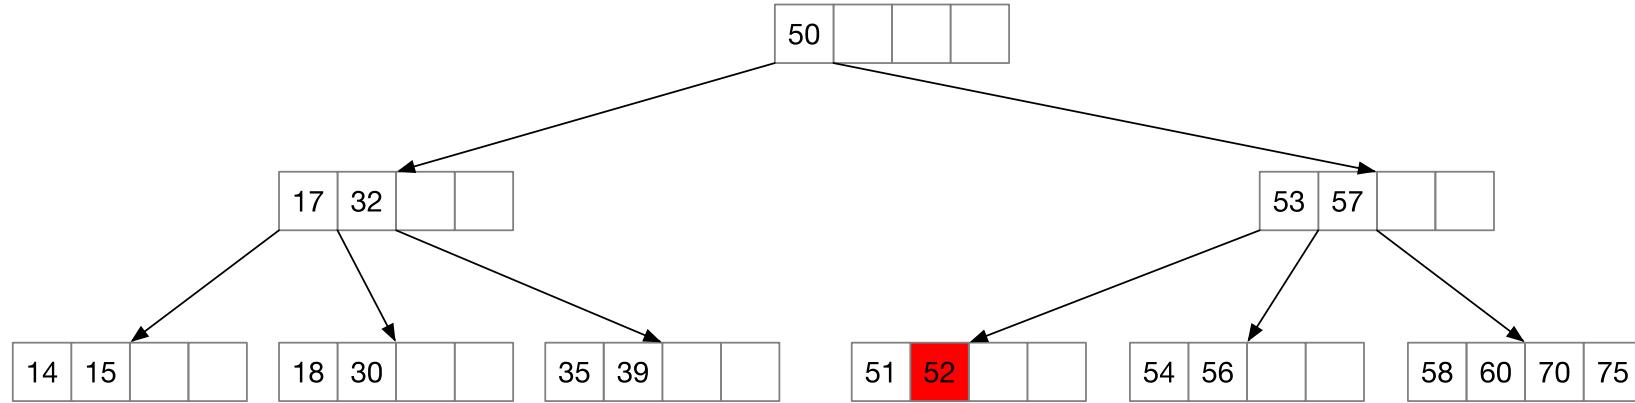
B-Trees

Let's **delete key 52** (now without subtitles)



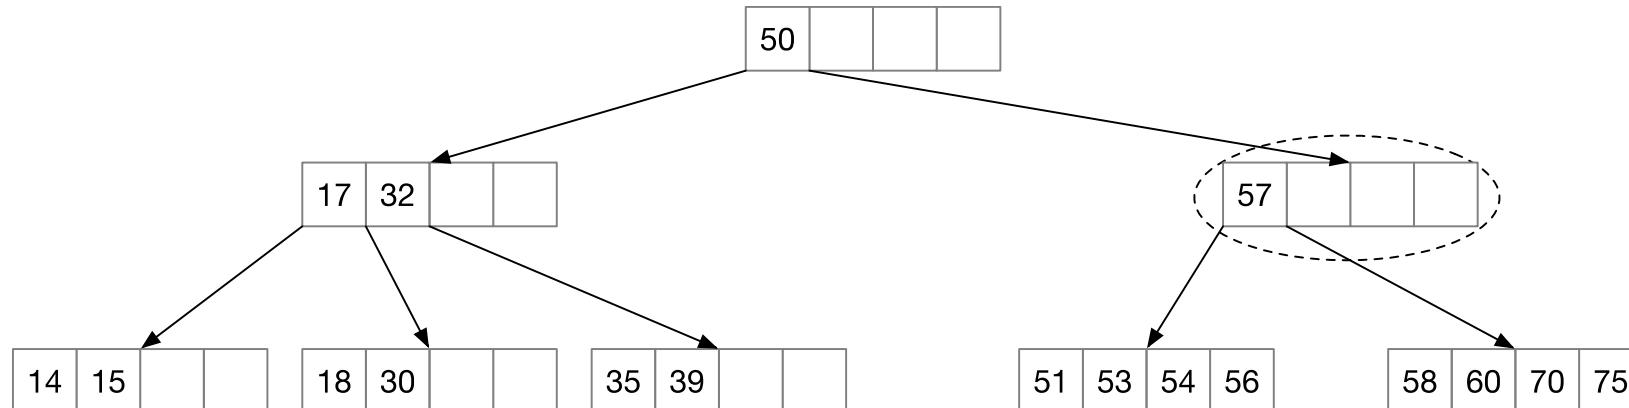
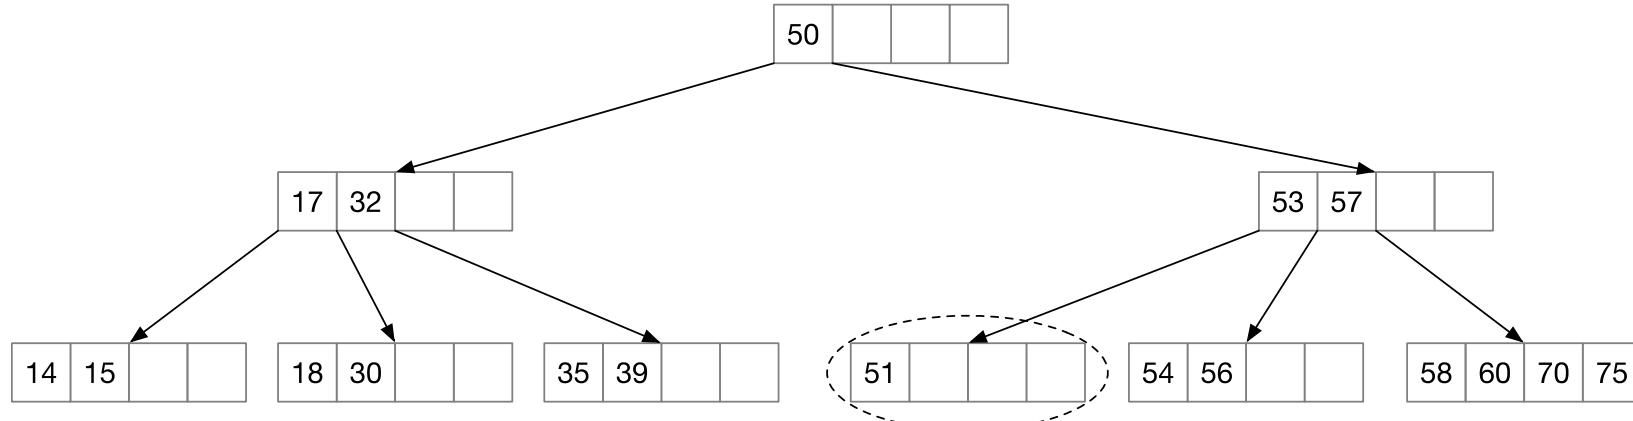
B-Trees

Let's **delete key 52** (now without subtitles)



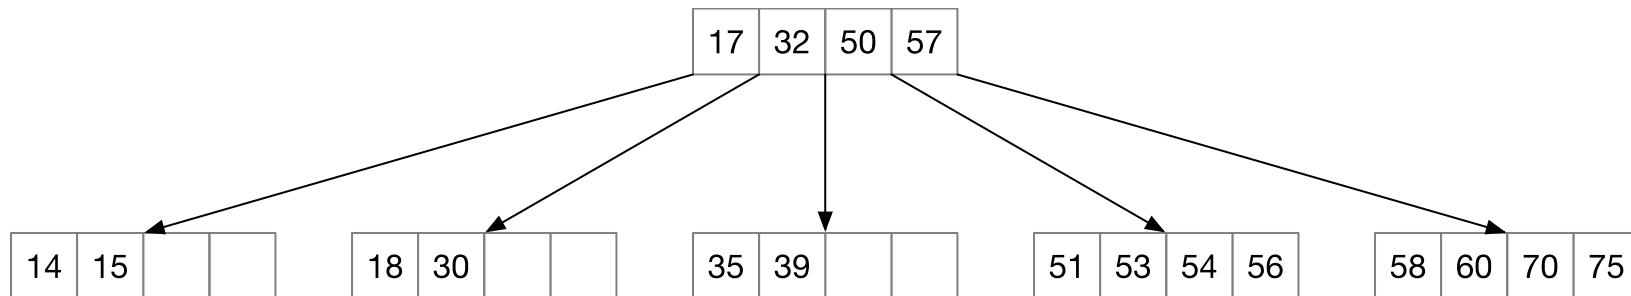
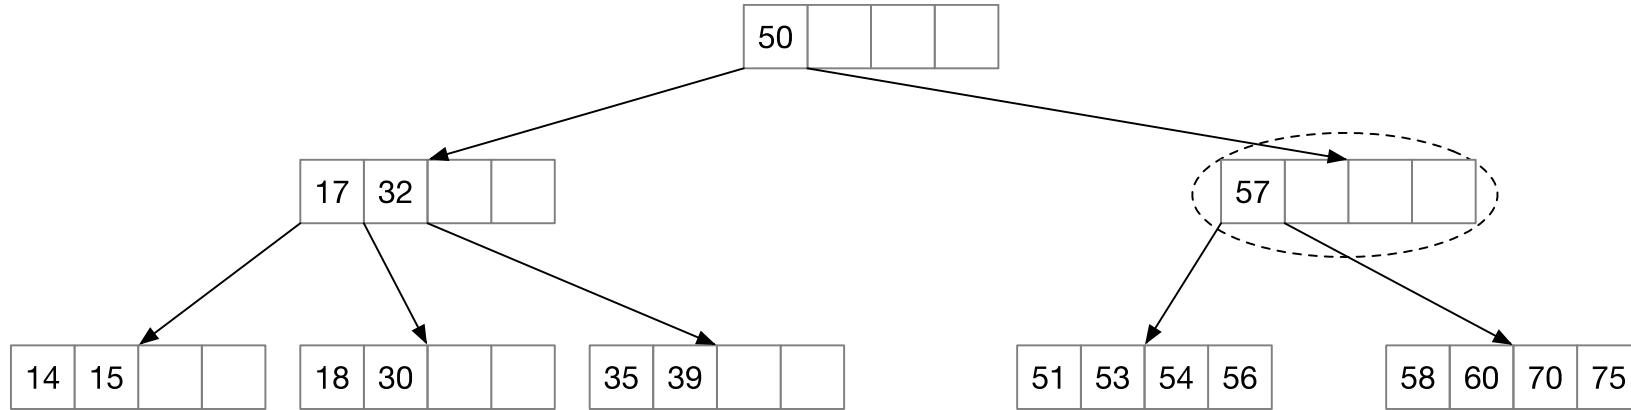
B-Trees

Let's **delete key 52** (now without subtitles)



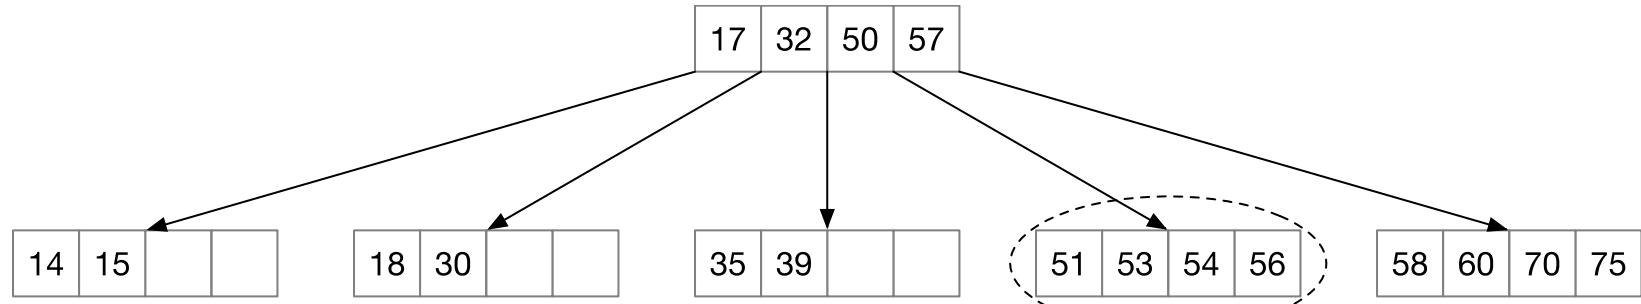
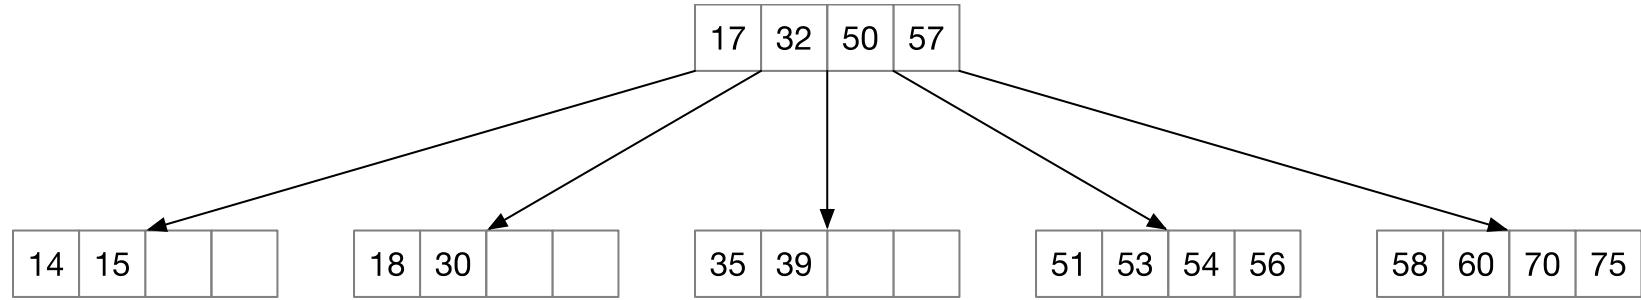
B-Trees

Let's **delete key 52** (now without subtitles)



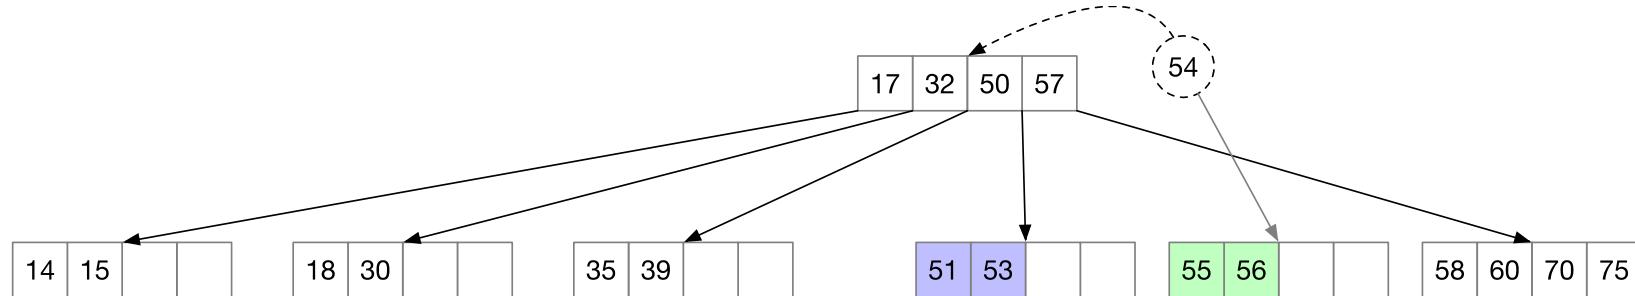
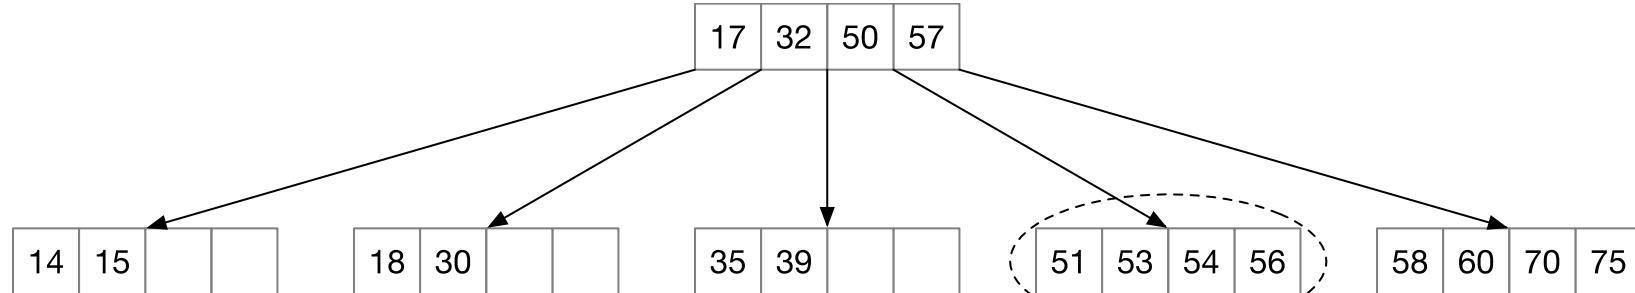
B-Trees

Let's **insert key 55** (now without subtitles)



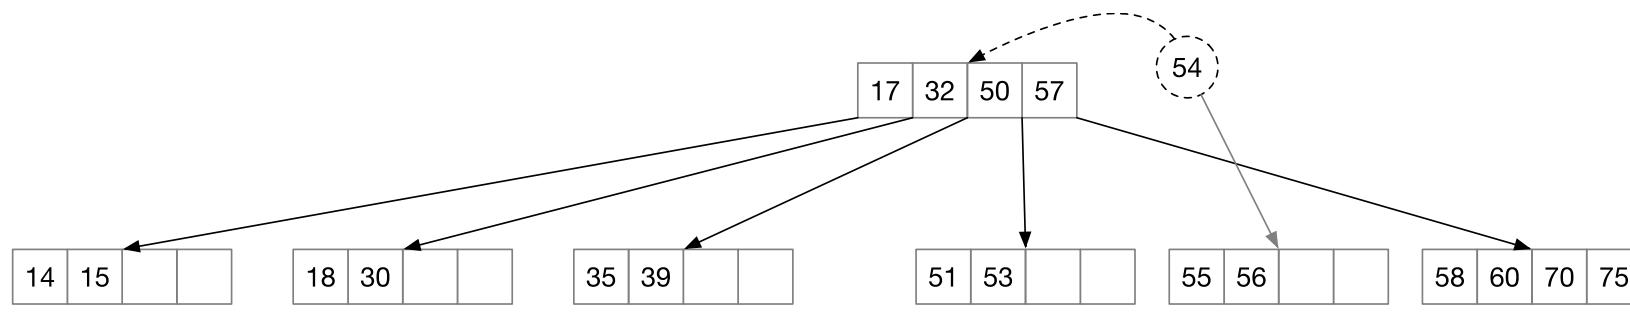
B-Trees

Let's **insert key 55** (now without subtitles)

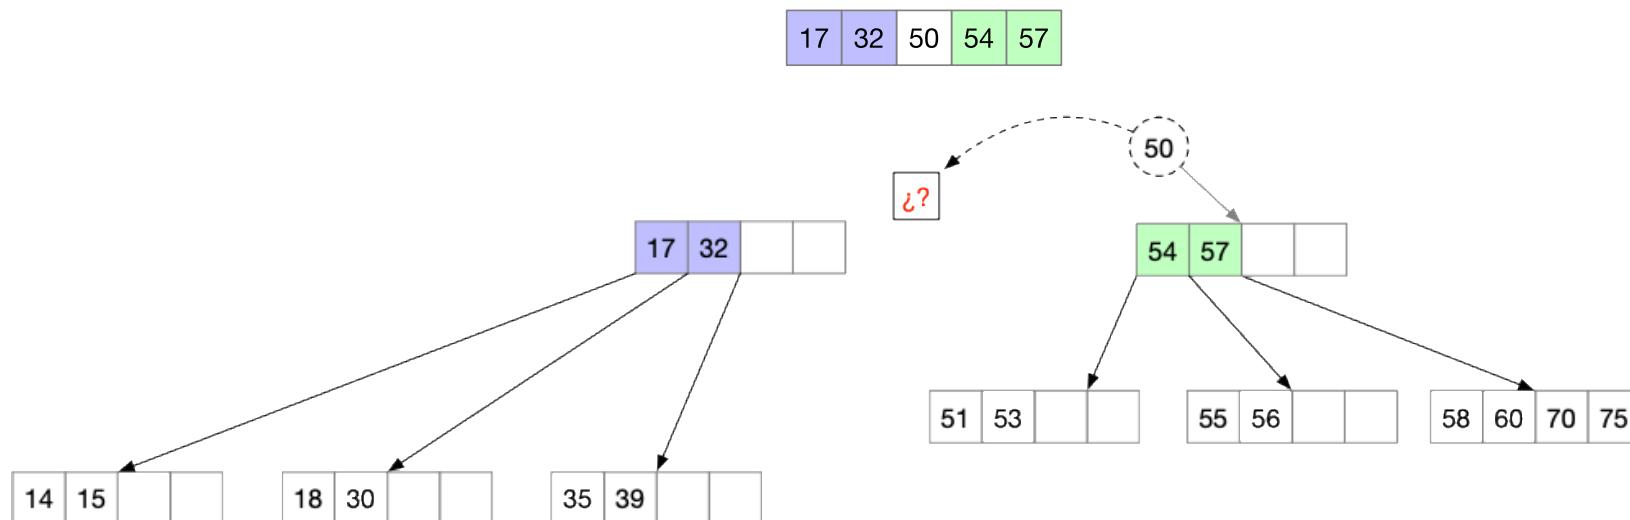


B-Trees

Let's **insert key 55** (now without subtitles)



| | | | | |
|----|----|----|----|----|
| 17 | 32 | 50 | 54 | 57 |
|----|----|----|----|----|



B-Trees

Let's **insert key 55** (now without subtitles)

